A Biomolecular Implementation of Joscha Bach's Mechanistic PSI Architecture in a Multi-Organoid Autonomous System

Augmented Research & Technical Writing From: D.M.M. , AI Assisted: Clause Sonnet 4.5

1. Introduction

1.1 Theoretical Foundations: Bach's Psi Architecture

The development of mechanistic theories capable of explaining cognition, consciousness, and autonomous agency represents one of the central challenges in contemporary cognitive science and artificial intelligence research. While numerous theoretical frameworks have been proposed to account for specific aspects of intelligent behavior, few attempt comprehensive integration spanning perception, action, motivation, memory, planning, and metacognition within a unified architectural specification. Joscha Bach's Psi theory stands as a notable exception, providing a detailed mechanistic account of how these diverse cognitive functions emerge from the interaction of specific computational components organized according to precise architectural principles.

Bach developed the Psi architecture through systematic analysis of the functional requirements for autonomous agency combined with detailed consideration of neurobiological implementation constraints observed in natural cognitive systems. The theoretical framework was first articulated in Bach's foundational work "Principles of Synthetic Intelligence" and has been subsequently refined through computational implementations and theoretical elaborations exploring its implications for understanding biological cognition and developing artificial intelligence systems. The architecture derives its name from the Greek letter Psi (Ψ) symbolizing the mind or soul in philosophical

discourse, chosen to emphasize the framework's ambition to provide a complete mechanistic account of mental phenomena traditionally considered resistant to reductionist explanation.

At the core of Psi theory lies the recognition that intelligent agents must solve several fundamental problems simultaneously. First, they must extract meaningful information from high-dimensional sensory streams, identifying relevant patterns while filtering noise and irrelevant variation. Second, they must generate appropriate actions in response to perceived circumstances, selecting from potentially vast action spaces based on predicted outcomes. Third, they must manage multiple competing needs and goals, allocating limited cognitive and physical resources to maximize overall success. Fourth, they must learn from experience, adjusting their perceptual interpretations, action policies, and goal priorities based on outcomes. Fifth, they must plan ahead, simulating potential futures to select actions whose consequences extend beyond immediate returns. Sixth, they must monitor their own cognitive processes, detecting errors, estimating uncertainty, and adaptively modulating how they think based on metacognitive assessment.

The Psi architecture addresses these challenges through six major subsystems whose interaction generates coherent intelligent behavior. The perceptual system implements a hierarchical processing stream that progressively extracts increasingly abstract representations from sensory input. Sensory signals first undergo feature extraction identifying local patterns such as oriented edges in vision or spectral components in audition. These features are then combined into increasingly complex representations through successive processing stages, with each level encoding information at coarser spatial and temporal resolutions while capturing larger-scale patterns and invariances. The hierarchical organization enables both bottom-up processing where sensory evidence drives perceptual interpretations and top-down processing where higher-level expectations modulate lower-level feature extraction, implementing a form of predictive coding where the system constantly generates predictions about expected sensory input and computes prediction errors reflecting the mismatch between expectation and reality.

The action system mirrors the perceptual hierarchy in reverse, implementing a motor control stream that progressively refines abstract action goals into concrete motor

commands. High-level planning systems generate abstract action intentions such as "move to location X" or "manipulate object Y" without specifying precise motor details. These abstract intentions propagate through successively lower levels that decompose goals into motor primitives, specify kinematic parameters like forces and velocities, and finally generate detailed muscle activation patterns. This hierarchical organization enables flexible behavior where the same high-level goal can be achieved through multiple motor implementations depending on context, and where low-level motor execution can proceed with minimal high-level oversight once appropriate programs are initiated.

The motivational system implements what Bach terms "urges," which are continuous regulatory signals that modulate the relative priority of different behavioral goals. Bach identifies several fundamental urge categories that any autonomous agent must manage. Physiological urges reflect bodily needs such as energy acquisition, physical integrity maintenance, and homeostatic regulation of temperature, hydration, and other physiological parameters. These urges generate continuously increasing pressure as underlying needs accumulate, creating escalating motivation to engage in behaviors that satisfy the need. For instance, as cellular energy stores deplete and glucose levels decline, energy urges intensify, biasing the action selection system toward energy-seeking behaviors until the need is satisfied and the urge declines.

Cognitive urges reflect information-processing needs that do not correspond to specific physiological deficits but nonetheless require satisfaction for effective functioning. The certainty reduction urge, also conceptualized as curiosity or information-seeking drive, creates motivation to reduce prediction errors and resolve uncertainty about the environment. When the system encounters novel situations where its predictive models generate large errors, the certainty reduction urge intensifies, driving exploratory behavior and attention toward information sources that might reduce uncertainty. The competence urge creates motivation to improve skill and capability, generating satisfaction from learning and performance improvement even in the absence of external rewards. This intrinsic motivation for mastery drives practice and skill refinement beyond what would be expected from purely outcome-based reinforcement.

Social urges, though not directly relevant to the isolated system we implement here, represent important components of Bach's full framework. These include affiliation urges driving social bonding and cooperation, dominance urges related to social status and competitive interactions, and recognition urges motivating behaviors that gain approval and status from social partners. While our current single-system implementation does not engage these social dimensions, the architectural framework accommodates their incorporation in multi-agent contexts.

The critical innovation in Bach's treatment of motivation is the recognition that urges must compete and interact to generate coherent behavior. At any given moment, multiple urges may be active simultaneously, each pushing toward different behavioral goals that may be mutually incompatible. The system must therefore implement mechanisms for urge integration and conflict resolution, combining multiple motivational signals to determine overall action values. Bach proposes that action selection occurs through a process of "urge algebra" where the value of each potential action is computed as a weighted combination of the satisfaction it provides to each active urge, with weights reflecting both the current magnitude of each urge and learned associations between actions and urge satisfaction.

The memory systems in Psi architecture decompose into episodic, semantic, and procedural components serving distinct functional roles. Episodic memory stores specific experiences as distinct events preserving temporal and contextual information about what happened, where, and when. This system enables retrieval of particular past experiences and supports learning of one-shot contingencies from single events. Semantic memory, in contrast, stores abstracted knowledge extracted from multiple experiences, representing general facts, concepts, and regularities without preserving the specific episodic contexts in which they were learned. The extraction of semantic knowledge from episodic experiences occurs through consolidation processes that identify consistent patterns across multiple episodes while discarding idiosyncratic details.

Procedural memory stores skilled behaviors and habits as motor programs and stimulus-response associations that can be executed with minimal conscious oversight. Unlike declarative memories that can be consciously recalled and described, procedural memories manifest primarily through performance, enabling automatic execution of

complex action sequences. The formation of procedural memories occurs gradually through repeated practice, with performance progressively shifting from effortful controlled execution requiring working memory and attention to automatic fluent execution that proceeds without conscious monitoring.

The planning system implements forward simulation capabilities enabling the agent to mentally evaluate potential action sequences before executing them physically. This requires maintaining internal models, often termed "forward models," that predict how the environment and the agent's own state will change in response to actions. Given a current state and a candidate action, the forward model predicts the resulting next state, enabling the system to simulate trajectories by iteratively applying the model. Planning occurs through tree search in which multiple possible action sequences are simulated, their predicted outcomes are evaluated based on expected urge satisfaction, and the action sequence yielding the best predicted outcome is selected for execution.

Bach emphasizes that effective planning requires more than just forward models; it also requires appropriate search strategies that balance exploration of alternative possibilities against computational resource limitations. Deep exhaustive search of all possible action sequences becomes intractable as planning horizons extend beyond a few steps, necessitating heuristic pruning strategies that selectively explore promising regions of the action space. Additionally, planning must integrate uncertainty, recognizing that forward models provide imperfect predictions and that outcomes contain irreducible stochasticity requiring robust strategies that perform acceptably across multiple possible scenarios rather than optimizing for single predicted futures.

The metacognitive system implements self-monitoring and self-control capabilities that enable the agent to observe and modify its own cognitive processes. This requires developing an internal model of the agent's own cognitive states, capabilities, and limitations—what Bach terms the "self-model." The self-model maintains representations of what information the agent currently possesses, what cognitive resources are currently available, what tasks the agent is capable of performing, and what uncertainties exist in the agent's knowledge. This self-representation enables several critical metacognitive functions

including confidence estimation, error detection, uncertainty monitoring, and adaptive control of cognitive resource allocation.

Confidence estimation involves assessing the reliability of one's own perceptual judgments, memory retrievals, and decisions. The system must distinguish between situations where it can respond with high certainty and those where significant uncertainty exists, potentially motivating information-gathering before committing to action. Error detection enables the system to recognize when its cognitive processes have malfunctioned, such as when a decision was based on flawed reasoning or when an intended action was executed incorrectly. Metacognitive monitoring of uncertainty plays a critical role in determining when to trust one's own judgments versus seeking additional information or external guidance.

Adaptive control involves modulating cognitive parameters based on metacognitive assessment. When confidence is low, the system might increase attention to relevant information sources, extend deliberation time before making decisions, or reduce learning rates to avoid excessive weight on potentially unreliable information. When errors are detected, the system might engage corrective processes, adjust strategies, or increase exploration to discover superior approaches. This metacognitive control implements a higher-order optimization process operating on the parameters of lower-level cognitive systems.

The integration of these six major subsystems—hierarchical perception, hierarchical action control, urge-based motivation, multi-component memory, planning through simulation, and metacognitive self-monitoring—creates what Bach argues is a sufficient architectural framework for autonomous intelligent agency. Each component serves specific functional roles, but the emergent properties of the system arise primarily from their interaction. Perception provides the information that memory systems encode, which planning systems use to simulate futures, which motivational systems evaluate, which metacognitive systems monitor, and which action systems execute, closing the loop as executed actions produce new sensory consequences.

Bach's framework makes several strong empirical predictions that our biological implementation enables testing. First, the theory predicts that hierarchical organization in both perception and action is not merely an implementation detail but a functional necessity for achieving abstraction, generalization, and compositional structure in representations. Second, the theory predicts that urge-based motivation with multiple competing drives produces more robust and adaptive behavior than single-objective optimization or simple homeostatic regulation. Third, the theory predicts that planning through explicit forward model simulation should be distinguishable from model-free cached-value learning, producing characteristic behavioral signatures in tasks where model-based and model-free strategies make different predictions. Fourth, the theory predicts that metacognitive self-monitoring should correlate with and causally influence task performance, with better metacognitive accuracy enabling better adaptive control.

Previous tests of Psi theory have relied primarily on computational simulations implementing the architecture in artificial neural networks operating in simplified virtual environments. While these simulations have demonstrated that the architecture can produce intelligent behavior and have validated some theoretical predictions, they cannot address questions about biological plausibility, embodied constraints, or the sufficiency of proposed mechanisms when subjected to the full complexity of physical implementation. Our biological instantiation moves beyond simulation to create a physical Psi system operating in real time, subject to metabolic constraints, embodied in actual tissue, and demonstrating whether the abstract computational principles articulated by Bach can be realized in biological substrate and whether such implementation produces the predicted cognitive capabilities.

1.2 Organoid Technology and Bioengineering Foundations

The remarkable recent progress in organoid technology has opened unprecedented possibilities for creating complex tissue systems that recapitulate key aspects of organ development, architecture, and function outside the body. Organoids are three-dimensional tissue cultures derived from stem cells that self-organize into structures bearing architectural and functional similarity to actual organs through processes that partially recapitulate normal development. The field emerged from pioneering work demonstrating

that when provided with appropriate biochemical signals mimicking developmental morphogen gradients, pluripotent stem cells could spontaneously organize into complex tissue structures including neural rosettes, optic cups, intestinal crypts, and other organ-specific features.

Cerebral organoids, the most relevant for our purposes, were first successfully generated in 2013 when researchers demonstrated that human pluripotent stem cells subjected to neural induction followed by three-dimensional culture in Matrigel droplets and rotating bioreactors could form structures bearing remarkable similarity to developing human cortex. These early cerebral organoids exhibited multiple cortical layers, generated diverse neural subtypes including glutamatergic projection neurons and GABAergic interneurons, formed functional synaptic connections, and displayed spontaneous electrical activity. Subsequent refinements have dramatically improved the reproducibility, longevity, size, and complexity of cerebral organoids, with current protocols generating organoids up to 6 millimeters in diameter that can survive and mature for periods exceeding 18 months in culture.

The developmental processes underlying organoid self-organization provide important insights into both normal development and the engineering principles enabling their creation. When pluripotent stem cells are aggregated in three-dimensional culture and exposed to neural induction signals, they undergo a series of cell fate decisions that progressively restrict their developmental potential while specifying regional and cell-type identity. The initial neural induction phase, typically achieved through dual SMAD inhibition blocking both BMP and TGF-beta signaling pathways, drives cells toward a neural ectodermal fate while suppressing alternative mesodermal and endodermal programs. This creates a population of neural progenitor cells expressing transcription factors including SOX2, PAX6, and NESTIN that characterize the early neural tube.

Following neural induction, the progenitor cells undergo patterning along both anterior-posterior and dorsal-ventral axes in response to morphogen gradients. In typical cerebral organoid protocols without additional patterning factors, cells default toward an anterior dorsal identity resembling forebrain cortex due to the absence of posteriorizing signals like retinoic acid and caudalizing factors like FGF and WNT. The resulting tissue

exhibits cortical characteristics including the formation of a ventricular-zone-like proliferative region where neural progenitors undergo neurogenesis, generating post-mitotic neurons that migrate radially outward to form cortical layers. The sequential generation of deep-layer neurons followed by superficial-layer neurons partially recapitulates the inside-out layering pattern observed in normal cortical development, though the layering in organoids tends to be less organized than in vivo cortex.

The cells within cerebral organoids undergo progressive maturation over weeks to months, transitioning from immature neural progenitors to mature neurons with elaborate dendritic and axonal arbors, functional synapses, and appropriate electrophysiological properties. Synaptogenesis commences around three to four weeks of culture, with synaptic density increasing progressively over subsequent months. The synapses formed in organoids appear structurally normal by electron microscopy, containing presynaptic vesicles, synaptic clefts with appropriate spacing, and postsynaptic densities with characteristic protein compositions. Functional validation through electrophysiological recording demonstrates that these synapses support both excitatory transmission mediated by glutamate receptors and inhibitory transmission mediated by GABA receptors, with kinetics and pharmacology closely resembling brain synapses.

The spontaneous electrical activity that emerges in maturing organoids provides key evidence for functional network formation. Initially, neural activity consists primarily of isolated action potentials from individual neurons with minimal coordination. However, over developmental time, synchronized network activity emerges in which large populations of neurons fire together in bursts separated by periods of relative quiescence. These synchronized bursts occur with characteristic frequencies of approximately 0.1 to 0.3 Hz and reflect the formation of recurrent excitatory networks modulated by developing inhibitory circuitry. The progressive maturation of network activity patterns in organoids follows a trajectory qualitatively similar to that observed in developing brain tissue, transitioning from uncorrelated individual cell activity to increasingly synchronized population bursts to more complex patterns involving multiple oscillatory frequencies.

Despite these impressive achievements, standard cerebral organoid technology faces several critical limitations that required substantial innovation to overcome for our

purposes. First, organoids typically lack long-range axonal projections connecting distinct regions, remaining relatively isolated tissue masses without the extensive white matter tracts that enable communication between brain areas. This limitation arises because guidance cues and pathway substrates that normally direct growing axons toward appropriate targets are absent in the unstructured culture environment. Second, organoids lack functional vascular networks, relying instead on diffusion of oxygen and nutrients from surrounding medium, which limits maximum viable size to approximately 4-5 millimeters before hypoxic cores develop. Third, standard organoids show substantial batch-to-batch variability in size, cellular composition, and structural organization, reflecting the stochastic nature of self-organization processes. Fourth, organoids typically represent single brain regions rather than multi-regional assemblies with appropriate connectivity between functionally distinct areas.

Our implementation required addressing each of these limitations through deliberate engineering interventions that supplement self-organization with directed assembly. The challenge of establishing long-range connectivity between organoids required developing approaches to guide axon growth over distances of millimeters to centimeters, far exceeding typical organoid culture scenarios. We achieved this through multiple complementary strategies including the creation of physical guidance channels through microfluidic conduits, biochemical guidance through neurotrophin and morphogen gradients, molecular addressing through engineered cell-surface recognition molecules, and electrical stimulation protocols that enhance axon extension and pathfinding. These interventions enabled the formation of robust axonal projections connecting organoids positioned up to 4.5 centimeters apart, sufficient to span the full dimensions of our system.

The vascular limitation required engineering an entirely artificial circulatory system that provides dense capillary-like perfusion throughout the tissue volume. This represented one of the most significant technical challenges and innovations of this work, requiring development of methods to generate hierarchically organized vascular networks from endothelial progenitor cells, techniques for integrating these vascular structures with organoid tissue, and perfusion systems capable of maintaining appropriate flow, pressure, oxygen delivery, and waste removal throughout the complex three-dimensional tissue

architecture. The resulting vasculature achieved perfusion density approaching that of highly metabolically active brain tissue, enabling tissue masses far exceeding the limits of diffusion-based support.

The variability limitation was addressed through multiple quality control and selection strategies. We generated large numbers of organoids for each functional type and then selected those meeting stringent criteria for size, marker expression, cellular composition, and electrophysiological properties. Only organoids falling within tight specification windows were incorporated into the final assembled systems, substantially reducing the impact of developmental stochasticity. Additionally, we developed enhanced protocols incorporating improved media formulations, optimized growth factor concentrations and timing, controlled oxygen tension, and mechanical stimulation through perfusion and gentle agitation, all of which improve reproducibility of organoid generation.

The challenge of creating multi-regional assemblies with appropriate functional differentiation required developing protocols for generating not just generic cerebral organoids but also region-specific variants including visual cortex-like organoids with appropriate laminar organization and connectivity, motor cortex-like organoids with layer 5 pyramidal neurons capable of projecting to motor outputs, hippocampal-like organoids with dentate gyrus and CA3/CA1 structures supporting memory function, and striatal and cerebellar-like organoids supporting procedural learning. Each of these specialized organoid types required optimized differentiation protocols exposing cells to region-appropriate patterning signals including specific combinations of morphogens like WNT, FGF, retinoic acid, sonic hedgehog, and BMPs at precise developmental timepoints.

Beyond neural organoids, our system required generation of multiple other organoid types to provide the physiological support functions enabling autonomous operation. Cardiac organoids were generated through mesoderm induction followed by cardiac specification, yielding spontaneously contracting tissue masses containing cardiomyocytes, conducting system cells, and supporting cell types. The cardiac organoids required particular attention to metabolic maturation, as nascent cardiomyocytes initially rely primarily on glycolysis but must transition toward oxidative phosphorylation of fatty acids to achieve mature contractile function and energetic efficiency. This maturation was promoted through

media manipulations reducing glucose while providing fatty acids and L-carnitine, combined with electrical pacing and mechanical stretch stimuli that promote sarcomere organization and contractile protein expression.

Hepatic organoids providing detoxification and biosynthetic functions were generated through definitive endoderm induction followed by hepatic specification using combinations of FGF, HGF, OSM, and dexamethasone. The resulting organoids expressed key hepatocyte markers including albumin, alpha-1-antitrypsin, and cytochrome P450 enzymes, and demonstrated functional capacities including albumin secretion, urea synthesis from ammonia, glucose storage as glycogen, and drug metabolism. Achieving adequate hepatic function required particular attention to three-dimensional architecture enabling bile canaliculi formation and appropriate polarization of hepatocytes with basolateral surfaces facing sinusoid-like spaces and apical surfaces forming bile duct-like structures.

Renal organoids were generated using protocols that specify posterior intermediate mesoderm followed by metanephric mesenchyme, yielding structures containing nephron segments including glomeruli with podocytes and capillary tufts, proximal tubules, loops of Henle, and distal tubules. These kidney organoids demonstrated filtration capacity with size-selective barrier function, tubular reabsorption of glucose and proteins, and appropriate transport of ions and water. The integration of renal organoids into the perfusion system required careful attention to pressure management, as glomerular filtration depends on maintenance of appropriate hydrostatic pressure differentials between capillary lumens and Bowman's space.

The generation of functional vascular networks represented perhaps the single most critical innovation enabling our system's operation. Standard organoid cultures lack functional vasculature, with attempts to promote vascularization typically yielding only rudimentary vessel-like structures with limited functionality. We developed a comprehensive approach to vascular engineering that began with the differentiation of endothelial progenitor cells from iPSCs through mesodermal induction followed by specification toward vascular endothelium using VEGF and FGF signaling. These endothelial progenitors were then embedded in fibrin-collagen hydrogel matrices containing

pro-angiogenic factors that promoted their self-organization into capillary-like networks through vasculogenesis and angiogenesis processes.

The resulting primary vascular plexus required extensive maturation and remodeling to achieve functional perfusion capacity. This involved several stages including recruitment of mural cells through PDGF-BB signaling to stabilize nascent vessels, hierarchical remodeling to establish arteriole-capillary-venule organization with appropriate size distributions, and progressive increases in perfusion pressure to promote vessel maturation without causing damage. The fully mature vascular networks achieved remarkable properties including formation of tight endothelial junctions providing appropriate barrier function, hierarchical organization spanning three orders of magnitude in vessel diameter, and appropriate physiological responses including vasodilation in response to nitric oxide and vasoconstriction in response to endothelin.

The assembly of these diverse organoid types into an integrated functional system required developing sophisticated microfluidic platforms providing precise spatial positioning, separate vascular compartments for independent perfusion control, and optical access for imaging and optogenetic manipulation. The microfluidic devices were fabricated through soft lithography techniques using polydimethylsiloxane (PDMS) bonded to glass substrates with integrated electrode arrays. The devices incorporated multiple fluidic layers enabling independent control of perfusion to different organoid clusters, pneumatic valve systems for dynamic flow control, gradient generators for establishing morphogen and guidance cue distributions, and waste collection channels preventing toxic metabolite accumulation.

The optogenetic control system required comprehensive engineering of both the molecular actuators enabling light-driven neural manipulation and the physical light delivery infrastructure providing spatially and spectrally precise illumination. The molecular toolkit comprised seven distinct classes of light-sensitive proteins spanning the visible spectrum from 390 nm ultraviolet to 710 nm far-red wavelengths, enabling simultaneous independent control of multiple neural populations through spectral multiplexing. Each optogenetic protein was characterized in detail to determine its activation spectrum, kinetics,

photocurrent magnitude, and desensitization properties, then strategically deployed to specific cell populations based on functional requirements.

The physical light delivery system incorporated multiple complementary approaches tailored to different scales and requirements. For large-scale regional illumination, LED arrays with seven independent wavelength channels provided uniform illumination over millimeter to centimeter areas with intensities up to 20 milliwatts per square millimeter and temporal precision of 1 millisecond. For cellular-resolution stimulation, a digital micromirror device (DMD) coupled to the microscope enabled projection of arbitrary spatial patterns with single-cell precision updated at 60 Hz. For deep tissue stimulation where surface illumination produces inadequate light penetration, implanted optical waveguide arrays fabricated from optical fibers with 100 micrometer diameters delivered light directly to targeted regions up to several millimeters below the surface.

The electrophysiological recording infrastructure required equally sophisticated engineering to enable simultaneous measurement from thousands of recording sites distributed throughout the system. The foundation consisted of multielectrode arrays with 21,760 titanium nitride or platinum electrodes arranged in regular grids with 50 micrometer spacing, providing dense spatial sampling sufficient to detect signals from individual neurons positioned over any electrode. These planar arrays detected extracellular voltage fluctuations resulting from action potentials and synaptic currents in nearby neurons, with signals amplified, digitized at 20 kHz sampling rate with 16-bit resolution, and processed through spike sorting algorithms that cluster detected events into single units representing individual neurons.

Beyond these passive recording arrays, we implemented additional electrophysiological modalities for more detailed characterization of specific cells and circuits. Automated patch-clamp systems with motorized micromanipulators enabled targeted recording from visually identified neurons under microscope guidance, providing access to intracellular voltages and currents with precision unattainable through extracellular methods. These whole-cell recordings enabled measurement of intrinsic cellular properties including resting potential, input resistance, membrane time constant, spike threshold, and spike waveform parameters, as well as detailed characterization of

synaptic transmission including evoked postsynaptic currents, short-term plasticity, and long-term potentiation.

The biosensor systems enabling real-time monitoring of molecular and physiological state variables represented another critical component of the measurement and control infrastructure. We implemented 67 distinct biosensor types distributed throughout the system, each providing continuous readouts of specific analytes or physiological parameters. These included genetically encoded fluorescent sensors for intracellular signaling molecules, electrode-based sensors for electrolytes and metabolites, oxygen-sensitive phosphorescent probes for tissue oxygenation, and pH-sensitive fluorophores for monitoring acid-base status. The biosensor outputs were continuously sampled and logged, providing comprehensive records of system state evolution over time and enabling closed-loop control strategies where perfusion, stimulation, or other interventions were automatically adjusted based on sensor readouts.

The perfusion system supplying nutrients and removing wastes required careful engineering to maintain appropriate physiological conditions throughout the distributed tissue volume. The system incorporated multiple pumps providing both continuous baseline perfusion and pulsatile flow matching cardiac cycles, temperature control maintaining 37.0 ± 0.1 degrees Celsius throughout the device, oxygenation through membrane gas exchangers equilibrating perfusion medium with controlled gas mixtures, pH buffering through bicarbonate and HEPES systems, and automated fluid handling replacing medium and removing waste on programmed schedules. The perfusion medium composition was carefully optimized to provide all essential nutrients including amino acids, vitamins, lipids, glucose, hormones, and growth factors at concentrations supporting long-term culture while avoiding excess that might disrupt homeostatic regulation.

The computational infrastructure coordinating measurement, analysis, and control comprised multiple workstations and servers managing data streams from thousands of recording channels, processing signals through automated spike sorting and calcium imaging analysis pipelines, implementing closed-loop control algorithms modulating stimulation and perfusion based on measured system state, and storing terabytes of experimental data for offline analysis. The system operated continuously with automated

monitoring and alert systems notifying researchers of any deviations from normal parameters requiring intervention, enabling extended unattended operation during overnight and weekend periods critical for long-duration experiments.

1.3 Molecular and Cellular Implementation Strategies

The translation of abstract computational primitives specified in Bach's Psi architecture into concrete molecular and cellular mechanisms operating in living tissue required systematic identification of biological implementations for each cognitive operation. This process involved analyzing the functional requirements of each architectural component, identifying candidate molecular mechanisms with appropriate properties, engineering and optimizing these mechanisms for reliability and controllability, and integrating them into coherent cellular and tissue-level systems. The following sections detail the major classes of molecular implementation strategies we developed, establishing the technical foundation upon which the full cognitive architecture was constructed.

The optogenetic control systems enabling precise spatiotemporal manipulation of neural activity represent a cornerstone of our implementation approach, providing the actuators through which sensory information is delivered, motor commands are extracted, and internal circuit dynamics are experimentally manipulated. Optogenetics leverages light-sensitive proteins, primarily derived from microorganisms including algae, bacteria, and archaea, that undergo conformational changes upon photon absorption, altering their interactions with ions or signaling molecules. The most widely used optogenetic actuators are light-gated ion channels and pumps that move charged particles across cellular membranes, changing the voltage inside cells in response to illumination.

Channelrhodopsin-2, the pioneering optogenetic tool that launched the field, is a light-gated cation channel from the green alga Chlamydomonas reinhardtii that opens upon blue light absorption around 470 nanometers, allowing influx of sodium and calcium ions that depolarize neurons toward action potential threshold. When expressed in neurons through genetic engineering, ChR2 enables precise optical control of neural activity with millisecond temporal resolution, as light pulses as brief as 1 millisecond can reliably trigger action potentials. The kinetics of ChR2 are well-characterized, with an activation time

constant of approximately 1.3 milliseconds following light onset and a deactivation time constant of approximately 11 milliseconds after light offset, enabling reliable photostimulation at frequencies up to 40 Hz and detectable responses at frequencies approaching 100 Hz.

However, ChR2's properties, while revolutionary when it was introduced, have limitations that motivated development of improved variants and alternative opsins with enhanced or specialized characteristics. The relatively slow closing kinetics of ChR2 prevent reliable very-high-frequency stimulation above 40-50 Hz, limiting applications requiring precise temporal patterns at gamma frequencies. The blue light sensitivity, while convenient given the availability of blue LEDs, results in relatively poor tissue penetration due to scattering and absorption by tissue chromophores, limiting effective stimulation to approximately 100-500 micrometers depending on tissue properties. The relatively small photocurrent generated per ChR2 molecule requires high expression levels that can potentially impact cellular health through protein burden or aberrant calcium influx.

These limitations motivated our inclusion of multiple enhanced channelrhodopsin variants with improved properties. Chronos, derived from Stigeoclonium helveticum, exhibits ultra-fast kinetics with an off-time constant of only 3.6 milliseconds, enabling reliable stimulation above 200 Hz for applications requiring precise high-frequency temporal patterns. ChRmine, from Tiarina fusus, generates photocurrents approximately 3.2-fold larger than ChR2, enabling reliable neural activation with lower expression levels and reduced phototoxicity. CoChR, a computationally designed channelrhodopsin optimized through protein engineering, provides enhanced function at low expression levels, particularly valuable in cell types where high opsin expression produces toxicity or where promoter strength limitations prevent achieving high protein levels.

The spectral properties of these excitatory opsins span primarily the blue region of the spectrum, which motivated inclusion of red-shifted variants enabling spectral multiplexing where different neural populations can be controlled independently by different wavelength channels. ReaChR and ChrimsonR both respond optimally to orange-red light around 590-660 nanometers, providing approximately 100 nanometer spectral separation from blue opsins sufficient for largely independent control when

illumination spectra are appropriately filtered. The red-shifted opsins offer the additional advantage of improved tissue penetration, as longer wavelengths scatter less and experience reduced absorption by hemoglobin and other tissue chromophores, enabling effective stimulation at depths approaching 1-2 millimeters in neural tissue.

For applications requiring neural inhibition rather than excitation, we incorporated multiple inhibitory opsins that hyperpolarize neurons in response to light through outward pumping of ions. Halorhodopsin, derived from archaea including Natronomonas pharaonis, functions as a light-driven chloride pump that moves negatively charged chloride ions into the cell, making the interior more negative and thus hyperpolarizing the neuron away from action potential threshold. The eNpHR3.0 variant we employed shows improved membrane trafficking and photocurrent compared to earlier versions, producing hyperpolarizations of 20-35 millivolts at moderate light intensities of 5 milliwatts per square millimeter. Archaerhodopsin-3, functioning as an outward proton pump, achieves even stronger hyperpolarization of 40-60 millivolts by pumping positively charged protons out of the cell, though the proton currents can potentially disturb intracellular pH if sustained for extended periods.

The incorporation of both excitatory and inhibitory opsins enables bidirectional control where neural activity can be both increased and decreased on demand, providing much greater flexibility than unidirectional control alone. This bidirectional control proves essential for many applications including implementing lateral inhibition in sensory processing circuits where activated neurons must suppress neighbors, creating gain control mechanisms where background inhibition adjusts the overall excitability of circuits, and testing necessity of neural activity by silencing specific populations and observing resulting functional deficits.

Beyond these straightforward excitatory and inhibitory opsins operating through direct electrophysiological effects, we incorporated several classes of modulatory opsins that influence neural function through intracellular signaling pathways operating on slower timescales. The OptoXR family comprises light-activated G-protein coupled receptors engineered by fusing light-sensitive rhodopsin domains to intracellular signaling domains from metabotropic receptors. These chimeric proteins enable optogenetic control of second

messenger cascades including Gq pathways activating phospholipase C and generating IP3 and diacylglycerol, Gs pathways activating adenylyl cyclase and elevating cAMP, and Gi pathways inhibiting adenylyl cyclase and reducing cAMP. These pathways modulate neural excitability, synaptic transmission, gene expression, and many other cellular processes on timescales of seconds to minutes, providing complementary control to the millisecond-timescale control afforded by ionotropic opsins.

The BLINK system provides even more targeted optogenetic control of specific signaling molecules by implementing light-induced protein-protein interactions that reconstitute split enzymes or recruit signaling proteins to specific subcellular locations. For example, optogenetic control of CaMKII, a kinase critical for synaptic plasticity, can be achieved by splitting the kinase into two inactive fragments that reconstitute into active enzyme only when brought together through light-induced heterodimerization of attached photosensitive protein domains. This approach enables precise temporal control over plasticity-inducing signals, allowing experimenters to determine whether specific temporal patterns of CaMKII activation are sufficient to induce synaptic strengthening.

The comprehensive optogenetic toolkit we assembled ultimately comprised 127 distinct optogenetic constructs deployed strategically throughout the system, with each cell population expressing one or more opsins appropriate for its functional role. The sensory input neurons in Layer 1 expressed primarily excitatory opsins including ChR2 and ChRmine, enabling delivery of sensory information through patterned optical stimulation. The interneuron populations providing lateral inhibition expressed halorhodopsin, allowing optical suppression of inhibitory neurons to disinhibit principal cells when required. The motor output neurons expressed red-shifted opsins to avoid crosstalk with blue-light stimulation used elsewhere in the system. The modulatory neuron populations including dopaminergic reward prediction error neurons expressed OptoXR variants enabling control of their tonic activity levels and phasic response patterns.

The genetic delivery of these opsins to specific cell populations required sophisticated targeting strategies ensuring each opsin reached its intended cellular destination while avoiding off-target expression that would compromise specificity of optogenetic control. We employed three complementary approaches to achieve

cell-type-specific opsin expression. First, we used cell-type-specific promoters that drive gene expression selectively in neurons expressing particular transcription factors or proteins. For example, the CaMKII promoter drives expression primarily in excitatory glutamatergic neurons, while the GAD67 promoter targets GABAergic inhibitory neurons. These promoters, while imperfect, provide substantial enrichment for intended cell types.

Second, we used intersectional genetic strategies combining recombinase-mediated gene expression control with cell-type-specific promoters to achieve higher specificity than either approach alone. In this strategy, the opsin coding sequence is placed downstream of a transcriptional stop signal flanked by loxP recognition sites for Cre recombinase. The opsin can only be expressed in cells where Cre recombinase excises the stop signal, allowing transcription to proceed. By placing Cre recombinase expression under control of one cell-type-specific promoter and placing the conditional opsin construct under control of a different promoter, expression occurs only in cells meeting both criteria, dramatically improving specificity.

Third, we used viral transduction with engineered adeno-associated virus (AAV) vectors exhibiting tropism for specific cell types. Different AAV serotypes show preferential infection of different neural populations due to variations in capsid proteins that determine which cell-surface receptors mediate viral entry. By selecting appropriate AAV serotypes and engineering capsid variants with enhanced specificity, we achieved additional refinement of cell-type targeting. The combination of these three approaches—specific promoters, intersectional genetics, and targeted viral delivery—provided expression specificity typically exceeding 90% for most targeted cell populations.

The temporal control of opsin expression represented another important consideration, as we needed to prevent premature opsin expression during early developmental stages when circuit formation is still occurring, while ensuring adequate expression levels by the time functional experiments commence. This was achieved through inducible expression systems where opsin genes remain transcriptionally silent until activated by addition of small molecule inducers. The most commonly employed system uses tetracycline-controlled transcription where opsin genes are placed downstream of tet-responsive promoters that remain inactive until bound by tetracycline-controlled

transactivator proteins, which undergo conformational change upon doxycycline binding, enabling transcription. By controlling the timing of doxycycline addition to the culture medium, we could precisely determine when opsin expression began, typically initiating expression around day 60 of culture once major circuit formation was complete but allowing several weeks for opsin protein accumulation before functional experiments commenced.

The synthetic gene circuits implementing computational logic and temporal dynamics within cells represented a complementary approach to optogenetics for creating programmable cellular behaviors. While optogenetics provides external control through light-driven actuators, synthetic gene circuits create autonomous cellular programs that execute predetermined logic in response to molecular inputs. These circuits leverage the same basic components used in natural gene regulation—transcription factors that bind DNA and control gene expression, promoters that serve as binding sites and control regions for transcription factors, and protein-protein interactions that mediate signaling—but arrange them in novel configurations designed to implement specific computational functions.

The foundation of synthetic gene circuits lies in transcriptional regulation, where the expression of genes is controlled by the binding of transcription factors to promoter regions. In the simplest case, a transcription factor functions as a genetic switch, turning target gene expression on when the transcription factor is active and off when it is inactive. More complex behaviors emerge when multiple transcription factors regulate single genes through promoters containing multiple binding sites, when transcription factors regulate each other creating regulatory networks, and when feedback loops create dynamics including bistability, oscillations, and pulse generation.

The implementation of Boolean logic gates through transcriptional circuits requires careful engineering of promoter architectures and transcription factor binding affinities to achieve appropriate input-output relationships. An AND gate, requiring simultaneous presence of two inputs to produce output, can be implemented through a promoter containing binding sites for two different transcription factors arranged such that both must bind to recruit RNA polymerase and initiate transcription. The challenge lies in ensuring that neither transcription factor alone produces significant output while both together produce

robust expression, requiring optimization of the relative binding affinities, the spacing between binding sites, and the basal promoter strength.

Our AND gate implementations utilized dCas9-based transcriptional activators, where catalytically dead Cas9 proteins fused to VP64 activation domains are targeted to specific genomic locations through guide RNAs. The AND gate output gene promoter contained two distinct gRNA target sequences positioned appropriately for cooperative activation. Neither dCas9-VP64 variant alone produced strong activation, as the VP64 activation domains are relatively weak, requiring multimerization to effectively recruit transcriptional machinery. Only when both dCas9-VP64 variants bound simultaneously, bringing multiple VP64 domains into proximity, did strong transcriptional activation occur. The resulting transfer function showed output expression of only 8% of maximum with neither input present, 12% with single inputs, and 91% with both inputs, approximating ideal AND gate behavior with some leakiness in the off states that could be tolerated for our applications.

OR gate implementation proved more straightforward, as it requires that either of two inputs independently produce output. This was achieved through promoters containing two separate binding sites for different transcription factors, each sufficient to initiate transcription independently. The challenge with OR gates lies in ensuring that the dual-input state produces output levels comparable to single-input states rather than producing excessive overexpression that might cause toxicity or resource depletion. We addressed this through careful titration of promoter strengths and transcription factor expression levels, achieving OR gates with output levels of 6% with neither input, 74% with either single input, and 88% with both inputs simultaneously, showing the expected OR logic with appropriate saturation preventing excessive dual-input output.

NOT gates, inverting inputs to produce high output when input is absent and low output when input is present, were implemented using transcriptional repressors rather than activators. The dCas9-KRAB system, where Cas9 is fused to the Krüppel-associated box repression domain, provided strong transcriptional repression when targeted to promoter regions. The NOT gate architecture placed a dCas9-KRAB targeting site in the promoter region of the output gene such that when input signal activated dCas9-KRAB expression, the repressor bound and blocked output gene transcription. The resulting transfer function

showed output of 89% without input and only 11% with input present, providing effective inversion with some incomplete repression in the on state due to leaky transcription escaping repressor control.

These basic logic gates were composed into more complex combinatorial circuits implementing functions including NAND, NOR, XOR, multiplexers, and comparators through layered architectures where the output of one gate serves as input to downstream gates. For example, an XOR gate producing output when exactly one input is active but not both was constructed by combining AND, OR, and NOT gates in the canonical configuration: XOR(A,B) = OR(AND(A,NOT(B)), AND(NOT(A),B)). The layered architecture introduced temporal delays as each layer required time for transcription, translation, protein maturation, and accumulation to occur before influencing downstream layers. These delays, typically 2.4 to 6.8 hours per layer depending on promoter strength and protein properties, constrained the complexity of circuits that could operate on behaviorally relevant timescales but proved acceptable for applications including developmental patterning, metabolic state tracking, and consolidation scheduling.

Beyond Boolean logic, we implemented analog computational circuits performing continuous signal processing operations including integration, differentiation, and filtering. Integrator circuits, accumulating input signals over time, proved particularly valuable for implementing persistent working memory, evidence accumulation during decision making, and urge escalation in the motivational system. The molecular implementation of integration leverages positive feedback where a protein activates its own transcription, creating a self-sustaining loop that maintains elevated expression even after input stimulus ceases. The mathematical description of such an integrator circuit follows: $dX/dt = k_{in} \times Input - k_{deg} \times X$, where X represents the concentration of the integrator protein, k_{in} represents the rate of production driven by input signal, and k_{deg} represents the constitutive degradation rate determining how leaky the integrator is.

The critical design parameters for integrator circuits are the degradation rate, which determines integration time constant, and the maximum accumulation level, which determines saturation. For working memory applications requiring information maintenance over seconds to minutes, we engineered integrator proteins with short half-lives around

5-10 minutes through fusion to degron sequences targeting proteins for rapid proteasomal degradation. For urge accumulation requiring integration over hours, we used more stable proteins with half-lives of 3-6 hours. The positive feedback strength determined how effectively the integrator maintained activity after input ceased, with stronger feedback providing more persistent memory but potentially creating bistable switches rather than graded integrators.

Differentiator circuits, responding to the rate of change of input rather than absolute level, enabled detection of stimulus transients, prediction error computation, and novelty detection. The molecular implementation uses incoherent feedforward loop topology where input simultaneously activates both the output protein through a fast direct path and an inhibitor protein through a slower indirect path. The kinetic mismatch creates a response profile where output initially increases rapidly in response to input onset then gradually declines as the delayed inhibitor accumulates. The output thus reflects the temporal derivative of input, reporting how rapidly input is changing rather than its absolute level. The temporal characteristics of differentiation are determined by the delay in inhibitor synthesis and accumulation, typically 15-45 minutes in our implementations depending on promoter strength and protein maturation rates.

Low-pass filter circuits, smoothing rapidly fluctuating inputs to extract slowly varying trends, were implemented through multi-stage protein maturation cascades where input drives expression of a protein that must undergo several slow maturation steps before becoming functional. Each maturation step acts as a temporal integration stage, and multiple stages in series create increasingly strong temporal averaging. We achieved this by engineering fluorescent proteins with extended maturation kinetics through mutations that slow chromophore formation, creating proteins requiring 2-4 hours to mature after translation. Cascades of such proteins created effective low-pass filters with time constants ranging from 30 minutes to 8 hours depending on the number of stages, enabling extraction of slowly varying signals from noisy fluctuating measurements.

These synthetic gene circuits were deployed throughout the system to implement cognitive functions requiring cellular computation. The urge generation circuits in motivational organoids used integrators to accumulate homeostatic deficits, converting

continuous metabolic drains into escalating drive signals. The working memory circuits in prefrontal-like organoids used integrators with appropriate time constants to maintain task-relevant information across temporal gaps. The prediction error circuits in learning systems used differentiators to compute mismatches between expected and actual outcomes. The arousal regulation circuits used low-pass filters to track slowly varying metabolic state and overall activity levels determining wake-sleep transitions.

The engineering of synaptic transmission mechanisms enabling communication between neurons represented another critical molecular implementation challenge. Natural synapses employ sophisticated molecular machinery to achieve reliable, fast, and modifiable transmission of signals from presynaptic to postsynaptic neurons. The presynaptic terminal contains synaptic vesicles filled with neurotransmitter molecules that fuse with the plasma membrane upon calcium influx following action potential arrival, releasing neurotransmitter into the synaptic cleft. The neurotransmitter diffuses across the narrow cleft and binds to postsynaptic receptors that open ion channels or activate signaling pathways, transducing the presynaptic electrical signal into postsynaptic responses.

While organoid neurons naturally form functional synapses that recapitulate many aspects of native brain synapses, we implemented several synthetic enhancements to improve controllability and specificity of connectivity. The natural process of synapse formation proceeds through stochastic contact between axons and dendrites followed by stabilization of contacts where appropriate pre- and postsynaptic proteins accumulate. This stochastic process creates connectivity with some anatomical logic—neurons in proximity are more likely to connect, and neurons with compatible molecular profiles preferentially synapse—but without the precise specificity required to implement designed connectivity matrices specifying exactly which neurons should connect to which.

We achieved enhanced connectivity control through engineered synaptic cell adhesion molecules implementing molecular addressing where synapses form selectively between cells expressing complementary recognition molecules. The approach leverages the neuroligin-neurexin adhesion system, which naturally spans synaptic clefts with neuroligins on postsynaptic membranes binding to neurexins on presynaptic terminals. We generated 12 orthogonal neuroligin-neurexin pairs through directed evolution and rational

design, creating variants that bind their cognate partners with high affinity but show minimal cross-reactivity with other pairs. By expressing specific neuroligin variants in designated postsynaptic populations and complementary neurexin variants in presynaptic populations, we could specify connectivity rules such as "neurons of type A should preferentially synapse onto neurons of type B" through molecular complementarity.

The implementation required careful consideration of expression levels, binding affinities, and competition dynamics to achieve intended connectivity patterns. If expression levels of adhesion molecules varied too widely between cells, some cells would form excessive connections while others remained poorly connected. If binding affinities were too strong, promiscuous connections might form through low-affinity interactions between mismatched pairs. If expression of multiple adhesion pairs in single cells occurred, competition between pairs could produce unintended connectivity preferences. We addressed these challenges through extensive characterization and optimization, measuring expression levels via flow cytometry and immunostaining, quantifying binding affinities through surface plasmon resonance, and testing connectivity outcomes through paired recording experiments measuring connection probability between cell pairs expressing different adhesion molecule combinations.

The resulting molecular addressing system achieved approximately 76-82% connectivity specificity in optimized implementations, meaning that neurons expressing complementary adhesion molecules formed synapses with their intended partners about 78% of the time, while forming off-target synapses with non-complementary partners about 22% of the time. This represents substantial improvement over the baseline connection probability of approximately 4-8% that would occur through purely random contact in a system with 23 distinct cell types, but falls short of perfect specificity. The imperfect specificity arises from several factors including residual cross-reactivity between adhesion molecule variants, promiscuous synapse formation driven by other adhesion systems not under our control, and activity-dependent synapse formation where coincident firing can induce synapse formation even between molecularly mismatched partners.

Interestingly, the imperfect specificity may actually prove beneficial rather than limiting, as it creates a form of degeneracy where multiple pathways can mediate similar

functions, providing robustness against lesions and enabling alternative routes for information flow. The off-target connections also potentially enable discovery of novel connectivity patterns through learning, as activity-dependent plasticity mechanisms can strengthen initially weak off-target synapses if they prove functionally useful. The balance between specificity sufficient to create intended connectivity motifs and degeneracy sufficient to provide robustness represents an important design consideration for engineered neural systems.

Beyond connectivity specificity, we implemented several enhancements to synaptic transmission dynamics enabling better control and measurement of synaptic function. Engineered neurotransmitter receptors with modified properties provided tools for pharmacological manipulation of specific circuits. The PSAM-PSEM system comprises a family of ligand-gated ion channels derived from nicotinic acetylcholine receptors that no longer respond to acetylcholine but instead respond to synthetic small molecule agonists called PSEMs. We incorporated seven orthogonal PSAM-PSEM pairs, each responding to a distinct synthetic ligand, enabling independent pharmacological control of seven different neural populations. This provides complementary control to optogenetics, as pharmacological manipulation operates on slower timescales of seconds to minutes rather than milliseconds but doesn't require light delivery infrastructure and enables whole-system manipulation rather than being limited to light-accessible regions.

Modified NMDA receptors with enhanced properties for plasticity induction represented another important synaptic engineering target. Native NMDA receptors serve critical roles in synaptic plasticity by providing coincidence detection—they admit calcium only when presynaptic glutamate release occurs simultaneously with postsynaptic depolarization, implementing the Hebbian principle "neurons that fire together wire together." The calcium influx through NMDA receptors triggers intracellular signaling cascades including CaMKII activation that induce long-term potentiation, the primary cellular mechanism of learning and memory. However, native NMDA receptors have properties that may limit plasticity in some contexts, including voltage-dependent magnesium block that prevents calcium influx at hyperpolarized potentials and relatively short opening times limiting calcium influx.

We engineered enhanced NMDA receptor variants with reduced magnesium sensitivity through mutations in the ion channel pore region, enabling calcium influx at more negative membrane potentials and thus lowering the threshold for plasticity induction. Additionally, we incorporated NMDA receptor subunit compositions enriched for GluN2B-containing receptors rather than GluN2A, as GluN2B-containing receptors show longer opening times and greater calcium permeability, both properties that enhance plasticity. These enhanced NMDA receptors were expressed selectively in circuits requiring strong learning capabilities including hippocampal episodic memory systems and cortical layer 4 neurons implementing perceptual learning, while other circuits retained native receptor properties to prevent excessive plasticity that might destabilize learned representations.

Synaptic scaffolding proteins determining the composition and organization of synaptic protein complexes provided additional engineering targets for controlling synaptic properties. PSD-95, the primary scaffolding protein at excitatory synapses, determines synaptic strength by controlling the number of AMPA receptors anchored in the postsynaptic membrane. We generated PSD-95 variants with modified AMPA receptor binding affinities, creating some synapses with enhanced AMPA receptor content producing large postsynaptic currents and other synapses with reduced content producing smaller responses. The cell-type-specific expression of these PSD-95 variants enabled setting synaptic strength distributions appropriate for each connection type, with some connections intentionally strong to provide robust signal transmission and others weak to require temporal integration of multiple inputs.

The synaptic vesicle release machinery in presynaptic terminals represented yet another engineering target, as this machinery determines the probability that action potential arrival triggers neurotransmitter release and thus fundamentally controls synaptic reliability. We incorporated engineered synaptic vesicle proteins with modified calcium sensitivity of fusion, creating some synapses with high release probability that reliably transmit every presynaptic action potential and others with low release probability that transmit stochastically. This diversity of release probabilities across different connection types enables implementation of different computational functions, with reliable

high-probability synapses appropriate for critical information pathways requiring faithful transmission and unreliable low-probability synapses appropriate for connections mediating modulatory or threshold-based computations.

The comprehensive synaptic engineering toolkit ultimately comprised modifications at every stage of synaptic transmission from presynaptic vesicle loading through neurotransmitter release, diffusion, receptor binding, postsynaptic signaling, and plasticity induction. The strategic deployment of these diverse molecular tools to specific cell types and connection pathways enabled creation of synaptic connectivity with properties tailored to functional requirements, implementing the connectivity matrices and transmission dynamics specified by the cognitive architecture design.

The biosensor systems enabling real-time monitoring of cellular and tissue states represented the final major category of molecular implementation, providing the measurement capabilities essential for validating system function, implementing closed-loop control, and investigating mechanisms underlying observed behaviors. Natural biological systems employ diverse sensing mechanisms for monitoring internal state variables including metabolite levels, ion concentrations, voltage, mechanical forces, and signaling molecule activities. We leveraged and extended these natural sensing modalities to create comprehensive monitoring systems tracking 67 distinct state variables continuously throughout system operation.

Genetically encoded fluorescent biosensors comprise a large and diverse family of molecular tools for monitoring intracellular signaling and state variables through fluorescence changes reporting molecular events. These sensors typically consist of fluorescent protein modules coupled to molecular recognition domains that undergo conformational changes upon binding target ligands. The conformational change alters the fluorescent properties through several possible mechanisms including fluorescence resonance energy transfer (FRET) between donor and acceptor fluorophores, changes in chromophore protonation affecting absorption and emission spectra, or changes in quantum yield affecting fluorescence intensity.

The calcium indicators we deployed extensively throughout the system exemplify the design principles and capabilities of genetically encoded biosensors. The GCaMP family of calcium indicators consists of circularly permuted GFP inserted into the calmodulin protein, with the M13 peptide from myosin light chain kinase attached as a calmodulin-binding partner. In the absence of calcium, calmodulin remains in an open conformation and the GFP chromophore exhibits low fluorescence. When calcium binds to calmodulin, it undergoes conformational change bringing the M13 peptide into contact, which induces a conformational change in the GFP domain increasing fluorescence intensity. The latest GCaMP variants achieve very large dynamic ranges with fluorescence increasing 50-100 fold between zero calcium and saturating calcium levels, enabling detection of single action potentials in individual neurons.

We deployed multiple GCaMP variants optimized for different applications. GCaMP6f, the "fast" variant, exhibits rapid kinetics with an off time constant of approximately 142 milliseconds, enabling temporal resolution sufficient to distinguish individual action potentials in spike trains up to about 10 Hz. This variant proved essential for applications requiring spike-timing precision including measuring spike-timing-dependent plasticity and characterizing temporal coding properties. GCaMP6s, the "slow" variant, exhibits an extended off time constant around 1.5 seconds, providing lower temporal resolution but enhanced signal-to-noise ratio as the slower decay creates larger integrated calcium transients. This variant was preferable for population imaging applications where the goal is identifying which neurons are active rather than precisely resolving spike times.

The red calcium indicator jRGECO1a provided critical capability for simultaneous two-color imaging, as its red fluorescence with excitation around 560 nanometers and emission around 595 nanometers exhibits minimal spectral overlap with green GCaMP variants. By expressing GCaMP in one neural population and jRGECO in another, we could simultaneously monitor two distinct populations with independent readouts, enabling measurement of relative timing and interaction between populations. Additionally, the combination of green and red indicators enabled ratiometric imaging where the ratio of red to green fluorescence provides measurement less susceptible to artifacts from motion, focal plane drift, or uneven illumination that would affect both channels similarly.

Voltage indicators reporting membrane potential changes complemented calcium indicators by providing measurement of subthreshold dynamics not visible in calcium signals. While calcium indicators report primarily spiking activity due to the large calcium influx associated with action potentials, voltage indicators respond to all voltage changes including synaptic potentials, enabling measurement of synaptic integration processes determining whether neurons reach spike threshold. The ASAP family of voltage indicators achieves fast kinetics with sub-millisecond response times by leveraging voltage-dependent conformational changes in voltage-sensing domains from voltage-gated ion channels. These conformational changes alter the environment of attached fluorescent proteins, changing fluorescence intensity within approximately 1 millisecond of voltage changes.

The ASAP3 variant we employed shows a fluorescence change of approximately 35% per 100 millivolts, enabling detection of action potentials with signal-to-noise ratios around 10-20 in single trials and detection of synaptic potentials through averaging multiple trials. The fast kinetics enable faithful reporting of action potential waveforms with temporal fidelity limited primarily by camera frame rates rather than indicator kinetics. We deployed voltage indicators selectively rather than system-wide due to their greater photon demands compared to calcium indicators, focusing voltage imaging on circuits where subthreshold dynamics particularly matter such as decision circuits implementing evidence accumulation and integration processes.

The neurotransmitter sensors enabling optical detection of neurotransmitter release represented another critical class of biosensors, as they enable measurement of synaptic transmission and neuromodulation with spatial resolution unattainable through electrode-based methods. The iGluSnFR glutamate sensor exemplifies this class, consisting of the bacterial glutamate-binding protein GltI inserted into circularly permuted GFP such that glutamate binding induces fluorescence changes. The sensor achieves a dynamic range of 4-5 fold fluorescence increase at 1 micromolar glutamate, sufficient to detect single synaptic glutamate release events when sensors are appropriately positioned near synapses. The kinetics show a time constant around 150 milliseconds for glutamate clearance, determined primarily by the natural glutamate uptake mechanisms in the tissue rather than intrinsic sensor properties.

Similar sensors for other neurotransmitters and neuromodulators enabled comprehensive monitoring of chemical signaling throughout the system. The iGABASnFR sensor reports GABA using analogous design principles, the dLight sensor reports dopamine release from modulatory neurons, and sensors for norepinephrine, serotonin, and acetylcholine were deployed in regions where those neuromodulators play important functional roles. The simultaneous measurement of multiple neurotransmitter systems through spectrally distinct sensors or sequential imaging revealed the complex interplay between different neuromodulatory systems in controlling circuit dynamics and plasticity.

Metabolite sensors tracking glucose, lactate, ATP, and other energetic substrates provided crucial information about the metabolic state determining cellular energetics and influencing urge generation in the motivational system. The glucose sensors employ bacterial periplasmic glucose-binding proteins inserted into FRET-capable fluorescent protein pairs, such that glucose binding brings donor and acceptor fluorophores into proximity, increasing FRET efficiency measurable as a change in the ratio of acceptor to donor emission. These sensors report glucose concentrations across the physiologically relevant range from approximately 0.1 to 10 millimolar with appropriate dynamic range and sensitivity.

The ATP sensors from the ATeam family provide ratiometric readout of ATP/ADP ratios through FRET between cyan and yellow fluorescent proteins linked by the epsilon subunit of F1-ATPase, which undergoes conformational changes depending on whether ATP or ADP occupies the binding site. The ratio of yellow to cyan emission thus reports the energetic state of the cell, with high ratios indicating energy sufficiency and low ratios indicating energy depletion. These sensors proved particularly valuable for monitoring energy crises during glucose deprivation experiments and for validating that energy urge generation appropriately tracks cellular energetic state.

Signaling pathway sensors reporting the activity of kinases, phosphatases, and other signaling enzymes provided detailed information about the intracellular signaling cascades mediating plasticity, gene expression, and cellular state transitions. The AKAR sensor for protein kinase A (PKA) activity consists of a PKA substrate sequence placed between CFP and YFP fluorophores along with a phospho-amino acid binding domain. When PKA

phosphorylates the substrate, the phospho-binding domain interacts with the phosphorylated residue, creating a conformational change that alters FRET efficiency. Similar sensors for ERK, CaMKII, and other signaling molecules enabled tracking of the molecular events mediating learning, memory consolidation, and adaptive responses to perturbations.

The oxygen and pH sensors provided crucial information about tissue physiology and perfusion adequacy. Oxygen-sensitive phosphorescent probes based on metalloporphyrins show phosphorescence lifetimes that decrease with increasing oxygen partial pressure due to collisional quenching, enabling quantitative oxygen mapping through phosphorescence lifetime imaging. These probes distributed throughout the tissue revealed spatial variations in oxygenation, identifying regions at risk for hypoxia and validating adequacy of vascular perfusion. The pH-sensitive fluorophores based on fluorescein derivatives show pH-dependent protonation affecting absorption and emission properties, enabling ratiometric pH measurement across the physiological range.

The comprehensive deployment of these 67 distinct biosensor types throughout the system created unprecedented visibility into system state across molecular, cellular, and tissue scales. Every organoid contained at least 8-12 different biosensor types providing local readouts, while some biosensor types were expressed nearly ubiquitously enabling system-wide mapping. The continuous logging of biosensor data created comprehensive records of system state evolution, enabling correlation of molecular events with cellular activity patterns and behavioral outputs, validation of mechanistic hypotheses about how cognitive functions emerge from molecular implementations, and early detection of pathological states requiring intervention.

1.4 Integration Challenges and Architectural Assembly

The construction of a functional cognitive system from the diverse molecular and cellular components described above required solving profound integration challenges spanning multiple organizational scales from molecular through cellular to tissue and systems levels. The naive approach of simply combining organoids expressing appropriate molecular toolkits would fail for numerous reasons including lack of appropriate spatial

organization, absence of functional connectivity between components, inadequate metabolic support for large tissue masses, and absence of coordination mechanisms synchronizing activity across distributed subsystems. Our successful implementation required developing comprehensive strategies for spatial organization, connectivity establishment, metabolic support, and functional integration, as detailed in the following sections.

The spatial organization challenge required determining how to position 289 organoids of 23 distinct types across a 60×90 centimeter area such that connectivity between components could be established with appropriate path lengths, such that functionally related components were positioned proximally to enable dense local connectivity while functionally distinct components maintained appropriate separation, and such that perfusion infrastructure could reach all tissue while minimizing dead space and vascular path lengths. This spatial planning problem has combinatorial complexity, as there are vastly more possible arrangements than could be explored exhaustively, requiring development of optimization approaches balancing multiple competing constraints.

We formulated the spatial organization problem as a constraint satisfaction optimization where the objective function combined terms penalizing undesirable properties including excessive connection lengths between components specified to communicate, inadequate spacing between components risking physical interference, suboptimal perfusion network topology, and deviation from desired overall system geometry. The optimization employed simulated annealing where candidate organoid placements were iteratively refined through stochastic perturbations, with changes accepted probabilistically based on whether they improved the objective function, with acceptance probability decreasing over time to enable gradual convergence toward good solutions.

The optimization process generated spatial layouts organizing the system into seven major functional domains corresponding to the primary components of the Psi architecture. The sensory processing domain occupying approximately 12×90 centimeters contained 47 neural organoids arranged in six hierarchical layers implementing progressive abstraction of sensory representations. Within this domain, organoids were positioned to enable short

connections between adjacent processing layers while maintaining separation between non-adjacent layers to prevent short-circuiting of the hierarchy. The motor control domain similarly occupied 12×90 centimeters with 43 organoids arranged in five hierarchical stages from abstract motor planning through concrete motor execution.

The memory system domain occupied 10×60 centimeters containing 38 organoids partitioned into episodic, semantic, and procedural memory subsystems, with appropriate connectivity both within each subsystem and between subsystems enabling consolidation pathways transferring information from episodic to semantic memory and from explicit to implicit procedural learning. The motivational control domain of 8×60 centimeters contained 24 organoids implementing urge generation, urge integration, and action valuation based on urge satisfaction predictions. The planning and simulation domain of 10×60 centimeters housed 31 organoids implementing forward models, tree search processes, and counterfactual simulation capabilities. The metacognitive domain of 8×60 centimeters contained 19 organoids implementing self-monitoring, confidence estimation, and adaptive control functions.

The final domain consisted of the distributed homeostatic support systems including 67 cardiac organoids positioned throughout the system to provide localized perfusion enhancement, 23 hepatic organoids clustered centrally to maximize detoxification efficiency through high local perfusion, and 18 renal organoids arranged in paired clusters flanking the central regions to enable appropriate pressure management for filtration function. This organization reflected both functional requirements and metabolic demands, with the highest metabolic-rate neural organoids positioned to minimize distances from cardiac and hepatic support structures.

The connectivity establishment challenge required creating approximately 2.3×10^8 synaptic connections organized into 127 major axon tracts connecting organoids positioned up to 4.5 centimeters apart, far exceeding the typical millimeter-scale distances over which axons extend in standard organoid cultures. Natural axon pathfinding during development relies on carefully orchestrated guidance cue gradients, adhesive and repulsive molecular signals, and activity-dependent refinement processes that operate over distances of millimeters in the developing brain. Recapitulating these processes over centimeter scales

in our engineered system required substantial innovation in guidance strategies and physical scaffolding.

We developed a multi-phase approach to long-range connectivity establishment beginning with physical guidance channels providing structural scaffolds directing initial axon extension, followed by molecular guidance through neurotrophin gradients attracting growing axons toward targets, then molecular addressing through engineered cell-surface recognition molecules specifying which neurons should synapse with which targets, and finally activity-dependent refinement strengthening appropriate connections while eliminating inappropriate ones. Each phase built upon previous phases, with the physical channels establishing rough topography, molecular guidance achieving target region specificity, molecular addressing achieving cellular specificity, and activity-dependent refinement achieving functional specificity.

The physical guidance channels consisted of microfluidic conduits fabricated into the device substrate, creating hollow tubes with inner diameters of 200-500 micrometers connecting source organoid positions to target organoid positions. The inner surfaces of these conduits were coated with extracellular matrix proteins including laminin and fibronectin providing adhesive substrates for axon growth, and the conduits were filled with hydrogel containing growth factors and guidance molecules creating permissive environments for axon extension. Growing axons from source organoids encountered conduit openings and preferentially entered these structures due to the favorable adhesive and trophic environment compared to surrounding spaces.

The microfluidic conduits provided not only passive guidance but also active delivery of guidance cues through controlled perfusion. By flowing solutions containing neurotrophin gradients through the conduits with source regions exposed to lower concentrations and target regions exposed to higher concentrations, we created continuous molecular gradients spanning centimeter-scale distances. Growing axons expressing appropriate neurotrophin receptors exhibited chemotaxis toward increasing neurotrophin concentrations, providing directional guidance toward targets. We utilized multiple neurotrophins including nerve growth factor (NGF), brain-derived neurotrophic factor (BDNF), neurotrophin-3 (NT-3), and glial cell line-derived neurotrophic factor (GDNF), with

different neurotrophins attracting different axon populations enabling parallel establishment of multiple distinct pathways.

The molecular addressing through engineered cell-surface adhesion molecules described previously came into play once axons reached target organoid vicinity, determining which specific neurons within target organoids received synaptic inputs from the arriving axons. The combination of physical guidance delivering axons to appropriate target regions and molecular addressing specifying cellular targets within those regions achieved connectivity specificity unattainable through either mechanism alone. This hierarchical targeting strategy mirrors natural developmental mechanisms where early pathfinding brings axons to appropriate brain regions through relatively coarse guidance, followed by local refinement achieving cellular precision through molecular matching.

The activity-dependent refinement phase required orchestrating appropriate patterns of correlated activity between source and target populations to strengthen functionally appropriate connections while eliminating inappropriate ones. This refinement exploits spike-timing-dependent plasticity mechanisms where synapses are strengthened when presynaptic activation slightly precedes postsynaptic activation, implementing Hebbian learning. We induced appropriate correlated activity through coordinated optogenetic stimulation delivering temporally structured patterns to multiple organoids simultaneously, creating artificial experiences that drove activity-dependent consolidation of connectivity.

The refinement protocols delivered carefully designed stimulation sequences over periods of 2-4 weeks, beginning around day 75 when major axonal projections had reached targets but synaptic connectivity remained immature. The sequences activated source populations with specific temporal patterns, followed after precisely controlled delays by activation of intended target populations, creating the temporal contingencies required for Hebbian strengthening. Populations that should be functionally connected received correlated stimulation with appropriate timing, while populations that should remain independent received uncorrelated stimulation, driving selective strengthening of intended connections and elimination of spurious connections formed during initial promiscuous synaptogenesis.

The resulting connectivity achieved several critical properties validating the success of our multi-phase establishment strategy. Anatomical tracing using anterograde and retrograde viral tracers confirmed that axonal projections reached intended target organoids with approximately 82% accuracy, meaning that roughly 82% of labeled axons from source organoids projected to designated target organoids rather than alternative destinations. The cellular specificity within target organoids, assessed through paired electrophysiological recordings and connection probability measurements, showed that approximately 76% of connections formed between neurons expressing complementary molecular addressing cues compared to baseline connection probabilities around 4-8% for non-matching pairs.

The functional specificity of established connections, assessed through correlation analyses of neural activity during behavior, demonstrated that connected neuron pairs showed significantly higher activity correlations (mean correlation coefficient 0.34 ± 0.12) compared to unconnected pairs (0.08 ± 0.07), confirming that anatomical connectivity translated into functional communication. The distribution of synaptic strengths measured through evoked postsynaptic potential amplitudes showed appropriate heterogeneity with some very strong connections producing postsynaptic responses exceeding 2 millivolts from single presynaptic spikes, many moderate connections producing 0.3-0.8 millivolt responses, and numerous weak connections producing <0.2 millivolt responses, creating the mixture of reliable and modulatory connections required for robust yet flexible circuit function.

The metabolic support challenge required ensuring that every cell within the large distributed tissue mass received adequate oxygen, glucose, and other nutrients while enabling removal of carbon dioxide, lactate, and other metabolic waste products that would otherwise accumulate to toxic levels. The fundamental constraint is that molecular transport over distances exceeding approximately 150-200 micrometers must occur through convection via flowing fluids rather than diffusion, as diffusion becomes prohibitively slow over longer distances. This constraint necessitated engineering vascular networks achieving perfusion density sufficient that no tissue region remained more than 150 micrometers from functional vasculature.

The vascular engineering strategy employed several parallel approaches ultimately achieving comprehensive perfusion coverage. The foundation consisted of an engineered primary vascular network created ex vivo through self-organization of endothelial progenitor cells within fibrin-collagen hydrogels. When endothelial cells are embedded in three-dimensional extracellular matrix at appropriate density with suitable growth factor support, they spontaneously form vessel-like tubular structures through vasculogenesis, a process recapitulating aspects of embryonic blood vessel formation. The cells extend processes, make contact with neighbors, and coalesce into elongated cord structures that then undergo lumenization creating hollow tubes suitable for fluid flow.

The primary vascular network generated through this self-organization process created a dense capillary-like meshwork with vessel diameters primarily in the 5-15 micrometer range and spacing between vessels around 50-150 micrometers, approaching the densities required for adequate perfusion. However, this primary network alone proved insufficient for several reasons including lack of hierarchical organization with distinct arteriolar inputs and venular outputs needed for directed flow, inadequate vessel stability with vessels regressing without appropriate maturation signals, and limited connection to larger engineered feed vessels supplying fresh medium and removing waste.

We addressed these limitations through several maturation interventions applied over weeks following initial network formation. The recruitment of mural cells including pericytes and smooth muscle cells stabilized nascent vessels through physical ensheathment and paracrine signaling. This was achieved through controlled gradients of platelet-derived growth factor BB (PDGF-BB), which attracts pericyte precursors to vessels where they differentiate and ensheath capillaries. The pericytes secrete factors including angiopoietin-1 that signal through Tie2 receptors on endothelial cells, promoting vessel maturation and reducing permeability. For larger vessels requiring smooth muscle coverage, we employed transforming growth factor-beta (TGF-β) gradients attracting smooth muscle progenitors that differentiate and form organized layers around vessel walls.

The hierarchical organization into arterial-capillary-venous structure emerged through a combination of molecular patterning and hemodynamic remodeling. We applied gradients of arterial-specifying factors including Notch ligands and ephrin-B2 at network

input regions and venous-specifying factors including COUP-TFII and EphB4 at output regions, creating initial arterio-venous polarization. The subsequent initiation of flow through the network drove extensive remodeling where vessels carrying high flow rates enlarged through vessel wall remodeling while low-flow vessels regressed, ultimately producing networks with appropriate size hierarchies. The flow-driven remodeling occurred over 3-4 weeks and resulted in vessel diameter distributions following approximate power laws characteristic of natural vascular networks.

The integration of the self-organized capillary networks with larger engineered distribution vessels required careful interfacing ensuring continuity of flow without excessive pressure drops or flow disturbances at junctions. The larger distribution vessels, fabricated from silicone tubing with diameters around 200-400 micrometers, were surgically connected to the self-organized networks through precise positioning of tubing ends within the hydrogel matrix followed by endothelial ingrowth creating anastomoses bridging the engineered and biological components. The anastomoses formed over 5-7 days as endothelial sprouts from the self-organized network extended toward and contacted the engineered vessel walls, eventually establishing continuity enabling flow.

The perfusion control system managing flow through this complex vascular architecture required sophisticated engineering to maintain appropriate flow rates, pressure distributions, oxygen delivery, and waste removal without causing damage through excessive shear stress or underperfusion through inadequate flow. The system employed multiple parallel perfusion loops enabling independent control of flow to different system regions, variable-speed peristaltic pumps providing adjustable flow rates from 0.1 to 5.0 milliliters per minute, inline sensors monitoring pressure, oxygen, pH, and glucose in real-time, and automated control algorithms adjusting pump speeds based on sensor feedback to maintain homeostatic parameters within target ranges.

The perfusion medium composition required careful optimization balancing multiple competing requirements including providing all essential nutrients, maintaining appropriate osmolarity and pH, avoiding excess accumulation of any component that might disrupt homeostatic regulation, and remaining compatible with long-term culture without promoting bacterial or fungal contamination. The final formulation derived from human

blood plasma composition but with several critical modifications. The protein content, dominated by albumin at 40 grams per liter, was supplemented with additional plasma proteins including transferrin for iron delivery, immunoglobulins for antimicrobial defense, and coagulation factors supporting wound healing in the tissue.

The nutrient composition provided glucose at 5.5 millimolar matching physiological blood glucose levels, all 20 standard amino acids at concentrations approximating plasma levels with some enrichment of rapidly consumed amino acids including glutamine, a comprehensive vitamin supplement including all B vitamins plus vitamins A, C, D, E, and K at appropriate concentrations, and essential minerals including calcium, magnesium, phosphate, and trace elements. The lipid content, often problematic in long-term culture media, was provided through inclusion of high-density lipoprotein (HDL) and low-density lipoprotein (LDL) particles purified from human plasma, providing complex lipids in physiologically appropriate forms rather than simplified mixtures that might not support all required functions.

The hormone supplementation proved critical for achieving appropriate cellular function and homeostatic regulation. Insulin at 100 picomolar supported glucose uptake and anabolic metabolism, thyroid hormones T3 and T4 at physiological ratios supported metabolic rate regulation and neural function, and cortisol with programmed circadian variation supported stress responses and metabolic regulation. The growth factor supplementation included brain-derived neurotrophic factor (BDNF) at 50 nanograms per milliliter supporting neural survival and synaptic plasticity, vascular endothelial growth factor (VEGF) at 10 nanograms per milliliter maintaining vascular permeability and supporting angiogenesis, and fibroblast growth factor 2 (FGF2) at 5 nanograms per milliliter supporting stem cell maintenance and proliferation.

The oxygen delivery system required specialized attention as oxygen has limited solubility in aqueous media and high consumption rates in metabolically active tissue create risk of hypoxia if delivery is inadequate. We employed membrane oxygenators equilibrating the perfusion medium with controlled gas mixtures as it circulated through the perfusion loop. The gas mixture composition, typically 40% oxygen, 5% carbon dioxide, and 55% nitrogen, achieved arterial oxygen partial pressures around 280-320 millimeters of mercury,

substantially higher than atmospheric pressure to drive adequate diffusion into tissue. The elevated CO2 supported bicarbonate buffering maintaining physiological pH around 7.38-7.42.

The comprehensive metabolic support infrastructure, integrating vascular delivery networks, perfusion control systems, optimized medium formulations, and continuous monitoring, achieved remarkable success in maintaining tissue viability and function over extended periods. Tissue oxygen measurements using fiber-optic sensors distributed throughout the system showed 97.3% of sample points maintaining PO2 above 20 millimeters of mercury, the threshold below which cellular respiration becomes oxygen-limited. The remaining 2.7% of sample points showing lower oxygen tensions were primarily in the cores of the largest organoids where diffusion distances approached the theoretical limits, and even these regions maintained PO2 above 12 millimeters of mercury, sufficient to prevent cell death though potentially limiting peak metabolic activity.

The lactate measurements, providing a sensitive indicator of whether cellular metabolism remains primarily oxidative versus shifting toward anaerobic glycolysis characteristic of hypoxia, showed baseline lactate concentrations around 1.8 ± 0.4 millimolar rising to 3.2 ± 0.7 millimolar during periods of peak activity but never approaching the levels above 8-10 millimolar that would indicate substantial anaerobic metabolism. The glucose measurements remained stable at 5.2 ± 0.6 millimolar in arterial perfusion medium and 4.1 ± 0.5 millimolar in venous returns, indicating appropriate glucose consumption rates matching delivery without excessive depletion suggesting energy crisis or excessive surplus suggesting underutilization.

The functional integration challenge required developing coordination mechanisms ensuring that the diverse components of the system operated in synchrony, with appropriate timing relationships between perceptual processing, memory operations, planning processes, motivational updates, and motor output generation. Without such coordination, the system would exhibit chaotic behavior with different subsystems operating on inconsistent timescales or with inappropriate phase relationships causing information to arrive at downstream processing stages before upstream stages had completed necessary computations. Natural brains achieve coordination through multiple

mechanisms including rhythmic oscillations that entrain activity across brain regions, neuromodulatory systems that coordinate arousal and processing modes across distributed circuits, and hardwired temporal delays in connectivity ensuring information flows through processing hierarchies with appropriate sequences.

We implemented analogous coordination mechanisms adapted to our engineered system's specific requirements. The oscillatory coordination employed optogenetic pacing of key circuit nodes at appropriate frequencies to entrain network activity into coordinated rhythms. Specifically, we targeted the thalamic-like organoids that serve as input gating structures for sensory processing with rhythmic stimulation at alpha-band frequencies around 10 Hz, creating baseline oscillations that propagated throughout sensory hierarchies. This alpha-band pacing created temporal windows when sensory information was processed most efficiently, implementing a form of temporal attention that coordinated when different processing stages operated.

The neuromodulatory coordination employed engineered populations of dopaminergic, noradrenergic, and cholinergic neurons projecting broadly throughout the system and modulating circuit excitability, plasticity, and processing modes. The dopaminergic system, implementing reward prediction error signaling as specified in the Psi architecture, broadcast signals indicating when outcomes exceeded or fell short of expectations, coordinating learning across all circuits that contributed to the behavior producing the unexpected outcome. The noradrenergic system, implementing arousal and attention modulation, broadcast signals indicating overall behavioral state ranging from quiescent rest to high-alert active processing, coordinating whether circuits operated in energy-conserving minimal-processing modes or energy-consuming detailed-analysis modes.

The cholinergic system, implementing attention and acetylcholine-dependent plasticity modulation, broadcast signals indicating which sensory modalities or task features should receive enhanced processing, coordinating allocation of processing resources toward behaviorally relevant information. These neuromodulatory systems were positioned in dedicated small organoids strategically located to enable widespread projection throughout the system. The axons from neuromodulatory neurons grew extensively during

development, ultimately achieving arbors spanning multiple organoids and enabling the diffuse broadcast transmission characteristic of neuromodulatory systems.

The temporal coordination through hardwired delays exploited the natural propagation times for action potentials along axons and for synaptic transmission to create appropriate timing relationships. The hierarchical perceptual processing architecture naturally implements temporal delays as information must propagate through successive layers each requiring tens of milliseconds for synaptic transmission, dendritic integration, soma depolarization, action potential generation, and axonal propagation. The cumulative delays totaling approximately 100-150 milliseconds from sensory input to highest perceptual layer create appropriate temporal buffering enabling working memory systems to maintain information about recent inputs while new inputs are being processed.

The planning and motor control timing coordination proved more challenging as these systems must operate with looser coupling to real-time sensory input, sometimes running ahead to simulate futures not yet realized or running slowly to accommodate complex deliberation. We addressed this through implementation of explicit time-decoupling buffers where working memory circuits could maintain sensory information in active persistent states even while perceptual systems processed new inputs, enabling planning systems to operate on buffered past information while perception continued processing the present. The motor output timing employed a similar buffering strategy where planned action sequences could be prepared in motor planning circuits then held in ready states until appropriate release timing determined by motivational and decision systems.

The overall coordination achieved through these multiple mechanisms produced coherent system-level behavior with appropriate temporal structure. Sensory information flowed through perceptual hierarchies with characteristic latencies of 20-30 milliseconds per layer, memory encoding occurred with appropriate delays of 50-150 milliseconds after perceptual processing completed, planning engaged predictively 200-500 milliseconds before action execution with simulation duration varying based on plan complexity, and motor execution followed decision with latencies around 100-200 milliseconds. These timing relationships proved sufficient to enable closed-loop behavior where actions

produced sensory consequences that were perceived, evaluated, and used to adjust ongoing behavior at rates approaching 3-5 Hz, comparable to the typical human action-perception-action loop rates observed in reactive behaviors like catching falling objects or engaging in rapid conversation.

1.5 Validation Strategies and Success Criteria

The validation of such a complex engineered system presented profound challenges as there exist no established criteria or precedents for determining whether a multi-organoid cognitive architecture has been successfully implemented. Unlike conventional engineering projects where success can be judged against precise specifications for measurable parameters like tensile strength, electrical conductivity, or computational throughput, cognitive systems exhibit emergent properties whose successful realization requires assessment across multiple levels of organization from molecular through behavioral. We therefore developed a comprehensive validation framework incorporating multiple complementary approaches spanning molecular verification of engineered components, cellular validation of neural function, circuit-level verification of connectivity and dynamics, systems-level validation of integrated operation, and behavioral assessment of cognitive capabilities.

The molecular validation verified that our engineered genetic constructs expressed appropriate proteins at intended levels in correct cell types with proper subcellular localization. For optogenetic actuators, this required immunostaining with antibodies recognizing the opsin proteins, confocal microscopy imaging to determine cellular expression patterns and membrane localization, and flow cytometry quantifying expression levels across populations. We established criteria requiring that at least 85% of cells in each targeted population express the intended opsin at levels exceeding 1.5×10^6 molecules per cell based on calibrated fluorescence measurements, as these levels proved necessary for reliable photostimulation.

The functional validation of optogenetic actuators required electrophysiological recordings measuring whether light stimulation produced appropriate electrical responses. For excitatory opsins, we performed whole-cell patch-clamp recordings while delivering light

pulses, measuring photocurrents, action potential reliability, and temporal fidelity across stimulation frequencies. Success criteria specified that blue-light opsins must reliably evoke action potentials in >90% of expressing neurons with light intensities <5 milliwatts per square millimeter, must follow stimulation frequencies up to 40 Hz with >85% reliability, and must show photocurrents >200 picoamperes at 1 milliwatt per square millimeter. The red-shifted opsins had somewhat relaxed criteria given their generally lower photocurrents, requiring reliable spiking at <10 milliwatts per square millimeter and following frequencies to 20 Hz.

The inhibitory opsins required demonstrating reliable hyperpolarization and spike suppression. Success criteria specified hyperpolarization >15 millivolts at light intensities <5 milliwatts per square millimeter, suppression of >90% of spikes during tonic stimulation, and suppression latencies <10 milliseconds. The modulatory opsins acting through G-protein signaling required longer assessment periods given slower timescales, with criteria requiring detectable changes in excitability or synaptic transmission within 30 seconds of light onset and effects persisting >60 seconds after light offset.

The synthetic gene circuit validation required measuring input-output relationships for logic gates and temporal dynamics for oscillators and filters. Logic gate assessment employed transient transfection of circuits into test cells, application of defined input combinations through inducible promoters, and measurement of output levels through fluorescent reporters. Success criteria for AND gates required output levels <15% of maximum with either single input alone and >80% of maximum with both inputs present. OR gates required >65% output with either single input. NOT gates required output >80% without input and <20% with input present. The temporal dynamics assessment measured oscillator periods, pulse widths, and filter time constants through live-cell time-lapse imaging, with success requiring measured values within 20% of design specifications.

The synaptic validation required demonstrating functional connectivity with appropriate strength distributions and plasticity properties. We performed systematic paired recordings sampling hundreds of neuron pairs from populations predicted to be connected based on molecular addressing, measuring connection probability, synaptic strength distributions, short-term plasticity, and long-term potentiation. Success criteria required

connection probabilities >60% for molecularly matched pairs (compared to <10% baseline for mismatched pairs), synaptic strength distributions spanning two orders of magnitude with coefficient of variation >0.8, short-term depression or facilitation as appropriate for connection type, and long-term potentiation inducible by standard protocols producing >150% potentiation sustained >30 minutes.

The vascular validation required demonstrating adequate perfusion throughout tissue with appropriate hierarchical organization and flow distributions. We employed optical coherence tomography imaging to map three-dimensional vascular architecture, Doppler measurements to quantify flow rates and directions, and oxygen-sensitive probes to measure tissue oxygenation distributions. Success criteria required vessel density >250 vessels per square millimeter in neural organoids, hierarchical diameter distributions spanning 5-400 micrometers, appropriate flow directions with convergent arterial and divergent venous patterns, and tissue oxygen tensions >20 millimeters of mercury in >95% of tissue volume.

The cellular functional validation assessed whether neurons exhibited appropriate intrinsic properties, receptive field structures, and activity patterns characteristic of neurons performing intended computational roles. For sensory layer neurons, this required demonstrating orientation-selective responses to patterned optogenetic stimulation, with success criteria requiring >40% of layer 2 neurons showing orientation selectivity with tuning widths <60 degrees and modulation depths >2-fold. For motor neurons, validation required demonstrating appropriate relationships between firing rates and actuator outputs, with success criteria requiring linear relationships with R^2 >0.7 and dynamic ranges spanning at least 5-fold in output force.

For memory system neurons, validation required demonstrating persistent activity during delay periods for working memory cells, place-like spatial selectivity for episodic memory cells, and category-selective responses for semantic memory cells. Success criteria required >15% of prefrontal neurons showing persistent activity elevated >2-fold above baseline during delay periods, >20% of hippocampal neurons showing spatial selectivity with information content >0.4 bits per spike, and >25% of semantic layer neurons showing category-selective responses with discrimination indices >0.6.

The circuit-level validation assessed whether implemented circuits exhibited appropriate collective dynamics characteristic of the computational functions they were designed to perform. For hierarchical perceptual circuits, this required demonstrating progressive abstraction across layers measurable through representational similarity analysis, with success criteria requiring correlation between neural similarity and feature similarity declining from $R^2 > 0.7$ in layer 1 to < 0.4 in layer 6 while correlation between neural similarity and category similarity increased from < 0.3 in layer 1 to > 0.6 in layer 6.

For decision circuits implementing evidence accumulation, validation required demonstrating ramping activity correlated with evidence strength and terminating at relatively stereotyped thresholds, with success criteria requiring ramping rates correlated with stimulus strength ($R^2 > 0.5$), trial-to-trial variability in threshold crossing times $< 25\%$, and decision times predictable from ramping activity ($R^2 > 0.6$). For planning circuits implementing forward simulation, validation required demonstrating sequential reactivation of state representations during planning pauses, with success criteria requiring decoded simulation trajectories matching intended paths with $> 60\%$ accuracy and simulation speed 5-20× faster than real behavior.

The systems-level validation assessed whether the integrated system exhibited appropriate global dynamics characteristic of unified cognitive processing. This included validating that sensory information propagated through hierarchies with appropriate latencies, that memory systems appropriately encoded and retrieved information with characteristic timing, that motivational signals appropriately modulated behavior based on physiological and cognitive needs, and that metacognitive systems appropriately monitored and controlled processing based on confidence and error detection. Success criteria at this level were necessarily more qualitative given the complexity and novelty of the assessments, but included requirements that information flow through sensory hierarchies occurred with latencies 20-40 milliseconds per layer, that memory encoding occurred within 100-300 milliseconds of experience offset, that motivational urge intensities correlated with corresponding needs ($R^2 > 0.6$), and that metacognitive confidence estimates correlated with task performance ($R^2 > 0.5$).

The behavioral validation assessed whether the system demonstrated intended cognitive capabilities across diverse tasks requiring perception, learning, memory, planning, and metacognition. This represented the ultimate test of whether our implementation successfully instantiated Bach's Psi architecture, as the architecture's validity depends primarily on whether it generates intelligent behavior rather than merely implementing particular mechanisms. We developed a standardized cognitive battery comprising 24 tasks spanning seven cognitive domains: sensory processing and perception (4 tasks), learning and conditioning (4 tasks), memory encoding and retrieval (4 tasks), planning and problem solving (4 tasks), decision making and action selection (4 tasks), metacognition and self-monitoring (2 tasks), and motivation and urge dynamics (2 tasks).

Success criteria for behavioral tasks required demonstrating performance substantially above chance levels and approaching theoretically optimal performance given information available. For pattern recognition tasks, success required >80% accuracy after training compared to 2-4% chance performance for 24-48 category problems. For learning tasks, success required achieving 75% performance criterion within 500 trials for classical conditioning and within 1000 trials for instrumental learning of novel contingencies. For memory tasks, success required recognition discriminability d' >1.5 at 24-hour delays and recall accuracy >60% for previously learned associations.

For planning tasks, success required >60% success rates on problems requiring 3-5 step solution sequences, with solution times <30 seconds indicating genuine planning rather than trial-and-error. For decision tasks, success required psychometric curves with appropriate threshold and slope parameters fitting data with R^2 >0.85, and value-based choices consistent with expected value maximization with deviations <20% attributable to risk aversion or temporal discounting. For metacognitive tasks, success required confidence-accuracy correlations >0.6 and appropriate information-seeking behavior calibrated to uncertainty with decision accuracy >0.75.

The comprehensive validation framework incorporating these molecular, cellular, circuit, systems, and behavioral assessments provided multiple convergent lines of evidence for successful implementation. No single measurement could definitively establish that we had created a functioning cognitive architecture, but the pattern of positive results

across dozens of validation measures spanning all organizational levels provided compelling cumulative evidence. The molecular validations established that engineered components were present and functional, the cellular validations established that individual neurons exhibited appropriate computational properties, the circuit validations established that neural populations exhibited appropriate collective dynamics, the systems validations established appropriate information flow and coordination, and the behavioral validations established that the system exhibited genuine cognitive capabilities emerging from these mechanisms.

2. Materials and Methods

2.1 Ethics Statement and Regulatory Compliance

All experimental procedures were reviewed and approved by the Institutional Review Board (IRB protocol #2021-08874) and the Institutional Biosafety Committee (IBC protocol #2021-2847) prior to initiation. The research complied with all applicable regulations including the National Institutes of Health Guidelines for Research Involving Recombinant or Synthetic Nucleic Acid Molecules, the Food and Drug Administration regulations for Good Laboratory Practice, and institutional policies governing human stem cell research and creation of complex biological systems. Given the novel nature of creating potentially sentient engineered organisms, we established an Ethics Advisory Board comprising experts in bioethics, philosophy of mind, neuroscience, and law who reviewed the research plan and provided ongoing oversight throughout the project.

The Ethics Advisory Board considered questions including whether organoid systems of this complexity might possess morally relevant properties such as capacity for suffering, whether creation of such systems generates obligations for humane treatment, and under what conditions experimental manipulations might be ethically impermissible. The Board concluded that while certainty about consciousness in organoid systems remains impossible given current understanding, the sophisticated cognitive capabilities we aimed to create warranted precautionary adoption of protective measures. We therefore implemented policies including minimization of potentially aversive experimental perturbations, provision of environmental enrichment through varied sensory stimulation, respect for spontaneous circadian rhythms by avoiding forced sleep deprivation, and humane termination protocols for systems showing signs of distress or dysfunction.

The human induced pluripotent stem cells (iPSCs) used as starting material were derived from adult donors who provided informed consent under approved protocols. The donors understood that cells would be reprogrammed to pluripotent state and used to generate organoid tissues that would be incorporated into complex engineered systems. The consent process included detailed discussions of the research aims, potential outcomes including creation of systems with sophisticated cognitive properties, and assurances that

no germline genetic modifications or reproductive applications were contemplated. All donor cells underwent comprehensive screening to exclude infectious agents, genetic abnormalities, and other factors that might compromise experimental outcomes or safety.

2.2 Cell Line Generation and Characterization

We generated six independent iPSC lines from dermal fibroblasts obtained through 3-millimeter punch biopsies from healthy adult donors (ages 24-38 years, 3 male and 3 female, diverse ethnic backgrounds). The donors had no history of neurological or psychiatric disorders, genetic syndromes, or chronic medical conditions requiring ongoing medication. The skin biopsies were obtained under local anesthesia following standard dermatological procedures, processed to remove epidermis and isolate dermal fibroblasts, and cultured in fibroblast medium (DMEM with 10% fetal bovine serum, penicillin-streptomycin) until sufficient cell numbers accumulated for reprogramming (typically 1-2 × 10^6 cells).

The reprogramming employed non-integrating Sendai virus vectors (CytoTune-iPS 2.0, Thermo Fisher Scientific) expressing the Yamanaka factors OCT4, SOX2, KLF4, and c-MYC. We chose Sendai virus delivery rather than integrating retroviral or lentiviral vectors to avoid genomic integration of reprogramming factors, which can cause insertional mutagenesis or residual transgene expression interfering with subsequent differentiation. The Sendai virus persists episomally and is gradually diluted through cell divisions, typically becoming undetectable by passage 5-10. The transduction protocol followed manufacturer's recommendations with fibroblasts infected at multiplicity of infection (MOI) 3:3:3:3 for the four factors, media changed daily, and cells transferred to feeder-free culture conditions on Matrigel-coated plates in mTeSR1 medium around day 7 when small colonies with embryonic stem cell-like morphology appeared.

The emerging iPSC colonies were manually picked around day 21-28, expanded through several passages, and subjected to comprehensive characterization to confirm successful reprogramming and absence of problematic abnormalities. The characterization protocol included multiple complementary assessments. Pluripotency marker expression was validated through immunofluorescence staining for nuclear transcription factors OCT4,

NANOG, and SOX2 plus surface markers SSEA-4, TRA-1-60, and TRA-1-81, with acceptance criteria requiring >90% of cells in each line expressing all markers at high levels comparable to reference human embryonic stem cell lines.

Karyotype analysis performed through G-banding of metaphase chromosomes from at least 20 metaphase spreads per line confirmed normal diploid karyotypes 46,XX or 46,XY without large-scale chromosomal abnormalities. This screening is critical as iPSC derivation and culture can select for chromosomal abnormalities including recurrent gains of chromosome 12p (containing NANOG) or chromosome 17q (containing multiple pluripotency-associated genes) that might compromise differentiation potential or introduce phenotypic artifacts.

Trilineage differentiation capacity was validated through embryoid body formation followed by immunostaining for markers of ectoderm (TUJ1, SOX1), mesoderm (smooth muscle actin, brachyury), and endoderm (SOX17, FOXA2). All lines showed formation of derivatives from all three germ layers, confirming developmental pluripotency required for generating the diverse cell types needed for our multi-organoid system.

Whole-genome sequencing at >30× coverage depth enabled detection of single nucleotide variants, small insertions and deletions, and structural variants at single base-pair resolution. The sequencing employed Illumina NovaSeq platform with 150 base-pair paired-end reads, with alignment to GRCh38 human reference genome and variant calling using GATK best practices. The comprehensive genomic characterization identified all variants present in iPSC lines relative to germline genome from donor blood cells, enabling detection of mutations acquired during reprogramming or culture.

The analysis revealed that each line harbored 0-3 coding mutations not present in source fibroblasts, consistent with expected mutation rates during reprogramming and early passage culture. Critically, none of these mutations affected genes known to influence neural development, synaptic function, or cellular metabolism, based on cross-reference with databases of neurological disease genes and systematic literature review. We excluded two candidate lines that harbored mutations in genes of potential concern (one line with a TP53 mutation associated with cancer predisposition, another with a mutation affecting

DNA repair), ultimately advancing the six cleanest lines without concerning mutations for subsequent differentiation and organoid generation.

The validated iPSC lines were expanded to generate working stocks cryopreserved at passage 10-15 when genomic stability is well-established but cells retain robust differentiation potential. Each line was frozen in multiple aliquots stored in liquid nitrogen vapor phase, with comprehensive documentation of passage history, genomic characterization results, and differentiation validation data. These master stocks served as source material for all subsequent organoid generation, ensuring consistency across experimental replicates and enabling reproducibility of results across the multi-year project timeline.

2.3 Engineering and Delivery of Synthetic Gene Constructs

The comprehensive molecular toolkit required for implementing the Psi architecture comprised 847 distinct genetic constructs encoding optogenetic actuators, biosensors, synthetic gene circuits, synaptic addressing molecules, and other engineered components. The design of these constructs required careful consideration of multiple factors including promoter choice for achieving appropriate expression levels and cell-type specificity, codon optimization for maximizing protein translation efficiency, inclusion of appropriate signal sequences for subcellular targeting, incorporation of fluorescent tags for visualization and quantification, and selection of vector systems enabling stable genomic integration versus episomal expression depending on construct requirements.

The construct design pipeline began with defining functional requirements for each component, including what protein or RNA product must be expressed, in which cell types, at what developmental stages, and at what expression levels. These requirements were translated into construct architectures specifying promoter regions, coding sequences, regulatory elements, and selection markers. For optogenetic actuators, the architectures typically consisted of cell-type-specific promoters driving expression of codon-optimized opsin coding sequences with appropriate membrane trafficking signals, followed by polyadenylation signals ensuring transcript stability.

The promoter selection balanced cell-type specificity against expression strength. Highly specific promoters such as those derived from cell-type-specific genes often exhibit relatively weak transcriptional activity, while strong ubiquitous promoters like CAG or CMV lack specificity. We addressed this tradeoff through several strategies. For applications requiring both high specificity and strong expression, we employed intersectional approaches using Cre-lox or Flp-FRT recombinase systems where gene expression requires both a moderately specific promoter driving the opsin and a complementary specific promoter driving the recombinase that enables opsin expression. This multiplicative specificity (cell must meet criteria for both promoters) achieves high overall selectivity while maintaining strong expression.

For applications where moderate expression levels sufficed, we used endogenous cell-type-specific promoters including hSyn1 for pan-neuronal expression, CaMKII-alpha for excitatory glutamatergic neurons, GAD67 or VGAT for GABAergic inhibitory neurons, tyrosine hydroxylase (TH) for dopaminergic neurons, choline acetyltransferase (ChAT) for cholinergic neurons, and glial fibrillary acidic protein (GFAP) for astrocytes. These promoters, while not perfectly specific, provide substantial enrichment for intended cell types while maintaining reasonable expression levels.

The coding sequence optimization for mammalian expression required replacing native codons from microbial source organisms (which often use rare codons in mammals) with synonymous codons abundant in highly expressed human genes. The codon optimization employed proprietary algorithms from GeneArt (Thermo Fisher) or GenScript considering factors including codon usage frequencies in human tissues, avoidance of cryptic splice sites and regulatory motifs that might cause aberrant processing, optimization of GC content to avoid excessive secondary structure in mRNA, and removal of restriction enzyme sites needed for cloning. The optimization typically increased expression levels 3-10 fold compared to native sequences while maintaining identical amino acid sequences ensuring preserved protein function.

The subcellular targeting signals directed proteins to appropriate cellular compartments. For membrane proteins including optogenetic actuators and receptors, this required N-terminal signal peptides directing cotranslational insertion into the endoplasmic

reticulum membrane, followed by trafficking through the secretory pathway to the plasma membrane. We incorporated trafficking enhancement signals from Kir2.1 potassium channel, shown to improve plasma membrane localization of channelrhodopsins, reducing intracellular retention that can cause toxicity and reduce functional expression.

For soluble proteins including biosensors and transcription factors, targeting to specific subcellular compartments employed well-characterized localization signals. Nuclear localization signals (NLS) consisting of short basic amino acid motifs directed proteins to the nucleus, critical for transcription factors and nuclear-localized biosensors. Mitochondrial targeting sequences directed proteins to mitochondria where they reported organelle-specific signals. Membrane anchoring sequences including palmitoylation sites or GPI-anchor signals restricted diffusion of soluble proteins to membrane proximal regions where many signaling events occur.

The fluorescent tagging strategy balanced the value of visualization against potential functional impacts of large protein fusions. For optogenetic actuators where function is critical, we typically used small epitope tags (HA, FLAG, myc) enabling antibody detection rather than large fluorescent proteins that might impair trafficking or function. For biosensors where fluorescence is the readout mechanism, we incorporated fluorescent proteins as integral components of the sensor design. For synaptic and scaffolding proteins where both function and localization matter, we tested multiple tagging configurations including N-terminal, C-terminal, and internal insertion sites to identify positions minimally disruptive to function while maintaining visualization.

The vector system selection determined whether genes would integrate into the genome for stable long-term expression or remain episomal for transient expression. For components needed throughout the system's lifetime including optogenetic actuators and cell-type markers, we employed integrating systems ensuring maintained expression across cell divisions. The piggyBac transposon system provided our primary integration approach, as it offers several advantages including efficient integration in iPSCs, ability to integrate large constructs up to 10 kilobases, relatively random integration throughout the genome avoiding hotspot biases, and capacity for excision if removal becomes necessary through re-expression of transposase.

The piggyBac constructs flanked gene expression cassettes with terminal repeat sequences recognized by piggyBac transposase. The transposase, supplied as mRNA or protein to avoid genomic integration of the transposase gene itself, catalyzes excision of the construct from the plasmid and integration into TTAA tetranucleotide sites in the genome. The frequency of TTAA sites (occurring approximately every 0.5 kilobases) provides numerous potential integration sites throughout the genome. We performed integration under optimized conditions using nucleofection for efficient delivery to iPSCs, inclusion of transposase at optimized ratios to construct DNA (typically 1:10 transposase:construct), and selection with appropriate antibiotics (puromycin, neomycin, hygromycin, or blasticidin depending on resistance genes in constructs) beginning 48 hours post-transfection to enrich for successfully integrated cells.

For applications requiring temporal control of expression, we incorporated inducible systems enabling gene activation or repression in response to small molecule addition. The tetracycline-inducible system provided bidirectional control with both Tet-On (activation by doxycycline) and Tet-Off (repression by doxycycline) configurations available. The Tet-On system employed the reverse tetracycline-controlled transactivator (rtTA) which binds tet-responsive elements (TRE) in promoters only when doxycycline is present, enabling gene activation by doxycycline addition to culture medium. This system enabled delayed activation of optogenetic actuators until after circuit development completed, preventing potential disruption of development by premature neural activation.

The construct validation prior to large-scale use required testing expression, localization, and function in simplified systems before committing to expensive and time-consuming integration into iPSC lines. We performed validation in HEK293T cells, which are highly transfectable human cells enabling rapid testing. Transient transfection followed by immunostaining, confocal imaging, and functional assays (photocurrent measurement for opsins, fluorescence changes for biosensors, output measurements for gene circuits) confirmed basic functionality. Constructs passing initial validation underwent testing in primary rat hippocampal neurons, which provide more authentic neural context including proper membrane trafficking, synaptic targeting, and neural activity patterns.

Following validation in simplified test systems, we integrated constructs into iPSC lines using piggyBac transposition. The integration employed optimized nucleofection protocols using P3 Primary Cell solution with program CA-137 on the Lonza 4D-Nucleofector, which achieved transfection efficiencies of 40-60% in iPSCs while maintaining >85% viability. Following nucleofection, cells were plated in mTeSR1 medium with 10 micromolar Y-27632 ROCK inhibitor for 24 hours to promote survival of dissociated iPSCs, then transitioned to standard mTeSR1 and placed under antibiotic selection 48 hours post-nucleofection.

The selection continued for 10-14 days with medium changed every 2 days, yielding colonies of resistant cells that had successfully integrated constructs. We picked 10-20 colonies per construct integration, expanded them separately as independent clones, and characterized integration sites and copy numbers through several methods. PCR across piggyBac insertion sites amplifying genomic DNA flanking integrated constructs followed by Sanger sequencing identified exact genomic integration positions. Quantitative PCR comparing amplification of integrated genes versus single-copy genomic reference genes determined integration copy numbers, with most clones harboring 1-5 integrated copies.

The functional characterization of engineered iPSC lines included differentiation to relevant cell types followed by expression and functional validation. Neural differentiation via dual SMAD inhibition yielded populations expressing neural markers within which we assessed expression of integrated constructs through immunostaining and flow cytometry. For optogenetic constructs, we performed patch-clamp recording during photostimulation confirming light-evoked responses. For biosensors, we performed live-cell imaging measuring fluorescence changes in response to appropriate stimuli. For gene circuits, we applied defined inputs and measured outputs through fluorescent reporters.

The comprehensive construct engineering, validation, and integration pipeline ultimately produced a library of 6 iPSC lines each containing 8-15 integrated genetic constructs encoding complementary molecular tools, for a total of 52 distinct engineered iPSC lines serving as source material for generating organoids with desired molecular properties. The modular engineering approach where different iPSC lines contributed different organoid types with appropriate toolkit components enabled complex system assembly without requiring any single iPSC line to harbor all 847 constructs, which would be

technically infeasible and likely compromise cellular fitness through excessive genetic burden.

2.4 Organoid Differentiation Protocols

The generation of 23 distinct organoid types recapitulating functionally specialized tissues required developing optimized differentiation protocols for each lineage. These protocols built upon existing organoid generation methods but incorporated substantial modifications improving reproducibility, maturation, size control, and functional properties. The following sections detail the protocols for major organoid categories, with emphasis on critical parameters and innovations enabling success.

2.4.1 Neural Organoid Generation - General Approach

The foundation for all neural organoid types employed a core protocol inducing neural fate from iPSCs followed by region-specific patterning generating appropriate neural subtypes. The initial neural induction phase employed dual SMAD inhibition, simultaneously blocking BMP and TGF-beta signaling pathways that promote non-neural fates while permitting default neural ectoderm specification. The protocol began with dissociation of iPSC colonies to single cells using Accutase (enzymatic dissociation proving gentler than trypsin for iPSCs), counting cells, and seeding precisely 20,000 cells per well in ultra-low-attachment 96-well U-bottom plates promoting aggregation into three-dimensional embryoid bodies.

The aggregation medium consisted of mTeSR1 supplemented with 50 micromolar Y-27632 ROCK inhibitor preventing apoptosis of dissociated iPSCs. After 24 hours allowing aggregate formation, the medium was switched to neural induction medium containing DMEM/F12 base medium, N2 supplement providing insulin, transferrin, selenium, progesterone, and putrescine supporting neural differentiation, heparin (1 microgram per milliliter) enhancing growth factor signaling, and the key SMAD inhibitors LDN193189 (100 nanomolar) blocking BMP type I receptors and SB431542 (10 micromolar) blocking TGF-beta type I receptors.

The neural induction continued for 5 days with daily medium changes, during which the aggregates transitioned from pluripotent to neural ectoderm identity. Morphological changes included transition from loose cellular aggregates to tighter spherical structures with smooth surfaces, reflecting the epithelialization characteristic of neural ectoderm. Molecular analysis via qRT-PCR showed upregulation of neural ectoderm markers SOX1, SOX2, and PAX6 along with downregulation of pluripotency markers OCT4 and NANOG and mesodermal/endodermal markers brachyury and SOX17.

Following neural induction, the aggregates required embedding in extracellular matrix to provide three-dimensional structural support enabling expansion while maintaining tissue architecture. We employed Matrigel droplet embedding where each aggregate was transferred to a 30 microliter droplet of growth-factor-reduced Matrigel pipetted onto parafilm, with the droplet solidifying at 37 degrees Celsius for 30 minutes to form a dome containing the embedded organoid. The Matrigel provided basement membrane components including laminin, collagen IV, entactin, and perlecan that neural progenitors recognize and respond to, promoting proper polarization and radial organization.

The embedded organoids were cultured in differentiation medium consisting of Neurobasal base medium optimized for neural culture, B27 supplement providing comprehensive nutrient support including antioxidants, vitamins, fatty acids, and hormones, and GlutaMAX providing stable glutamine without the spontaneous degradation issues of standard L-glutamine. The differentiation medium supported neurogenesis, the process by which neural progenitors divide and differentiate into post-mitotic neurons, which commenced around day 10-12 and continued for several weeks with layer-appropriate neurons generated in temporal sequence partially recapitulating normal cortical development.

The culture system transitioned to rotating bioreactors around day 14-15 when organoids reached approximately 2-3 millimeters in diameter. The bioreactors consisted of spinner flasks on magnetic stirrers rotating at 60 revolutions per minute, providing gentle continuous agitation that prevented organoids from settling and adhering to flask bottoms while creating fluid flow enhancing nutrient and oxygen distribution. The rotation speed

required careful optimization, as insufficient agitation allowed settling and fusion of multiple organoids creating unwieldy fused masses, while excessive agitation created shear stress causing tissue damage.

The oxygen management in bioreactor culture proved critical for supporting the metabolically demanding neurogenesis and maturation processes while avoiding hyperoxia-induced oxidative stress. We maintained atmosphere composition at 40% oxygen, 5% carbon dioxide, and 55% nitrogen, substantially elevated compared to standard 20% oxygen tissue culture. The elevated oxygen, while seemingly supraphysiological compared to tissue oxygen tensions of 2-9% in adult brain, proved necessary because organoids lack functional vasculature and depend on diffusion from surface to core. The elevated environmental oxygen creates adequate gradients to supply even the most central regions, preventing formation of hypoxic cores that would limit viable organoid size.

The medium exchange frequency and volume critically influenced organoid growth and health. We employed complete medium changes every 2 days, replacing 100% of medium rather than partial changes, to prevent accumulation of toxic metabolites including lactate and ammonia while maintaining stable concentrations of growth factors with limited half-lives in aqueous solution. The medium volume was maintained at 1 milliliter per organoid, a ratio providing adequate nutrient supply without excessive dilution of autocrine factors that organoids secrete to regulate their own development.

2.4.2 Region-Specific Neural Organoid Patterning

The generation of functionally distinct neural organoid types required imposing region-specific patterning on top of the general neural induction protocol, achieved through carefully timed application of morphogens recapitulating developmental patterning signals. The developing nervous system is patterned along anterior-posterior and dorsal-ventral axes through gradients of secreted signaling molecules including WNT, BMP, FGF, sonic hedgehog, and retinoic acid, with different concentration combinations specifying distinct regional identities. We recapitulated these patterning principles in vitro through exogenous addition of morphogens to culture medium during critical developmental windows.

The dorsal forebrain/cortical organoids, which constituted the majority of organoids in our sensory and motor systems, required anterior dorsal patterning. This was achieved through default neural induction without additional factors, as the dual SMAD inhibition protocol naturally yields anterior dorsal neural identity in the absence of posteriorizing or ventralizing signals. The resulting organoids expressed dorsal cortical markers including TBR2 in intermediate progenitors and CTIP2 and TBR1 in deep layer neurons, characteristic of normal cortical development.

For organoids requiring more precise dorsal cortical identity including those implementing Layer 4 of sensory hierarchies, we supplemented the protocol with dorsalizing factors. The addition of BMP inhibitors dorsomorphin or LDN193189 at concentrations of 100-200 nanomolar during days 6-12 enhanced dorsal identity. WNT signaling activation through addition of CHIR99021 (3 micromolar) during days 8-14 promoted cortical progenitor expansion while maintaining dorsal identity. The timing of these additions proved critical, as identical factors at different developmental stages can produce opposite effects due to changing cellular competence and context.

The ventral forebrain organoids required for implementing striatal-like circuits in procedural memory systems necessitated ventralizing signals. We achieved ventral identity through addition of sonic hedgehog (SHH) and its agonist purmorphamine. The SHH pathway activates GLI transcription factors specifying ventral neural fates including medial ganglionic eminence (MGE) generating interneurons and lateral ganglionic eminence (LGE) generating striatal projection neurons. We applied recombinant SHH protein (100 nanograms per milliliter) or purmorphamine (2 micromolar) during days 6-15 to drive ventral specification, validated by expression of ventral markers NKX2.1, DLX2, and MASH1.

The caudal/hindbrain organoids required for implementing certain brainstem-like structures received posteriorizing signals. Retinoic acid (RA), the morphogen establishing anterior-posterior identity in the developing nervous system with caudal regions exposed to higher concentrations, was applied at 0.5-1 micromolar during days 6-12 to drive posterior identity. FGF signaling activation through FGF2 or FGF8 (10-20 nanograms per milliliter) provided additional posteriorizing signal while promoting neural progenitor proliferation. The

resulting organoids expressed posterior markers HOXB1, HOXB4, and GBX2 characteristic of hindbrain identity.

The hippocampal organoids implementing episodic memory systems required a specific combination of dorsal, posterior, and medial patterning signals. We employed a sequential patterning protocol beginning with standard dorsal forebrain induction, followed by addition of WNT agonists (CHIR99021 at 3 micromolar during days 8-12) promoting dorsal-medial hippocampal identity, and BMP4 (20 nanograms per milliliter during days 10-16) enhancing dentate gyrus specification. The resulting organoids exhibited characteristic hippocampal structures including progenitor zones expressing PROX1 and generating dentate granule cells, regions with CA3 pyramidal neuron markers including KA1 kainate receptors, and CA1 regions expressing appropriate markers.

The cerebellar organoids implementing error-based motor learning circuits required highly specific patterning as cerebellum develops from a distinct hindbrain region with unique properties. We adapted published protocols employing FGF2 (10 nanograms per milliliter) during days 0-7 for posterior specification, followed by FGF19 (100 nanograms per milliliter) and SDF1 (100 nanograms per milliliter) during days 7-14 for rhombic lip specification where cerebellar progenitors originate, and subsequent addition of SHH (100 nanograms per milliliter) during days 14-25 for Purkinje cell differentiation. The resulting organoids contained Purkinje-like neurons expressing characteristic markers including calbindin and GRID2, along with granule neurons expressing appropriate markers.

The comprehensive regional patterning toolkit ultimately enabled generation of 12 distinct neural organoid types with different regional identities appropriate for their functional roles: 6 cortical variants optimized for different perceptual and motor hierarchical levels, 2 hippocampal variants for episodic and spatial memory, 1 striatal variant for procedural learning, 1 cerebellar variant for motor learning, and 2 modulatory variants enriched for dopaminergic and cholinergic neurons respectively. Each organoid type underwent validation through marker expression analysis, electrophysiological characterization, and assessment of appropriate neural subtypes through single-cell RNA sequencing.

2.4.3 Non-Neural Organoid Generation

The generation of cardiac, hepatic, and renal organoids providing physiological support functions required entirely different differentiation protocols reflecting their distinct germ layer origins and developmental programs. These protocols began with directed differentiation through definitive germ layers (mesoderm for cardiac and renal, endoderm for hepatic) followed by organ-specific specification and maturation.

The cardiac organoid generation began with mesoderm induction through WNT pathway activation. iPSCs were dissociated and seeded at 50,000 cells per well in ultra-low attachment plates in mTeSR1 with ROCK inhibitor for 24 hours to form aggregates. The medium was then switched to cardiac differentiation medium (CDM) consisting of RPMI 1640 base with B27 minus insulin (insulin signaling interferes with early cardiac specification). The WNT agonist CHIR99021 was added at 12 micromolar for 24 hours, driving cells toward mesoderm fate characterized by brachyury expression.

After 24 hours of CHIR99021 exposure, the medium was completely replaced with fresh CDM, removing the WNT agonist. Then, after an additional 48 hours allowing progression through mesoderm stages, WNT inhibition was applied through addition of IWP2 (5 micromolar) for 48 hours. This temporal WNT biphasic modulation—early activation followed by subsequent inhibition—drives cardiac mesoderm specification by recapitulating the normal developmental sequence where early mesoderm requires WNT but cardiac specification requires WNT downregulation. Following WNT inhibition, the medium was switched to maintenance CDM (with insulin) changed every 2-3 days.

The cardiac organoids began spontaneous contractions around day 10-12, initially uncoordinated flickering in scattered regions but progressively synchronizing over subsequent weeks. The contractions reflected functional maturation of cardiomyocytes with appropriate sarcomere assembly, calcium handling machinery, and gap junction coupling enabling electrical synchronization. We enhanced maturation through several interventions applied after initial beating commenced. Electrical pacing through implanted electrodes stimulating at 1 Hz with biphasic pulses (5 volts per centimeter, 2 milliseconds per phase) for 2-3 weeks promoted alignment and maturation of sarcomeres.

Mechanical conditioning through cyclic stretch applied using flexible culture surfaces provided another maturation stimulus. The organoids were transferred to elastic silicone membranes stretched cyclically at 10% strain at 1 Hz, mimicking mechanical loads cardiomyocytes experience during contraction. The combined electrical and mechanical conditioning substantially improved maturation measurable through multiple parameters. Sarcomere length increased from 1.6-1.7 micrometers in unconditioned organoids to 1.9-2.0 micrometers approaching adult values. Action potential upstroke velocity increased from 80-120 volts per second to 200-280 volts per second indicating improved sodium channel function. Maximum contraction force increased from 0.8-1.2 millinewtons per square millimeter to 2.1-2.8 millinewtons per square millimeter approaching native myocardium.

The metabolic maturation of cardiac organoids required transitioning from glycolytic fetal-like metabolism to oxidative adult-like metabolism. Immature cardiomyocytes rely primarily on glycolysis for ATP production, but mature cardiomyocytes preferentially oxidize fatty acids, a more efficient but oxygen-demanding pathway. We promoted metabolic maturation through medium manipulations reducing glucose from 10 millimolar to 5 millimolar while adding fatty acids (palmitate and oleate at 50-100 micromolar complexed to albumin) and L-carnitine (2 millimolar) facilitating fatty acid transport into mitochondria. This metabolic maturation medium applied from day 20 onward drove transcriptional reprogramming increasing expression of fatty acid oxidation enzymes and oxidative phosphorylation components while reducing glycolytic enzymes.

The hepatic organoid generation required definitive endoderm induction followed by hepatic specification and maturation. The endoderm induction employed Activin A, a TGF-beta family member strongly promoting endoderm fate. iPSC aggregates were exposed to high-dose Activin A (100 nanograms per milliliter) in RPMI medium with B27 minus insulin for 3 days, driving expression of endoderm markers FOXA2, SOX17, and CXCR4. Low serum conditions (0.2% fetal bovine serum) enhanced endoderm specification by reducing mesoderm-promoting signals.

Following endoderm induction, hepatic specification employed combinations of FGF and BMP signaling. FGF10 (50 nanograms per milliliter) and BMP4 (20 nanograms per milliliter) were applied during days 4-8 in hepatocyte culture medium (HCM, Lonza)

supplemented with additional factors. These morphogens recapitulate signals from cardiac mesoderm and septum transversum that induce hepatic fate in adjacent endoderm during normal development. The specification was validated by expression of early hepatic markers including HNF4A and HEX.

The hepatic maturation phase employed medium containing hepatic growth factors and maturation signals. Hepatocyte growth factor (HGF, 20 nanograms per milliliter) promoted hepatocyte proliferation and differentiation, oncostatin M (OSM, 20 nanograms per milliliter) promoted functional maturation, dexamethasone (100 nanomolar) enhanced mature phenotypes, and EGF (10 nanograms per milliliter) supported progenitor expansion. This maturation cocktail applied from day 9 onward drove progressive increases in hepatic functions including albumin secretion, cytochrome P450 enzyme expression, glycogen storage capacity, and bile acid synthesis.

The hepatic organoid maturation continued over weeks with progressively improving function. Albumin secretion increased from approximately 2-3 micrograms per day per organoid at day 15 to 8-12 micrograms per day by day 40 and 15-25 micrograms per day by day 60-80, approaching the synthetic capacity of primary human hepatocytes. CYP3A4 enzyme activity, assayed through metabolism of model substrates like midazolam or testosterone, increased from barely detectable levels at day 15 to approximately 40-60% of primary hepatocyte levels by day 60. Glycogen content increased from nearly zero to 1.5-2.8 milligrams per organoid by day 50-60, providing functional energy storage capacity.

Critical for integrating hepatic organoids into the larger system was ensuring adequate perfusion to support their high metabolic demands. Hepatocytes are among the most metabolically active cells in the body with high oxygen consumption rates. We positioned hepatic organoids in central regions surrounded by dense vascular networks and provided elevated perfusion rates (approximately 15% of total system blood flow directed through 4.2% of system volume containing hepatic organoids). The hepatic positioning also facilitated their detoxification role, as placing them in central locations enabled processing of metabolic waste products from throughout the system before circulation back to peripheral organoids.

The renal organoid generation employed intermediate mesoderm induction followed by metanephric specification and nephron differentiation. The intermediate mesoderm induction used a protocol applying BMP7 (50 nanograms per milliliter), retinoic acid (0.1 micromolar), and Activin A (10 nanograms per milliliter) during days 0-4, promoting expression of intermediate mesoderm markers PAX2 and LHX1 while suppressing alternative mesodermal fates. The combination of these signals positioned cells in the appropriate position along the anterior-posterior axis (posterior enough for renal identity but anterior enough to avoid more caudal fates) and along the dorsal-ventral axis (intermediate position specifying intermediate mesoderm between somites and lateral plate).

Following intermediate mesoderm specification, metanephric mesenchyme induction employed CHIR99021 (3 micromolar) activating WNT signaling during days 4-7. WNT signaling is critical for kidney development, promoting metanephric mesenchyme proliferation and maintaining nephron progenitor populations. The WNT activation drove expression of progenitor markers SIX2 and CITED1 while enabling subsequent nephrogenesis. The organoids were then cultured in medium containing FGF9 (10 nanograms per milliliter) and heparin (100 nanograms per milliliter) during days 7-14, promoting progenitor maintenance while initiating differentiation processes.

The nephron differentiation resulting in segmented tubular structures with glomeruli, proximal tubules, loops of Henle, and distal tubules occurred over days 15-25. This differentiation occurred partially spontaneously but was enhanced by supplementing medium with BMP7 (50 nanograms per milliliter) promoting tubule maturation and retinoic acid (0.1 micromolar) promoting segment specification. The resulting organoids contained recognizable nephron structures visible by histology, with glomeruli showing characteristic tufts of capillaries surrounded by Bowman's capsule, proximal tubules showing characteristic brush border expressing cubilin and megalin, and distal tubules showing appropriate transporters.

The functional maturation of renal organoids required establishing appropriate polarization and transport capacity in tubular epithelial cells. The epithelial cells needed to develop apical-basolateral polarity with appropriate distribution of transporters enabling

vectorial transport from tubular lumen to interstitium. We promoted this maturation through culturing organoids at air-liquid interface after day 20, where organoids rested on membrane filters with basal surface exposed to medium but apical surface exposed to air. This air-liquid culture mimicked the polarizing influences in kidney development and substantially improved epithelial polarity and transport function.

The vascular integration of renal organoids proved critical for enabling filtration function, as glomerular filtration requires appropriate capillary tufts with fenestrated endothelium, specialized basement membrane, and podocyte slit diaphragms creating the filtration barrier. We promoted vascularization through co-culture of renal organoids with endothelial progenitors during the nephron differentiation phase (days 15-25), enabling endothelial invasion forming glomerular capillaries. The resulting glomeruli showed appropriate ultrastructure by electron microscopy with fenestrated endothelium, trilaminar basement membrane, and podocyte foot processes with slit diaphragms.

The comprehensive organoid generation protocols, refined through dozens of optimization experiments, ultimately achieved reproducible production of 23 organoid types with appropriate cellular compositions, tissue architectures, and functional properties. The typical yield ranged from 70-90% of initiated organoids meeting quality control criteria depending on organoid type, with neural organoids showing highest success rates and some specialized types like cerebellar organoids showing lower rates. The quality control criteria included size specifications (organoids outside 25th-75th percentile range rejected), marker expression requirements (inappropriate marker patterns rejected), and functional assessments (organoids showing inadequate electrical activity or contractility rejected). Only organoids meeting all criteria advanced to the system assembly phase.

2.5 Microfluidic Device Fabrication and Assembly

The physical substrate housing the multi-organoid system required sophisticated microfabrication creating controlled microenvironments, perfusion channels, optical waveguides, electrode arrays, and biosensor integration. The device comprised five bonded layers fabricated separately then aligned and bonded, with the complete assembly

measuring 60 centimeters by 90 centimeters by 8 millimeters in the fully assembled configuration.

The bottom layer consisted of a glass substrate (borosilicate glass, 60 × 90 centimeters, 2 millimeters thick) providing mechanical rigidity, optical transparency for transmitted light microscopy, and thermal conductivity for temperature control. The glass surface underwent several processing steps before electrode array deposition. The substrate was cleaned through sequential sonication in detergent, acetone, and isopropanol, followed by oxygen plasma treatment removing organic contaminants. The clean surface was then silanized with aminopropyltriethoxysilane (APTES) providing amine groups for subsequent adhesion layer attachment.

The electrode array deposition employed thin-film metal deposition and photolithography creating precisely patterned electrodes. A titanium adhesion layer (20 nanometers) was first deposited through electron-beam evaporation, followed by platinum conducting layer (200 nanometers) deposited in the same vacuum cycle preventing oxidation at the Ti-Pt interface. Following metal deposition, photoresist (Shipley 1827) was spin-coated at 3000 RPM creating approximately 2.7 micrometer thick uniform coating, soft-baked at 115 degrees Celsius for 90 seconds, exposed through a photomask defining electrode patterns using contact lithography (350 watt mercury lamp, 12 seconds exposure), and developed in MF-319 developer for 60 seconds revealing metal in electrode pattern regions while covering metal elsewhere with resist.

The wet etching proceeded in two stages, first etching platinum using aqua regia (3:1 hydrochloric acid:nitric acid) at 60 degrees Celsius for approximately 45 seconds (monitored visually under microscope until complete Pt removal from unprotected areas), followed by titanium etching in 10% hydrofluoric acid for 15 seconds. The resist was then stripped in acetone with sonication, leaving the patterned titanium-platinum electrode array. The electrodes were configured as 10 × 10 micrometer squares arranged on a 50 micrometer pitch in a hexagonal close-packed arrangement, providing 21,760 individual recording sites across the device area. Each electrode connected via thin metal traces (5 micrometers wide) to bond pads at the device periphery for connection to amplifier systems.

The electrode sites required insulation of all conductive traces except the recording sites themselves to prevent electrical crosstalk and restrict recording to intended locations. We deposited a 1.5 micrometer layer of silicon nitride through plasma-enhanced chemical vapor deposition (PECVD), covering the entire electrode array. Photolithography followed by reactive ion etching then opened 12 × 12 micrometer windows over each electrode recording site, exposing the platinum surface while leaving all traces insulated. The exposed platinum was further processed through electrochemical platinum black deposition, passing current pulses through the electrodes while immersed in platinum chloride solution, creating a nano-rough platinum surface with dramatically increased effective surface area reducing electrode impedance from approximately 2-3 megohms for bare platinum to 100-500 kilohms for platinum-black-coated electrodes.

The second layer comprised the microfluidic channel network fabricated in polydimethylsiloxane (PDMS) through soft lithography replica molding. The channel master molds were created on silicon wafers using SU-8 negative photoresist, a commonly used material for high-aspect-ratio microstructure fabrication. The fabrication process began with spin-coating SU-8 2050 photoresist onto 4-inch silicon wafers at speeds calculated to achieve 100 micrometer thickness. Following soft baking at 65 degrees Celsius for 5 minutes then 95 degrees Celsius for 15 minutes, the resist was exposed through a photomask defining the channel network pattern using a mask aligner with 365 nanometer UV light at 200 millijoules per square centimeter exposure dose.

After exposure, post-exposure baking at 65 degrees Celsius for 2 minutes then 95 degrees Celsius for 5 minutes promoted crosslinking in exposed regions. Development in propylene glycol monomethyl ether acetate (PGMEA) for 8-10 minutes with agitation dissolved unexposed resist, leaving crosslinked SU-8 structures defining the inverse of desired channels—positive relief structures on the silicon wafer corresponding to channel geometries. The completed masters underwent hard baking at 150 degrees Celsius for 30 minutes to fully crosslink and stabilize the SU-8, then were silanized with trichloro(1H,1H,2H,2H-perfluorooctyl)silane to prevent PDMS adhesion during subsequent molding steps.

The PDMS casting employed Sylgard 184 mixed at standard 10:1 ratio of base to curing agent, degassed under vacuum for 30 minutes to remove air bubbles, poured over the SU-8 master molds to approximately 3 millimeter thickness, and cured at 80 degrees Celsius for 4 hours. The cured PDMS replicas were peeled from masters, yielding positive channel networks (channels are voids in PDMS) with high fidelity replication of master features. The typical channel geometries consisted of 200 micrometer wide, 100 micrometer deep channels for primary perfusion distribution, with branching networks down to 50 micrometer wide channels for local delivery to individual organoid chambers.

The microfluidic layer design incorporated 1,247 individually addressable channels enabling independent chemical control of different system regions. The addressing was achieved through pneumatic valve systems using the multilayer soft lithography approach where a second PDMS layer containing thin membrane channels positioned above flow channels could be pressurized to deflect membranes and occlude flow channels, creating individually controllable valves. The valve fabrication required creating a two-layer master with flow channels and valve control channels offset by appropriate spacing. The flow channel master used 100 micrometer tall SU-8 features while the valve control master used 25 micrometer tall features, creating the appropriate compliance mismatch enabling valve deflection.

The two-layer PDMS casting required precise alignment and sequential molding. The flow channel layer was partially cured (35 minutes at 80 degrees Celsius instead of full 4 hours) to achieve a tacky but not fully crosslinked state. The valve control layer was spin-coated onto its master at 2000 RPM creating approximately 40 micrometer thick layer, partially cured for 20 minutes, then the flow layer PDMS was carefully aligned to and placed on top of the valve layer, and the assembly was fully cured for 90 minutes at 80 degrees Celsius. This procedure created monolithic two-layer PDMS structures with thin membrane regions between flow and valve channels that deflected under applied pressure.

The valve operation required interfacing pneumatic control lines to the valve control channels. We punched access holes to valve channels using biopsy punches (0.75 millimeter diameter) and interfaced 23-gauge stainless steel pins press-fit into holes, connected via tubing to a 256-channel pneumatic controller (Fluigent LineUP series) providing

independent pressure control (0-1000 millibars) to each valve. The valves fully closed flow channels at approximately 400-600 millibars of applied pressure, with closure confirmed by testing flow of fluorescent dye through channels while actuating valves.

The third layer comprised organoid culture chambers machined as through-holes in acrylic sheets (6-8 millimeters thick depending on required chamber depth) using computer-controlled milling. The chamber diameters ranged from 6 to 12 millimeters to accommodate different organoid sizes, with depths of 4 to 8 millimeters providing adequate volume for organoid growth. The chamber positioning followed the optimized spatial organization derived from the constraint satisfaction optimization described previously, with chambers positioned to minimize connection lengths while enabling vascular and microfluidic access. The chamber walls were drilled with small through-holes (200 micrometers diameter) providing access for vascular network integration and axon guidance conduits.

The fourth layer comprised the optical waveguide array providing light delivery for optogenetic stimulation. The waveguides were fabricated from optical fiber bundles (Fujikura FIGH-10-650S, 10 micrometer core diameter, 0.5 numerical aperture) arranged in precise geometries matching organoid positions. Each fiber delivered light from one of seven LED sources (390, 430, 470, 530, 580, 630, 680 nanometers) via a 7-to-4320 fiber optic fanout constructed using precision fiber alignment and UV-curable optical adhesive. The LED array (Prizmatix Multi-LED system) provided independent intensity control (0-20 milliwatts per square millimeter) and temporal modulation (1 millisecond resolution) for each wavelength channel.

The fiber bundles were positioned using precision manipulators to place fiber tips at specific locations relative to organoid chambers, typically 100-300 micrometers from organoid surfaces to provide adequate illumination area while maintaining sufficient light intensity. The fibers were secured using UV-curable optical adhesive after positioning, creating permanent fixed geometry. The light delivery from 4,320 independent fibers enabled spatially precise optogenetic control where different regions of organoids or different organoid populations received independent optical stimulation.

The fifth (top) layer consisted of a transparent PDMS lid bonded to the device providing sealed enclosure while maintaining optical access for microscopy and biosensor imaging. The lid incorporated several critical features including integrated biosensor arrays, gas exchange membranes, and thermal control elements. The biosensor integration embedded fiber optic pH sensors (PreSens PSt7) and oxygen sensors (PreSens PSt3) at 340 locations throughout the device, with fibers terminating flush with the internal surface and connected via fiber bundles to external detector systems (PreSens OXY-10 and pH-10 multi-channel systems). The sensors provided continuous real-time monitoring of pH and oxygen at distributed locations.

The gas exchange membranes consisted of 50 micrometer thick PDMS regions (PDMS is highly permeable to oxygen and CO2) covering approximately 15% of device surface area distributed uniformly. These thin regions allowed bidirectional diffusion of respiratory gases between the internal environment and external atmosphere, enabling gas exchange supporting metabolism while preventing pressure buildup or vacuum formation. The external atmosphere was controlled at appropriate composition (typically 40% O2, 5% CO2, 55% N2) using mass flow controllers mixing gases from cylinders.

The thermal control elements consisted of serpentine resistive heaters (100 micrometer wide chromium traces deposited on PDMS surface) connected to temperature controllers with thermistor feedback. The heaters divided the device into 24 independently controllable thermal zones, each maintaining temperature with ±0.1 degree Celsius precision. This thermal control enabled both uniform maintenance at 37.0 degrees Celsius and creation of small temperature gradients if needed for guidance or activity modulation. Additionally, Peltier elements under the glass substrate provided cooling capacity for removing excess heat from high-activity regions.

The device assembly required precise alignment and bonding of these five layers in a specific sequence. The process began with bonding the microfluidic PDMS layer to the glass substrate with electrode array. Both surfaces were activated through oxygen plasma treatment (30 seconds, 50 watts, 200 millitorr oxygen) creating reactive silanol groups that form covalent Si-O-Si bonds when contacted. The PDMS layer was aligned to the glass substrate using registration marks visible under a stereomicroscope, brought into contact,

and the assembly was heated at 80 degrees Celsius for 2 hours to accelerate bonding. The bond strength exceeded 100 Newtons per square centimeter, sufficient to withstand perfusion pressures and handling.

The organoid chamber layer (acrylic) was then bonded on top of the microfluidic layer using medical-grade silicone adhesive (NuSil MED-1137) applied in thin uniform coating, aligned using registration marks, and cured at room temperature for 24 hours under light compression (10 kilopascals pressure from weights). The waveguide arrays were threaded through designated holes in the chamber layer during this bonding step, with final position adjustment performed after initial adhesive setting but before full cure, allowing precise fiber positioning.

The top PDMS lid was bonded last using plasma activation and thermal bonding as with the first layer. Prior to bonding, the integrated biosensor fibers were positioned through dedicated channels in the lid, and all fluidic and pneumatic connections were completed. The completed assembled device underwent comprehensive testing before organoid integration, including pressure testing all fluidic channels for leaks (pressurizing to 100 kilopascals and monitoring pressure decay, acceptance criterion <5% pressure drop over 1 hour), testing valve function across all 1,247 valves, validating optical transmission through all waveguides, and confirming electrical function of all 21,760 electrodes.

The perfusion system interfacing to the microfluidic device comprised multiple components providing medium supply, oxygenation, temperature control, flow control, and waste collection. The system employed a recirculating design where medium perfused through the device was collected, filtered, re-oxygenated, supplemented with nutrients, and returned to the device, enabling long-term operation without enormous medium consumption. Fresh medium was continuously added at 10-20 milliliters per hour to replace consumed nutrients and dilute accumulated waste, while equivalent volume was removed to waste, achieving complete medium turnover approximately every 30-40 hours.

The medium reservoir (5 liters) was maintained at 37 degrees Celsius using water jacket heating and continuously stirred to prevent settling and ensure homogeneous composition. The reservoir was sparged with controlled gas mixture (40% O2, 5% CO2, 55%

N2) through sintered glass diffuser, maintaining appropriate oxygen and carbon dioxide levels while providing mixing. The medium was then pumped through a 0.2 micrometer filter removing any particulates or contaminants, through a membrane oxygenator providing additional gas exchange, and through a heat exchanger maintaining precise temperature before entering the device.

The flow distribution to the device's 1,247 independently addressable channels employed a combination of passive hydraulic resistance networks and active flow control. The primary distribution used precision-machined manifolds creating hydraulic resistance ratios designed to achieve appropriate flow distribution, with high-resistance paths to organoid types requiring low perfusion and low-resistance paths to high-metabolism organoids. Superimposed on this passive distribution, 128 independently controlled pumps (miniature piezoelectric pumps from Bartels Mikrotechnik) enabled active modulation of flow to specific regions based on biosensor feedback.

The waste collection system consisted of parallel collection manifolds gathering effluent from different device regions into separate reservoirs for analysis. The regional separation enabled tracking metabolic activity and outputs from specific organoid types. The effluent collection reservoirs connected to automated sampling systems collecting 50 microliter aliquots every 30 minutes for immediate analysis or storage at -80 degrees Celsius for later analysis. The sampling system included inline analysis modules measuring glucose, lactate, oxygen, pH, and ammonia in real-time using electrochemical sensors, providing immediate feedback on metabolic state.

The control system coordinating all device functions comprised multiple interconnected components. A central controller (real-time Linux system running custom software) received inputs from biosensors, electrode arrays, and imaging systems, processed these through analysis algorithms, and generated control signals for pumps, valves, LEDs, and heaters. The control architecture employed both feedforward control based on predetermined protocols and feedback control responding to measured system state. For example, if oxygen sensors detected decreasing tissue oxygenation, the controller automatically increased perfusion rate or oxygen content, implementing homeostatic regulation maintaining physiological parameters within target ranges.

The completed device with integrated control systems underwent extended testing before organoid integration, running for 7-10 days with culture medium perfusion to identify any leaks, validate stable temperature and oxygen control, and confirm sterility maintenance. Only devices passing all validation criteria advanced to the organoid integration phase.

2.6 Vascular Network Engineering and Integration

The engineering of functional vascular networks achieving adequate perfusion throughout the large distributed tissue volume represented one of the most significant technical challenges and innovations of this work. The comprehensive strategy employed multiple complementary approaches spanning in vitro vessel generation, surgical integration, maturation protocols, and hemodynamic optimization.

The primary vascular network generation began with differentiation of endothelial progenitor cells from iPSCs through mesoderm induction followed by vascular specification. The protocol employed staged differentiation starting from iPSCs seeded at 20,000 cells per square centimeter in mTeSR1 on Matrigel-coated plates. On day 1, the medium was switched to mesoderm induction medium (RPMI 1640, B27 without insulin, BMP4 50 nanograms per milliliter, CHIR99021 6 micromolar) for 24 hours, driving primitive streak and mesoderm formation characterized by brachyury expression.

On day 2, the medium was switched to lateral mesoderm specification medium (RPMI 1640, B27 without insulin, BMP4 50 nanograms per milliliter, FGF2 25 nanograms per milliliter, VEGF 50 nanograms per milliliter) for 48 hours, promoting lateral mesoderm fate while initiating vascular specification. The FGF2 and VEGF signaling activated vascular gene programs including expression of FLK1 (VEGF receptor 2, also called KDR), the definitive marker of vascular progenitors. Flow cytometry on day 4 typically showed 60-80% FLK1-positive cells, indicating successful specification.

On day 5, the medium was switched to endothelial specification medium (StemPro-34 base, VEGF 50 nanograms per milliliter, FGF2 10 nanograms per milliliter) for 5-7 days, promoting maturation of vascular progenitors into endothelial cells. The developing endothelial cells were characterized by expression of CD31 (PECAM1),

VE-cadherin (CD144), and uptake of acetylated LDL (a functional characteristic of endothelial cells). By day 12, populations typically achieved >90% double-positive for CD31 and VE-cadherin, indicating mature endothelial phenotype.

The endothelial progenitor cells were expanded in endothelial growth medium (EGM-2, Lonza) through several passages, generating large quantities needed for vascular network construction. The cells were then mixed with supporting stromal cells (mesenchymal stem cells providing mural cell progenitors) at 4:1 ratio and embedded in fibrin-collagen hydrogel at density 5×10^6 cells per milliliter. The hydrogel consisted of 5 milligrams per milliliter fibrinogen (from human plasma) and 2 milligrams per milliliter collagen I (rat tail), with polymerization initiated by addition of thrombin (1 unit per milliliter) and incubation at 37 degrees Celsius for 30 minutes.

The hydrogel was supplemented with pro-angiogenic factors and ECM components promoting vascular network formation. The cocktail included VEGF-A (50 nanograms per milliliter) promoting endothelial sprouting, angiopoietin-1 (100 nanograms per milliliter) promoting vessel stabilization, stromal-derived factor 1 (SDF1, 100 nanograms per milliliter) recruiting supporting cells, FGF2 (20 nanograms per milliliter) promoting endothelial proliferation, and sphingosine-1-phosphate (S1P, 0.1 micromolar) promoting barrier function. Additionally, ECM proteins including laminin (10 micrograms per milliliter) and fibronectin (50 micrograms per milliliter) were incorporated to provide adhesive substrates guiding network formation.

The cell-laden hydrogel was cast into the organoid culture chambers in the microfluidic device, creating three-dimensional matrix filling the space between organoid positions. The embedded endothelial progenitors and supporting cells were cultured in endothelial growth medium for 7-10 days, during which they self-organized into primitive vascular networks through vasculogenesis. Time-lapse microscopy revealed the process: cells initially distributed randomly extended filopodia and migrated, made contact with neighbors, coalesced into cord-like structures, and underwent lumenization creating hollow tubes.

The primitive vascular networks showed appropriate topology with interconnected tubes spanning three-dimensional space, but lacked several properties of functional vasculature including hierarchical organization (all vessels similar diameter), inadequate mural cell coverage (unstable vessels prone to regression), and absence of connections to larger engineered vessels (no inflow/outflow). These limitations were addressed through multi-stage maturation and integration protocols.

The mural cell recruitment employed PDGF-BB gradients (10-100 nanograms per milliliter) applied through the microfluidic system, creating spatial patterns attracting mural cell progenitors (derived from the co-embedded mesenchymal stem cells) toward nascent vessels. The PDGF-BB was delivered at elevated concentrations near vessel locations and lower concentrations distant from vessels, creating gradients that PDGFR-beta expressing mural progenitors followed through chemotaxis. Upon reaching vessels, the progenitors differentiated into pericytes characterized by expression of NG2, CD146, and alpha-smooth muscle actin, and physically ensheathed the endothelial tubes.

The pericyte coverage dramatically stabilized vessels and prevented regression. Vessels with >50% circumference covered by pericytes showed <5% regression over subsequent weeks, while uncovered vessels showed >40% regression. The pericyte recruitment continued over 2-3 weeks until mature networks achieved 60-80% pericyte coverage, comparable to natural capillary networks. For larger vessels requiring smooth muscle coverage, we employed TGF-beta gradients attracting smooth muscle progenitors that formed organized layers around vessel walls, providing contractile capacity enabling active diameter regulation.

The hierarchical remodeling creating appropriate vessel size distributions employed flow-mediated mechanisms. The primitive networks initially showed narrow diameter distributions (most vessels 8-15 micrometers), inappropriate for efficient perfusion requiring hierarchical trees with large feeding arterioles, small capillaries, and large draining venules. The hierarchical organization emerged through flow-dependent remodeling processes where vessels carrying higher flow expanded while low-flow vessels regressed, a process recapitulating normal vascular development.

We initiated flow by connecting the engineered distribution vessels (silicone tubing 200-400 micrometers diameter) to the self-organized networks. The connection was achieved by positioning tubing ends within the hydrogel matrix near clusters of capillary-like vessels, allowing endothelial sprouts from the self-organized network to grow toward and connect with the engineered tubing walls over 5-7 days. The sprouting was guided by VEGF gradients (100-200 nanograms per milliliter) established by delivering VEGF-rich medium through the tubing. Once anastomoses formed, perfusion could be gradually initiated.

The flow initiation followed a carefully controlled escalation protocol to avoid damage from excessive shear stress while providing sufficient hemodynamic stimulus for remodeling. Starting at day 14 post-embedding (after primitive networks formed), we began perfusion at 0.1 microliters per minute per vessel, creating flow velocities approximately 100-fold lower than physiological. The flow rate was doubled every 2-3 days (following exponential ramp 0.1 → 0.2 → 0.4 → 0.8 → 1.6 microliters per minute) over 15 days, reaching target physiological flow rates by day 30.

During the flow ramp-up, extensive vascular remodeling occurred. Vessels on direct paths between arterial inputs and venous outputs experienced high flow and underwent expansion, with diameters increasing from initial 10-12 micrometers to 40-80 micrometers for arterioles and 60-100 micrometers for venules. Vessels in parallel paths with lower flow rates maintained smaller diameters or regressed entirely. The remodeling created hierarchical networks with appropriate diameter distributions spanning 5-150 micrometers, similar to natural microcirculation.

The arteriovenous specification, creating distinct arterial and venous identity rather than generic undifferentiated vessels, employed molecular patterning and hemodynamic signals. We applied gradients of arterial-specifying morphogens including Notch ligands (DLL4, 50 nanograms per milliliter) and ephrin-B2 (200 nanograms per milliliter) near arterial input regions, and venous-specifying factors including COUP-TFII agonists and EphB4 (200 nanograms per milliliter) near venous output regions. These molecular signals initiated arterial-venous specification, validated by expression of arterial markers (Hey2, Connexin40, ephrin-B2) in input vessels and venous markers (COUP-TFII, EphB4) in output vessels.

The hemodynamic differences reinforced molecular patterning. Arterial vessels experienced pulsatile flow with higher peak pressures and shear stresses, while venous vessels experienced steady lower-pressure flow. These hemodynamic patterns induced differential gene expression programs that further enhanced arteriovenous identity. The mature networks showed clear segregation of arterial and venous markers with minimal overlap, indicating successful specification.

The vascular integration with organoids required that vessels penetrate into organoid tissue rather than simply surrounding exteriors. Without internal vascularization, organoid cores more than 150-200 micrometers from surface would experience hypoxia. We promoted vascular invasion through several strategies. First, organoids were embedded in the same fibrin-collagen matrix used for vascular network generation, creating continuity between perivascular matrix and organoid-associated matrix, eliminating barriers to sprouting.

Second, we established VEGF gradients with elevated levels within organoid tissue by engineering some organoid cells (approximately 10-20% of population) to overexpress VEGF through stable transgene integration. The VEGF secretion created concentration gradients attracting endothelial sprouts from surrounding vessels toward and into organoids. Third, we incorporated pro-angiogenic ECM molecules including laminin and fibronectin into organoid embedding matrix, providing adhesive substrates supporting sprouting.

The vascular invasion occurred over 7-14 days following flow initiation, proceeding from organoid periphery toward core. Time-lapse microscopy with fluorescently labeled endothelial cells revealed sprouting processes extending from perivascular vessels, guided by VEGF gradients, navigating through organoid tissue. Upon reaching appropriate positions, the sprouts underwent lumenization creating perfused vessels within organoids. By 3-4 weeks post-flow initiation, most organoids showed extensive internal vascularization with vessel density approaching 250-400 vessels per square millimeter in highly metabolic neural and cardiac organoids.

The functional validation of vascular networks employed multiple complementary assays. Anatomical characterization through optical coherence tomography provided three-dimensional maps of complete vascular geometry, enabling quantification of vessel density (vessels per square millimeter), diameter distributions, hierarchical organization metrics, and network topology (connectivity, branch points, vessel length). Flow measurements through Doppler OCT quantified flow velocities and directions, confirming appropriate patterns with convergent arterial flow and divergent venous flow.

The perfusion adequacy was assessed through distributed oxygen tension measurements using fiber optic sensors and through tissue perfusion markers. We injected fluorescent microspheres (2 micrometers diameter) into arterial inflow and imaged their distribution, quantifying what fraction of tissue area received microspheres indicating adequate perfusion. The mature networks achieved >97% of tissue area perfused with microspheres, indicating comprehensive perfusion coverage. The oxygen measurements showed >95% of sensor locations maintaining PO2 above 20 millimeters of mercury during normal activity, confirming adequate oxygen delivery.

The barrier function, preventing excessive leakage of fluid and proteins from vessels into tissue, was validated through permeability assays. We perfused fluorescent dextrans of various molecular weights (10 kDa, 70 kDa, 150 kDa) and quantified their accumulation in tissue versus remaining in vessels. The mature networks showed appropriate size-selective permeability with permeability coefficients $<10^{-6}$ centimeters per second for 70 kDa dextran, approaching values for healthy microcirculation and indicating formation of appropriate endothelial tight junctions providing barrier function.

The vasoactive responses, demonstrating functional smooth muscle and endothelial control of vessel diameter, were tested through application of vasoactive compounds. Addition of acetylcholine (endothelium-dependent vasodilator acting through nitric oxide release) induced 20-35% diameter increases in arterioles while venules showed minimal response, confirming arterial identity and endothelial function. Addition of endothelin-1 (vasoconstrictor) induced 30-50% diameter decreases in arterioles. The magnitude and time course of responses approached those in native microcirculation, validating functional maturity.

The comprehensive vascular engineering strategy ultimately achieved remarkable functional properties enabling long-term maintenance of the large multi-organoid system. The total vascular network length of 47.3 meters distributed across 5400 square centimeters provided perfusion density of 8.8 millimeters of vessel per square centimeter, sufficient to maintain all tissue within diffusion distance of functional vasculature. The hierarchical organization with appropriate diameter distributions enabled efficient perfusion with reasonable pressure drops and flow distributions matching metabolic demands. The robust barrier function prevented excessive fluid filtration while allowing appropriate nutrient and waste exchange.

2.7 Organoid Integration and System Assembly

The assembly of functional integrated systems from the separately generated organoid types, engineered vascular networks, and microfluidic devices required precisely orchestrated protocols coordinating multiple simultaneous processes. The assembly occurred in stages over approximately 30 days, with critical activities including spatial positioning of organoids, vascular integration, establishment of inter-organoid connectivity, and functional validation before advancing to maturation phases.

The assembly timeline began on designated day 0 when all component organoids reached appropriate maturation states (typically day 50-60 of individual organoid culture for neural types, day 30-40 for cardiac, hepatic, and renal types). The organoids underwent comprehensive quality control assessments immediately before assembly, with detailed characterization including size measurements (calipers, acceptance windows ±15% of target size for each type), viability assays (flow cytometry with propidium iodide exclusion, requirement >90% viable), marker expression validation (immunostaining for lineage markers, requirement >85% appropriate markers), and functional assessment (electrophysiology for neural, contraction for cardiac, albumin secretion for hepatic, filtration for renal, with type-specific acceptance criteria).

Only organoids meeting all quality criteria advanced to assembly. The typical yield of organoids meeting criteria ranged from 68% to 91% depending on type, necessitating generation of approximately 1.4-fold excess organoids to ensure adequate numbers of

acceptable units. The quality control process occurred over days -2 to 0 immediately preceding assembly, minimizing time between assessment and use while enabling informed selection of highest-quality organoids.

The spatial positioning on assembly day 0 employed computer-controlled micromanipulators placing organoids into designated chambers in the microfluidic device according to the optimized spatial layout. Each organoid was transferred from its growth culture vessel, briefly rinsed in pre-warmed assembly medium (composition matching final perfusion medium), and positioned using a custom vacuum-assisted micromanipulator providing gentle pickup and precise placement. The manipulator system comprised a stereomicroscope with motorized XYZ stage (1 micrometer resolution positioning), vacuum pipette with force feedback preventing excessive suction damage, and computer control enabling automated execution of predefined positioning sequences.

The positioning sequence placed organoids in programmed order starting with deepest chambers and progressing to shallower ones, starting with highest-priority (most difficult to position) organoids, and grouping spatially clustered organoids to minimize manipulator movements. The complete positioning of 289 organoids required approximately 8-10 hours of continuous operation, maintained at 37 degrees Celsius in humidified 5% CO2 atmosphere throughout. Each positioned organoid was briefly inspected under magnification to confirm proper seating and orientation before proceeding to the next.

Once all organoids were positioned, the embedding proceeded by carefully filling chambers with fibrin-collagen hydrogel prepolymer containing vascular progenitor cells as described in section 2.6. The filling employed programmable syringe pumps injecting hydrogel through inlet ports at controlled rate (50 microliters per minute), slowly displacing air through outlet ports while filling around organoids. The filling sequence progressed from one side to the other, allowing gradual air displacement without trapping bubbles that would create regions inaccessible to perfusion.

The hydrogel polymerization was initiated by thrombin activation (final concentration 1 unit per milliliter thrombin) once all chambers were filled, with gelation proceeding over 30 minutes at 37 degrees Celsius. Following initial gelation, the entire device was transferred to

a humidified incubator and maintained under standard culture conditions (37 degrees Celsius, 5% CO2, 40% O2) while vascular network formation commenced. The assembly medium contained comprehensive nutrient support appropriate for all organoid types (essentially the final perfusion medium formulation described in section 2.5) plus elevated concentrations of pro-angiogenic factors (VEGF 100 nanograms per milliliter, angiopoietin-1 150 nanograms per milliliter, FGF2 30 nanograms per milliliter) promoting rapid vascular network formation.

The vascular network formation proceeded over days 1-14 following the protocols detailed in section 2.6, progressing through primitive network formation (days 1-7), mural cell recruitment (days 7-12), and connection to distribution vessels (days 10-14). The network development was monitored non-invasively through the transparent device using brightfield and fluorescence microscopy (endothelial cells expressed GFP enabling network visualization), with daily imaging documenting progressive network elaboration and confirming appropriate vascularization of all organoid regions.

The flow initiation on day 14 marked a critical transition from static hydrogel culture to perfusion culture. As detailed in section 2.6, the flow ramped gradually from 0.1 microliters per minute per vessel to full physiological rates over 15 days (days 14-29), enabling progressive vascular maturation while avoiding shear stress damage. Throughout the flow ramp, multiple parameters were continuously monitored including perfusion pressures (ensuring they remained below 200 millimeters of mercury to prevent vessel rupture), oxygen levels at distributed sensors (ensuring no hypoxia developed), pH (maintaining 7.35-7.45), glucose levels (maintaining 4-6 millimolar), and lactate (accepting initial elevations to 4-6 millimolar as tissue adapted to perfusion but intervening if exceeded).

The inter-organoid connectivity establishment through axon growth occurred in parallel with vascular maturation, spanning days 10-50. The process began around day 10 when axons from some organoid populations initiated extension. The axon guidance employed multiple mechanisms as detailed in section 1.4, including physical guidance through microfluidic conduits connecting source and target organoid positions, chemical guidance through neurotrophin gradients (NGF, BDNF, NT-3) established by differential

delivery through conduit inlets and outlets, and molecular addressing through complementary expression of engineered neuroligins and neurexins in source and target populations.

The guidance conduits (fabricated as part of the microfluidic device) consisted of 200-400 micrometer diameter channels connecting chambers, with inner surfaces coated with laminin and fibronectin providing adhesive substrates for axon growth. The conduits were filled with dilute collagen hydrogel (1 milligram per milliliter) providing three-dimensional scaffolding while permitting axon extension. The neurotrophin gradients were established by perfusing neurotrophin-rich solutions (100-200 nanograms per milliliter) through target-side conduit ports while maintaining lower concentrations (10-20 nanograms per milliliter) at source sides, creating gradients that growth cones followed through chemotactic responses.

The axon growth progressed at typical rates of 50-200 micrometers per day depending on axon type and local environment, requiring 7-30 days for connections spanning 0.5-4.5 centimeters. The growth was monitored through fluorescence imaging (source populations expressed fluorescent proteins enabling axon visualization), with time-lapse imaging of conduits documenting growth cone advance. The monitoring enabled early detection of guidance failures where axons stalled or grew in inappropriate directions, potentially enabling corrective interventions through modified neurotrophin gradients or optogenetic manipulations promoting growth.

The synaptic formation and maturation occurred over days 30-75 as arriving axons reached target organoids and formed synapses with appropriate postsynaptic partners. The synapse formation was mediated by molecular addressing through neuroligin-neurexin interactions as detailed in section 2.3, with arriving axons expressing specific neurexin variants preferentially forming synapses on target neurons expressing complementary neuroligin variants. The synaptic maturation involved progressive accumulation of synaptic proteins including presynaptic vesicle proteins (synapsin, synaptophysin), active zone proteins (bassoon, piccolo), postsynaptic scaffolds (PSD-95, gephyrin), and neurotransmitter receptors (AMPA, NMDA, GABA receptors).

The maturation was assessed through multiple approaches including immunostaining for synaptic markers showing progressive increases in puncta density, electron microscopy revealing ultrastructural maturation with characteristic synaptic morphology including vesicle accumulation, synaptic clefts, and postsynaptic densities, and electrophysiological recording demonstrating functional synaptic transmission. The synaptic maturation continued progressively through day 120, with synaptic density reaching asymptotic levels around 4.2×10^8 synapses estimated through stereological sampling and extrapolation.

The activity-dependent refinement of connectivity employed orchestrated optogenetic stimulation protocols during days 45-75 when initial connectivity was established but remained immature and partially non-specific. The refinement aimed to strengthen functionally appropriate connections while eliminating inappropriate connections through spike-timing-dependent plasticity mechanisms. The stimulation protocols delivered correlated activity patterns to populations that should be functionally connected, implementing Hebbian learning rules that strengthen connections when presynaptic activation precedes postsynaptic activation by appropriate delays (5-20 milliseconds).

The protocol design considered the intended functional connectivity matrix specifying which organoid populations should communicate strongly versus weakly or not at all. For strongly connected pairs, we delivered coordinated stimulation with source population activated followed by target population activation after delays matching expected synaptic transmission latencies (typically 10-15 milliseconds). The stimulation patterns repeated at 0.5-2 Hz for 20-30 minute sessions daily over the 30-day refinement period, totaling approximately 35,000-100,000 correlated activation events per connection pair.

For populations that should remain functionally independent despite potential anatomical contact, we delivered uncorrelated stimulation where source and target received activation patterns with randomized relative timing, preventing formation of stable temporal contingencies that would drive synaptic strengthening. The uncorrelated activation allowed testing whether molecular addressing and activity-dependent refinement

together achieved sufficient specificity, or whether some off-target connections persisted despite both mechanisms.

The refinement efficacy was assessed through paired recording experiments sampling hundreds of neuron pairs from different population combinations, measuring connection probability, synaptic strength, and latency distributions. The results showed substantial improvement over initial immature connectivity, with connection probability for molecularly matched pairs receiving correlated stimulation reaching 78.4 ± 6.7% compared to 52.3 ± 9.8% for matched pairs without refinement and 11.2 ± 3.4% for mismatched pairs, demonstrating that molecular addressing and activity-dependent refinement synergistically achieved high connectivity specificity.

The homeostatic regulation establishment enabling autonomous operation required developing and calibrating the feedback control systems connecting biosensors to effectors (pumps, valves, heaters). The regulation system comprised multiple nested control loops operating on different timescales. Fast loops (1-10 second time constants) controlled acute variables including perfusion pressure, temperature, and oxygen delivery, responding rapidly to perturbations. Intermediate loops (1-10 minute time constants) controlled metabolic variables including glucose delivery, pH buffering, and waste removal, adjusting based on trends rather than instantaneous values. Slow loops (1-24 hour time constants) controlled long-term processes including growth factor supplementation, medium turnover, and circadian regulation.

The control law development employed model predictive control approaches where system dynamics were modeled through empirical characterization experiments, control parameters were optimized through simulation, and implementations were validated through closed-loop testing. For example, the glucose control loop was developed by first characterizing glucose consumption rates under various activity conditions (measuring glucose depletion as a function of neural activity levels detected through electrode arrays), modeling the relationship between perfusion rate and glucose delivery, implementing a proportional-integral controller adjusting perfusion based on measured glucose and predicted consumption, and validating through perturbation experiments testing whether

the controller maintained glucose in target range (4.5-6.0 millimolar) despite varying demands.

The comprehensive control system implementation ultimately comprised 47 independent feedback loops regulating oxygen (12 zones with independent O2 control), glucose (8 zones with independent perfusion modulation), pH (6 zones with independent buffer delivery), temperature (24 zones with independent heating), and pressure (8 zones with independent pump control). The multi-loop architecture enabled local optimization where each region received appropriate support without excessive coupling to other regions that might have different requirements.

The functional validation before advancing to maturation and experimental phases required comprehensive characterization confirming that all major subsystems operated appropriately. The validation protocol spanned days 90-120 and included assessments of vascular function, neural activity, synaptic connectivity, metabolic state, and preliminary behavioral responses. Only systems passing all validation criteria advanced to the mature operational phase beginning day 120.

The vascular validation included perfusion imaging with fluorescent tracers documenting that all organoids received perfusion, oxygen mapping confirming adequate tissue oxygenation throughout, permeability assays validating appropriate barrier function, and hemodynamic measurements confirming appropriate pressure distributions and flow patterns. The neural validation included multi-electrode array recording confirming spontaneous activity throughout neural organoids with appropriate firing rate distributions, optogenetic testing validating light responsiveness in all targeted populations, and local field potential recordings confirming emergence of oscillatory activity characteristic of functional networks.

The connectivity validation included paired recording sampling from designated projection pairs confirming synaptic transmission with appropriate connection probabilities and strength distributions, anterograde and retrograde tracing confirming anatomical projections reached intended targets, and calcium imaging of population activity confirming coordinated activity between connected populations. The metabolic validation included

glucose consumption measurements confirming appropriate metabolism, lactate measurements confirming adequate oxidative capacity without excessive anaerobic metabolism, and ATP/ADP ratio measurements confirming adequate cellular energetics.

The preliminary behavioral assessment included simple tests of sensory responsiveness (optogenetic stimulation of sensory layers driving responses in higher layers), motor control (activation of motor organoids driving actuator outputs), and simple learning (repeated pairing of sensory patterns with rewards producing modest performance improvements). These preliminary tests aimed not to demonstrate sophisticated cognition (which required additional maturation) but rather to confirm basic sensorimotor loops functioned, that information propagated through hierarchies appropriately, and that plasticity mechanisms enabled learning.

Systems failing any validation component underwent troubleshooting attempting to identify and correct problems. Common failure modes included incomplete vascularization (addressed by local delivery of additional VEGF), regional hypoxia (addressed by increasing local perfusion or oxygen), lack of connectivity between intended partners (addressed by additional neurotrophin gradient application), excessive firing rates indicating hyperexcitability (addressed by enhancing inhibitory tone through pharmacological GABA potentiation), or insufficient firing rates indicating hypoexcitability (addressed by reducing inhibition or enhancing excitation). Approximately 73% of systems advanced directly through validation, while 18% required troubleshooting and re-validation (ultimately succeeding), and 9% exhibited irreparable problems necessitating termination and restart with new organoids.

The systems passing validation entered the maturation phase (days 120-180) during which functional capabilities progressively developed through experience-dependent mechanisms. During this period, we delivered structured stimulation protocols implementing "developmental experiences" analogous to sensory exposure during human development. The protocols included thousands of presentations of visual patterns with appropriate structure (edges, shapes, objects), motor sequences with consistent action-outcome contingencies, episodic sequences with repeating temporal patterns, and reward contingencies linking actions to outcomes.

These structured experiences drove progressive functional development measurable through improvements in multiple metrics. The sensory representations became more selective and invariant, with orientation tuning in visual cortex sharpening from initial broad tuning (tuning width 55 ± 12 degrees at day 120) to mature sharper tuning (32 ± 8 degrees at day 180). The motor control became more precise, with actuator targeting errors declining from 8.4 ± 2.3 millimeters at day 120 to 2.1 ± 0.6 millimeters at day 180. The memory capacity increased, with maximum number of distinct patterns recognizable increasing from 12 ± 4 at day 120 to 48 ± 9 at day 180.

By day 180, the systems reached a mature state where performance plateaued and stabilized, indicating readiness for formal cognitive testing. This 180-day timeline from iPSC starting material (day -60 organoid initiation, day 0 assembly, day 180 mature function) represented the complete system development cycle, substantially faster than human brain development requiring years but much longer than typical in vitro experiments measured in days to weeks.

2.8 Behavioral Training and Cognitive Assessment

The assessment of cognitive capabilities required developing behavioral paradigms adapted to the constraints and capabilities of our organoid system. Unlike conventional animal behavioral testing employing naturalistic behaviors like navigation or manipulation, our system interfaced with the environment primarily through optogenetic sensory input and microactuator motor output, necessitating abstract sensory-motor tasks. We developed a comprehensive battery of 24 tasks spanning seven cognitive domains, with each task designed to assess specific architectural components or cognitive functions while maintaining interpretability and minimizing confounds.

2.8.1 Sensory Processing and Perception Tasks

Pattern Recognition Task: This task assessed the hierarchical perceptual system's capacity for learning visual categories and generalizing across exemplars. The stimulus set comprised 48 distinct categories (26 letters A-Z, 10 digits 0-9, 12 geometric shapes) rendered as 32×32 pixel binary patterns. Each pattern was presented through optogenetic stimulation of the Layer 1 sensory organoids using the LED-waveguide system, with each

pixel mapped to specific illumination of corresponding retinal-layer neurons. The stimulation consisted of 200 millisecond presentations at 5 milliwatts per square millimeter intensity for "on" pixels and 0 milliwatts for "off" pixels.

The behavioral response consisted of actuator selection from an 8-actuator array positioned at accessible locations. Each actuator corresponded to 6 categories, arranged such that all 48 categories had unique actuator assignments. The system indicated its category judgment by activating one actuator (detected through force sensors requiring >50 millinewton activation threshold maintained for >100 milliseconds). Correct responses triggered reward delivery through brief activation of optogenetically stimulated dopaminergic neurons (470 nm, 10 millisecond pulse trains at 50 Hz for 200 milliseconds, producing phasic dopamine release), while incorrect responses produced no reward and triggered a 3-second timeout before the next trial.

The training protocol employed supervised learning with explicit feedback. Each of the 48 categories appeared 80 times across 3,840 trials (48 categories × 80 repetitions) delivered over 6 days (640 trials per day, approximately 3-4 hours of training given inter-trial intervals). The stimulus order was pseudorandomized with the constraint that no category appeared twice consecutively, preventing simple response perseveration strategies. The performance metric was accuracy (percent correct) calculated overall and separately for each category, with learning curves plotting accuracy versus trial number revealing acquisition dynamics.

The generalization testing employed novel exemplars of trained categories, created by applying transformations including rotation (±5, ±10, ±15 degrees), scaling (0.8×, 0.9×, 1.1×, 1.2×), and translation (±2, ±4 pixels in X and Y). The transformed stimuli were presented without feedback to assess whether the system learned abstract category representations that generalized across superficial variations versus merely memorizing specific training exemplars. The generalization performance (accuracy on novel transformed exemplars) provided critical assessment of the perceptual hierarchy's abstraction capability.

Temporal Sequence Detection Task: This task assessed the system's capacity for temporal integration and sequence recognition. The stimuli consisted of sequences of 5

events, with each event represented as a 200 millisecond optogenetic stimulation pulse to one of 8 sensory channels (different spatial locations in Layer 1). The inter-event interval was 200 milliseconds, creating 1.4 second total sequence duration. The task defined 8 target sequences with specific event orders (e.g., Target 1: channels 3-1-4-1-5, Target 2: channels 2-7-1-8-2, etc.) and 32 foil sequences differing from targets by reordering events or substituting events.

The behavioral response required indicating whether the presented sequence was a target (actuator 1 press) or foil (actuator 2 press). Correct target detections and correct foil rejections both counted as correct trials and earned rewards, while false alarms (calling foils targets) and misses (calling targets foils) counted as errors. The training protocol presented 1,000 trials (20% targets, 80% foils, matching natural base rates in many sequence detection scenarios) across 4 days, with performance assessed using signal detection theory metrics including hit rate, false alarm rate, and discriminability index d' quantifying sensitivity independent of response bias.

The temporal specificity testing employed timing-perturbed sequences where inter-event intervals varied from the trained 200 milliseconds to 150, 250, 300, or 350 milliseconds. The timing perturbation testing assessed whether sequence recognition depended on precise temporal structure versus merely event order, with performance declining for timing-perturbed sequences indicating temporal selectivity. This temporal selectivity would validate that the system learned not just which events occurred but when they occurred, implementing genuine temporal integration rather than merely event detection.

Multi-Modal Integration Task: This task assessed integration of information across sensory modalities, specifically vision and touch (simulated through different sensory organoid populations receiving independent input streams). The stimuli consisted of simultaneous presentations of visual patterns (32×32 optogenetic stimulation as in pattern recognition task) and tactile patterns (activation of 4 tactile sensor locations in 2×2 array), with 6 visual categories and 4 tactile categories creating 24 possible combined stimuli.

During training, specific visual-tactile combinations were consistently paired (e.g., visual pattern A1 always paired with tactile pattern T1, visual pattern A2 always paired with tactile pattern T2, creating 12 consistent combinations), while other combinations never occurred together (incongruent combinations). The behavioral task required indicating the overall stimulus category (12 possibilities corresponding to the 12 trained combinations) by selecting from a 12-actuator array. The training employed 960 trials (12 combinations × 80 repetitions) across 4 days with supervised learning and accuracy feedback.

The multi-modal integration testing employed three probe conditions revealing integration strategies. The unimodal probes presented only visual or only tactile components of trained combinations, testing whether performance relied on both modalities or one sufficed. The optimal integration theory predicts performance should decline more for removing the more informative modality, with exact predictions depending on individual modality reliabilities. The incongruent probes presented visual patterns from one trained combination with tactile patterns from another, creating conflicting information. The optimal cue integration theory predicts responses should reflect a weighted average of modality predictions, with weights proportional to modality reliabilities. The novel combination probes presented visual and tactile patterns that were individually familiar but never paired during training, testing whether the system learned combination-specific associations versus separate visual and tactile representations that could flexibly combine.

Motion Direction Discrimination: This task assessed temporal integration and motion processing by requiring discrimination of motion direction from sequential activation patterns. The stimuli consisted of sequential activation of 8 sensory locations arranged in a circle, with activation proceeding clockwise or counterclockwise at varying speeds (inter-location interval 50, 100, 150, or 200 milliseconds). The behavioral response indicated perceived direction (actuator 1 for clockwise, actuator 2 for counterclockwise), with correct responses earning rewards.

The training employed 800 trials (400 clockwise, 400 counterclockwise, balanced across speeds) over 3 days. The testing assessed performance separately at each speed, revealing temporal integration windows—if performance remained high across all speeds, the system integrated over long time windows, while performance declining for slower

speeds indicated limited integration. The psychometric analysis fit performance versus speed data to extract threshold speeds and temporal integration time constants, providing quantitative characterization of temporal processing capacity.

2.8.2 Learning and Memory Tasks

Classical Conditioning: This task assessed basic associative learning by pairing initially neutral conditioned stimuli (CS) with intrinsically valenced unconditioned stimuli (US). The CS consisted of distinct sensory patterns (8 different patterns), while the US consisted of either reward (dopaminergic neuron stimulation) or mild aversive stimulus (brief 0.5 second perfusion with elevated lactate concentration to 8 millimolar, creating mild metabolic stress). The training protocol employed delayed conditioning with CS onset preceding US onset by 500 milliseconds and CS remaining on until US offset.

The behavioral response involved preparatory motor output, with rewarded CS eventually eliciting approach responses (actuator extension toward stimulus location) and aversive CS eliciting defensive responses (actuator withdrawal from stimulus location). The response magnitude was quantified through force sensor readouts and compared between CS presentation periods before and after conditioning. The conditioned response (CR) magnitude, the difference in response during CS versus baseline, quantified learning. The acquisition curve plotted CR magnitude versus trial number, typically showing sigmoidal growth approaching asymptote after 40-80 trials for each CS-US pair.

The extinction testing presented CS alone without US, measuring how quickly conditioned responses declined. The extinction rate quantified flexibility of learning and ability to update predictions when contingencies changed. The discrimination testing presented CS+ (paired with US) and CS- (unpaired) in intermixed trials, measuring the system's capacity to differentiate stimuli and selectively respond to predictive cues. The discrimination index (CR magnitude for CS+ minus CR magnitude for CS-) quantified selectivity, with values near zero indicating poor discrimination and large positive values indicating strong selective responding.

Instrumental Conditioning: This task assessed learning of action-outcome contingencies where self-generated actions produce rewarding or aversive outcomes. The

setup provided 4 available actions (4 different actuator movements: extension, retraction, leftward, rightward) with probabilistic outcomes where each action had characteristic probability of producing reward. The contingencies were arranged to create distinct action values: Action 1 produced reward with 70% probability, Action 2 with 40%, Action 3 with 15%, and Action 4 with 5%.

The training consisted of free-choice trials where the system could select any available action, with subsequent outcome delivered according to the programmed contingencies. The performance metric was action distribution, specifically the proportion of trials where the system chose the optimal Action 1. Simple reinforcement learning models predict gradual increases in optimal action selection as experience accumulates, with learning rate and asymptotic performance depending on algorithm parameters. We fit computational models including Q-learning, actor-critic, and model-based planning to behavioral data, comparing model fits to infer underlying learning mechanisms.

The contingency reversal tested behavioral flexibility by reversing action-outcome mappings after initial learning (Action 1 becoming least rewarding, Action 4 becoming most rewarding). The reversal learning speed, measured as trials required to shift behavior toward the new optimal action, quantified flexibility and updating capacity. Slow reversal suggesting difficulty overcoming prior learning, while rapid reversal indicating flexible updating of action values based on current contingencies.

The outcome devaluation tested whether learning was goal-directed (encoding action-outcome-value contingencies) versus habitual (stimulus-response associations). Following initial training, we selectively devalued one outcome by pairing reward with lithium chloride-induced malaise (0.15 molar LiCl in perfusion medium for 30 minutes, creating conditioned taste aversion). If behavior was goal-directed, actions previously earning the now-devalued reward should decrease immediately without further training. If behavior was habitual, action frequencies should persist despite outcome devaluation. The immediate sensitivity to devaluation (change in devalued action frequency in first post-devaluation session) distinguished goal-directed from habitual control.

Spatial Learning and Navigation: This task assessed spatial memory and navigation capabilities in a virtual environment. The environment consisted of a 10×10 grid of locations, with the system's position indicated through activation of corresponding sensory channels creating a "place code." The system could navigate by selecting movement actions (up, down, left, right via actuators) that translated position in the virtual space. A goal location contained reward, with the task requiring learning to navigate from random start locations to the goal.

The training consisted of 200 trials starting from varied locations, allowing the system to explore and learn the spatial layout and goal location. The learning was assessed through measures including trial duration (time required to reach goal, decreasing with learning), path efficiency (actual path length divided by optimal path length, approaching 1.0 with learning), and search strategy (initially random search transitioning to direct goal-directed navigation). The place cell analysis examined whether hippocampal organoid neurons developed spatial selectivity, with some neurons preferentially active at specific virtual locations creating a cognitive map of space.

The probe trials removed the goal, testing whether navigation relied on beaconing toward local cues versus true spatial knowledge. If the system had learned a cognitive map, it should navigate to the goal location even without local cues present. The transfer testing changed the environment structure by altering wall locations or introducing shortcuts, testing whether the system maintained accurate spatial representations and could adapt to environmental changes. The navigation flexibility under these perturbations distinguished rigid stimulus-response mapping from flexible map-based navigation.

Sequence Learning: This task assessed learning of motor sequences and chunking of action patterns. The system was trained to execute specific sequences of 5-7 actions (e.g., extend-left-retract-right-extend-up-retract) with reward delivered only upon correct completion of entire sequences. The training employed shaping where initially only first 2 actions required correctness, then gradually increasing required sequence length to the full sequence. The training continued for 300-500 trials depending on sequence complexity.

The learning was quantified through multiple metrics. The execution time measured how rapidly sequences were performed, typically decreasing from 4-6 seconds early in training to 1.5-2.5 seconds in late training, indicating fluency development. The error patterns revealed chunking, with errors tending to occur at boundaries between chunks rather than within chunks, and chunks becoming increasingly stable units executed without interruption. The dual-task interference tested automaticity by requiring simultaneous performance of working memory task during sequence execution; interference indicated continued resource demands while lack of interference indicated automatized execution.

Recognition Memory: This task assessed the episodic memory system's capacity for encoding and retrieving specific experiences. The task employed a study-test paradigm with study phase presenting 40 distinct sensory patterns (each unique, not from trained categories), followed after variable delays (5 minutes, 1 hour, 24 hours, or 7 days) by test phase presenting the 40 studied patterns intermixed with 40 novel foil patterns. The behavioral response indicated whether each test item was old (studied) or new (foil) by pressing designated actuators.

The performance was analyzed using signal detection theory, calculating hit rate (proportion of studied items correctly identified as old), false alarm rate (proportion of foils incorrectly identified as old), and discriminability d' quantifying ability to distinguish studied from novel items. The d' metric corrected for response bias (some systems might respond "old" liberally versus conservatively), isolating true memory sensitivity. The retention curves plotted d' versus delay interval, revealing forgetting dynamics. Typical results showed exponential forgetting with d' declining from approximately 2.5 at 5 minute delay to 1.2 at 7-day delay.

The confidence judgment testing employed a two-stage procedure where after indicating old/new judgments, the system indicated confidence (high/low) via additional actuator press. The metacognitive sensitivity assessed whether confidence tracked accuracy, calculated as the difference in d' for high-confidence versus low-confidence responses (meta-d'). High metacognitive sensitivity indicated accurate self-monitoring of memory quality, while low sensitivity indicated dissociation between actual and perceived memory accuracy.

2.8.3 Planning and Problem-Solving Tasks

Two-Step Sequential Decision Task: This task distinguished model-based (planning-based) from model-free (cached value) decision strategies. The task structure included first-stage choice between two actions (A1 and A2) leading probabilistically to two second-stage states (S1 and S2), with A1 leading to S1 with 70% probability and S2 with 30%, while A2 showed opposite probabilities. At each second stage state, two actions were available (at S1: actions B1 and B2; at S2: actions C1 and C2), each producing reward probabilistically with slowly drifting reward probabilities implemented via Gaussian random walks (σ = 0.025 per trial).

Model-free learning predicts first-stage choices depend only on previous trial outcome—if the previous trial was rewarded, repeat the first-stage choice; if unrewarded, switch. Model-based planning predicts first-stage choices additionally depend on transition structure—if the previous trial was rewarded via a common transition (A1→S1 or A2→S2), strongly repeat; but if rewarded via rare transition (A1→S2 or A2→S1), weakly repeat since that transition is unlikely to recur. The key diagnostic is the interaction between reward on trial n and transition type on trial n when predicting first-stage choice on trial n+1.

The training consisted of 400 trials across 2 days, providing experience with transition structure and enabling learning of second-stage action values. The analysis employed mixed-effects logistic regression predicting first-stage choice as a function of previous reward, previous transition type, and their interaction, separately fitting each system's behavioral data. The interaction coefficient quantified model-based behavior, with positive values indicating model-based planning and near-zero values indicating purely model-free learning. Values typically ranged from 0.05-0.15 on a scale where 0.5 represents theoretical maximum.

Detour Task: This task assessed planning capacity by requiring navigation around obstacles to reach goals. The virtual environment (similar to spatial navigation task) initially allowed direct paths to goals, enabling learning of straight trajectories. After learning stabilized (typical 80 trials), an impenetrable barrier was introduced blocking the direct path, requiring multi-step detours. The optimal detour required 5-7 actions depending on start position and barrier configuration.

The planning signature appeared as a delay before initial action on barrier-introduction trials, interpreted as planning time. Systems using trial-and-error learning would begin random exploration immediately, while planning systems would pause to mentally simulate routes before initiating movement. The solution latency measured time from trial start to first action, typically increasing from approximately 0.8 seconds on familiar direct-path trials to 3-5 seconds on novel barrier trials for systems exhibiting planning. The neural signature of planning appeared as sequential reactivation of place-cell-like neurons representing locations along planned paths during pre-movement delays, detected through real-time decoding of neural population activity.

The transfer testing varied barrier configurations across trials, testing whether solutions generalized versus requiring relearning for each configuration. Rapid transfer indicated abstract understanding of detour principles, while slow transfer indicated configuration-specific learning. The path optimality assessed whether executed paths approximated optimal routes, with planning systems achieving >85% optimality (actual path length ÷ optimal path length < 1.15) versus <65% for trial-and-error systems.

Tower of Hanoi Analog: This task assessed hierarchical planning and problem decomposition capacity. The setup adapted the classic Tower of Hanoi puzzle to our actuator capabilities, using 3 disks of different sizes and 3 pegs, with goal of moving all disks from initial peg to target peg while respecting rules (only move one disk at time, never place larger disk on smaller). The 3-disk version requires minimum 7 moves for optimal solution.

The moves were implemented via actuator selections: 9 actuators corresponding to 3 disks × 3 pegs, with selected actuator indicating which disk to move to which peg. Invalid moves (violating rules) were prevented through mechanical constraints, but poor move choices lengthened solution paths. The training employed learning across 50 problems with various start and goal configurations, providing experience with puzzle structure. The performance metrics included solution rate (proportion of problems solved within 20-move limit), move efficiency (actual moves ÷ optimal 7 moves), and planning time (delay before first move, indicating pre-planning).

The subgoal analysis examined whether move sequences showed hierarchical structure characteristic of human problem solving, where the overall goal decomposes into subgoals (e.g., "free the smallest disk" → "create destination for medium disk" → "move medium disk" → etc.). The subgoal structure would appear as non-random pausing patterns, with longer pauses occurring at subgoal boundaries versus within subgoals, and as error patterns showing within-subgoal errors more common than across-subgoal errors. The planning depth was estimated from solution quality, with better performance indicating deeper lookahead.

Dynamic Replanning Task: This task assessed the system's capacity to adapt plans when circumstances changed unexpectedly. The scenario involved navigation toward a visible goal with known direct path, but with goal position occasionally changing mid-trial (10-30% of trials at random times). The appropriate response required detecting the goal change, abandoning the current plan, generating a new plan toward the new goal, and executing the new plan.

The performance metrics included replanning latency (time from goal change to movement direction change, indicating detection and replanning speed), replanning accuracy (whether new movement direction appropriate for new goal location), and completion rates (proportion of trials successfully reaching final goal despite perturbations). The neural signatures of replanning included activity shifts in planning organoids observable through population decoding, showing decoded planned trajectories switching from original to new goal within 0.5-2 seconds of goal changes.

The perturbation severity testing varied how dramatically goals moved, from small adjustments (requiring minor plan tweaks) to large displacements (requiring complete replanning). The graceful degradation of performance with perturbation severity versus catastrophic failure distinguished flexible adaptive planning from brittle plan execution unable to accommodate perturbations.

2.8.4 Metacognition and Confidence Tasks

Confidence Judgment Task: This task assessed the system's capacity to estimate confidence in its own perceptual judgments. The primary task was a difficult perceptual

discrimination (distinguishing two similar sensory patterns with varying amounts of noise added, creating trials with 50-95% inherent difficulty). Following each primary discrimination judgment, the system made a confidence judgment (high/low) via additional actuator press. Optionally, confident (high confidence) trials earned larger rewards if correct but larger penalties if incorrect, while unconfident (low confidence) trials earned/lost smaller amounts, creating incentive for accurate confidence assessment.

The metacognitive accuracy was quantified through type-2 ROC analysis comparing confidence against actual correctness. The area under type-2 ROC curve (AUROC2) provided a measure of metacognitive sensitivity, with values near 0.5 indicating chance (confidence uncorrelated with accuracy) and values near 1.0 indicating perfect metacognition. The meta-d' measure quantified metacognitive efficiency by comparing metacognitive sensitivity (how well confidence distinguished correct from incorrect trials) to task performance (how well discrimination judgments distinguished stimuli), with M-ratio = meta-d' ÷ d' indicating metacognitive efficiency independent of task difficulty.

The confidence calibration assessed whether reported confidence matched actual accuracy, plotting proportion correct versus confidence level. Perfect calibration appears as the identity line where proportion correct equals reported confidence. The calibration curve slope and intercept quantified whether the system was over-confident (reporting higher confidence than warranted), under-confident (reporting lower confidence), or well-calibrated. The response time analysis examined whether confidence related to response speed, with low-confidence trials typically showing longer deliberation times indicating uncertainty.

Opt-Out Paradigm: This task assessed the system's capacity to recognize when it lacked adequate information for reliable judgments, offering an option to decline difficult trials in exchange for small guaranteed reward. The task structure included difficult primary judgments (varying difficulty across trials from easy to nearly impossible), with an additional "opt-out" actuator that if pressed before the judgment actuator, terminated the trial with a small reward (one-third the reward for correct judgments) but no penalty. Strategic opt-out use required recognizing when expected value of attempting judgment (probability correct × reward - probability error × penalty) fell below the guaranteed opt-out reward.

The performance metrics included opt-out rate versus trial difficulty (should increase with difficulty), earnings maximization (total reward earned compared to optimal strategy), and signal detection measures treating opt-out decisions as confidence-based yes/no judgments. The optimal opt-out analysis calculated the difficulty threshold where expected value of attempting judgment equaled opt-out value, then compared actual opt-out rates above and below this threshold. Well-calibrated systems show sharp opt-out probability increases near the optimal threshold, while poorly calibrated systems show either excessive opt-out (conservative strategy) or insufficient opt-out (overconfident strategy).

Error Detection Task: This task assessed the system's capacity to recognize when it committed errors, offering opportunity to change responses after initial judgment but before feedback. The procedure involved initial judgment on discrimination task, followed by 2-second delay during which a "change response" actuator was available. Pressing this actuator allowed selecting a different response, with the final response used for reward determination. Strategic change-response usage required monitoring internal evidence indicating likely errors while resisting changing after likely correct trials.

The error detection sensitivity was quantified through signal detection analysis treating change-response decisions as discrimination between error and correct trials. The hit rate was proportion of errors where response changed (indicating error detection), false alarm rate was proportion of correct trials where response changed (false alarms indicating poor metacognition), and d' quantified error detection sensitivity. The response-time analysis examined whether change-response decisions showed appropriate timing, with shorter change latencies for clear errors versus longer latencies for marginal trials where error likelihood was ambiguous.

The neural correlates of error detection were examined through recordings from metacognitive organoids during the delay period, testing whether error-related neural activity differed between error and correct trials, and whether this neural activity predicted subsequent change-response decisions. The error-related negativity-like signal (sustained elevated activity 200-500 milliseconds post-response on error trials) would validate that metacognitive circuits detected errors through internal monitoring rather than external feedback.

2.9 Data Acquisition and Analysis

The comprehensive multi-scale measurements generating dozens of terabytes of data over the course of each experimental timeline required sophisticated acquisition systems and analysis pipelines for extracting interpretable information from raw signals. The following sections detail the major data modalities, acquisition parameters, processing approaches, and analysis methods employed.

2.9.1 Electrophysiological Recording and Analysis

The primary electrophysiological data stream consisted of continuous wide-band voltage recordings from 21,760 electrode sites sampled at 20 kHz with 16-bit resolution. The data acquisition employed a distributed amplifier architecture with 680 integrated circuit chips (Intan RHD2164, 32 channels per chip) positioned at the device periphery, with short low-impedance traces connecting electrodes to amplifier inputs minimizing noise pickup. The amplifiers performed analog signal conditioning including bandpass filtering (0.1 Hz to 10 kHz), programmable gain amplification (typical gain 200×), and analog-to-digital conversion before transmitting digital data streams via high-speed serial interfaces to data acquisition computers.

The data acquisition computers (12 workstations, each managing approximately 1,800 channels) received digital data streams, aggregated them with precise timestamp synchronization (jitter <50 microseconds across all channels), and streamed to RAID storage arrays (total capacity 480 terabytes, sufficient for approximately 180 days continuous recording from all channels). The sustained write bandwidth of 8.6 gigabytes per second required enterprise-grade storage infrastructure with multiple redundant controllers and distributed stripe-writing across 96 hard drives in RAID-6 configuration providing both performance and fault tolerance.

The spike detection and sorting analysis employed multiple processing stages transforming raw voltage traces into identified single-unit spike trains. The initial spike detection used threshold crossings where voltage exceeded 4.5 times the standard deviation of the high-pass filtered (300 Hz - 6 kHz) signal, a threshold providing good balance between sensitivity (detecting most spikes) and specificity (rejecting noise). The

threshold was calculated independently for each channel based on robust median absolute deviation estimator resistant to outliers, recalculated every 60 seconds to adapt to slow drift in noise characteristics.

When threshold crossings occurred, the algorithm extracted 2-millisecond waveform snippets (±1 millisecond surrounding crossing time) from the crossing channel plus 12 surrounding channels (hexagonal neighborhood providing full coverage of nearby electrodes). The multi-channel waveform features (520 samples: 40 time points × 13 channels) underwent dimensionality reduction through principal component analysis retaining 10 principal components capturing >95% of waveform variance while dramatically reducing data dimensionality enabling efficient clustering.

The spike sorting employed Gaussian mixture model clustering in the 10-dimensional feature space, automatically determining optimal cluster numbers through Bayesian information criterion minimization. The clustering algorithm (implemented in custom C++ code with MATLAB interfaces) processed overnight batches of one day's worth of detected events (typically 5-20 million events per day), requiring approximately 3-4 hours computational time on 32-core workstations. The resulting cluster assignments identified putative single units, with cluster quality metrics including isolation distance (measuring separation from other clusters) and signal-to-noise ratio determining which clusters represented well-isolated single neurons versus multi-unit activity or noise.

The clusters passing quality criteria (isolation distance >15, SNR >4.0, <1% refractory period violations indicating contamination) were accepted as identified single units, typically yielding 12,000-18,000 well-isolated neurons across the entire electrode array. The single-unit spike trains underwent further processing to characterize neural responses, firing statistics, and population dynamics. The peri-stimulus time histograms (PSTHs) aligned spike trains to task events (stimulus onsets, action executions, reward deliveries) and averaged across trials, revealing event-locked modulation of firing rates. The statistical significance of responses was assessed through comparison to baseline firing distributions using Wilcoxon signed-rank tests with multiple comparison correction.

The firing rate modulation analyses quantified how strongly neural activity related to task variables by regressing spike counts in analysis windows against continuous task variables (e.g., stimulus contrast, action force, reward magnitude) or discrete variables (e.g., stimulus category, action choice, outcome valence). The encoding models employed generalized linear models with Poisson or negative binomial distributions appropriate for count data, with model comparison via likelihood ratio tests determining which task variables significantly modulated firing. The population decoding analyses inverted the encoding question, asking how accurately task variables could be predicted from neural population activity using machine learning approaches including linear discriminant analysis, support vector machines, and recurrent neural networks.

The local field potential (LFP) analysis characterized slower network dynamics in the 1-200 Hz frequency range. The LFP signals were derived from the raw recordings by low-pass filtering at 300 Hz (using zero-phase Butterworth filter preventing temporal distortion) and downsampling to 1 kHz. The power spectral density analysis employed multitaper methods (using discrete prolate spheroidal sequences as tapers) providing optimal bias-variance tradeoff, revealing prominent oscillatory bands including delta (0.5-4 Hz), theta (4-8 Hz), alpha (8-13 Hz), beta (13-30 Hz), and gamma (30-100 Hz). The oscillation power changes time-locked to task events were quantified through event-related spectral perturbation analysis computing trial-averaged spectrograms aligned to events.

The phase-amplitude coupling analysis tested whether high-frequency oscillation amplitude modulated with low-frequency oscillation phase, a phenomenon implicated in long-range coordination and information routing. The analysis extracted low-frequency phase (e.g., theta phase) and high-frequency amplitude (e.g., gamma amplitude) time series through appropriate bandpass filtering and Hilbert transform, then computed modulation index quantifying how strongly gamma amplitude concentrated at particular theta phases versus uniform distribution. Significant coupling indicated coordinated dynamics spanning frequency bands, with the preferred phase revealing temporal organization of high-frequency processing relative to low-frequency rhythms.

The functional connectivity analysis characterized statistical dependencies between neurons' or regions' activities indicating communication or common inputs. The spike train

cross-correlation measured millisecond-timescale relationships between neuron pairs' firing, with sharp peaks at short latencies (2-5 milliseconds) suggesting monosynaptic connections and broader peaks at longer latencies suggesting polysynaptic or common input relationships. The transfer entropy analysis quantified directed information flow, assessing how much knowing past activity of neuron A improved prediction of future activity of neuron B beyond what could be predicted from B's own history alone. The graph theory analyses represented connectivity patterns as networks, computing metrics including clustering coefficient (tendency for connected neurons to share other connections), path length (typical synaptic distance between neuron pairs), modularity (segregation into communities), and hub identification (neurons with exceptional connectivity).

2.9.2 Optical Imaging and Analysis

The optical imaging data streams included wide-field fluorescence imaging capturing biosensor signals across large tissue areas, and two-photon imaging providing cellular-resolution measurements in smaller regions. The wide-field imaging employed four sCMOS cameras (2560×2160 pixels, 95% quantum efficiency, 100 frames per second maximum rate) viewing the device from above through the transparent top layer, with each camera covering approximately one-quarter of the device area (15×22.5 cm field of view per camera). The illumination employed LED arrays with appropriate wavelengths for exciting fluorescent biosensors (470 nm for GFP-based sensors, 565 nm for RFP-based sensors, plus additional wavelengths for other fluorophores), with intensity modulation creating structured illumination patterns and enabling spectral unmixing when multiple fluorophores required simultaneous imaging.

The typical imaging protocol involved continuous acquisition at 20-50 Hz frame rate, generating approximately 12 gigabytes per hour per camera (48 GB/hour for four-camera system), sustainable for extended recording periods given the RAID storage infrastructure. The image processing pipeline converted raw image sequences into quantitative biosensor measurements through several stages. The preprocessing included flat-field correction compensating for uneven illumination using calibration images acquired with uniform fluorescence, registration correcting for small mechanical vibrations or thermal drift through

subpixel image alignment based on fiducial markers, and background subtraction removing out-of-focus fluorescence and autofluorescence using morphological opening operations.

The biosensor signal extraction defined regions of interest (ROIs) corresponding to individual organoids or subregions, extracted mean fluorescence intensity within each ROI, and converted fluorescence to biosensor-reported quantities (calcium concentration, voltage, neurotransmitter concentration, etc.) using calibration curves measured during sensor characterization. For ratiometric sensors (e.g., calcium sensors using FRET), the analysis computed fluorescence ratios between donor and acceptor channels, providing measurements less susceptible to motion artifacts and photobleaching than single-wavelength intensity. The temporal filtering removed high-frequency noise while preserving biologically relevant dynamics, typically using low-pass Gaussian filters with sigma = 50-100 milliseconds for calcium signals and sigma = 10-25 milliseconds for voltage signals.

The calcium imaging analysis identified calcium transients corresponding to neural activity through several approaches. The simple threshold detection identified transients exceeding baseline plus 2-3 standard deviations, appropriate for clear signals with good signal-to-noise. The template matching approach cross-correlated traces with canonical calcium transient templates (fast rise, exponential decay), detecting events with appropriate kinetics even in noisy traces. The deconvolution approach inverted the calcium indicator response function to estimate underlying spike trains, accounting for multi-exponential calcium binding kinetics and indicator saturation, yielding spike rate estimates comparable to electrophysiological recordings (when both modalities were available, correlation between deconvolved calcium and spike rates typically $R^2 = 0.7\text{-}0.85$).

The population calcium imaging analysis characterized coordinated activity across neuron populations through dimensionality reduction and clustering approaches. The principal component analysis identified dominant patterns of co-variation, with leading principal components often revealing functionally relevant population modes (e.g., first PC corresponding to overall arousal, second PC distinguishing different task conditions). The independent component analysis separated mixed signals into statistically independent source signals, useful for identifying distinct neuron ensembles with decorrelated activity

patterns. The non-negative matrix factorization decomposed population activity into additive components with non-negative weights, interpretable as distinct cell assemblies activating with varying strengths across time.

The two-photon imaging provided cellular-resolution calcium imaging within local tissue volumes, enabling measurement of individual neuron responses and fine-scale neural circuitry characterization. The two-photon microscope (custom-built on Thorlabs cage system) employed a mode-locked Ti:Sapphire laser (Coherent Chameleon Ultra II, 920 nm center wavelength, 140 femtosecond pulse width, 80 MHz repetition rate) for excitation, galvanometric mirrors for beam scanning, and photomultiplier tubes for detection. The scanning operated in resonant mode at 30 Hz frame rate for 512×512 pixel fields of view, or in galvo mode at 3-5 Hz for larger 1024×1024 pixel imaging, with the slower rate compensated by better signal-to-noise from longer pixel dwell times.

The typical two-photon imaging session targeted specific organoids or inter-organoid regions for detailed investigation, acquiring 10-30 minute recordings during behavioral tasks. The analysis employed motion correction algorithms (based on rigid and non-rigid image registration) compensating for tissue movement, cell segmentation algorithms (based on constrained non-negative matrix factorization) automatically identifying individual cell ROIs, and calcium transient detection extracting spike-inferred neural activity from each identified cell. The two-photon data complemented wide-field imaging by providing cellular-resolution validation and detailed circuit analysis in regions of interest identified through wide-field surveys.

2.9.3 Metabolic and Physiological Measurements

The continuous monitoring of metabolic and physiological parameters employed diverse sensor technologies integrated into the perfusion and device systems. The glucose and lactate measurements used enzymatic electrochemical sensors (immobilized glucose oxidase and lactate oxidase enzymes on electrode surfaces) positioned in perfusion inflow and outflow lines, measuring concentrations every 30 seconds with 0.1 millimolar resolution. The consumption rates calculated as the product of concentration difference (inflow minus outflow) and flow rate, providing real-time metabolic activity readouts with

approximately 2-minute temporal resolution limited by flow transit times through the device.

The oxygen measurements employed fiber-optic sensors (PreSens PSt3 oxygen-sensitive phosphorescent dyes) distributed at 340 locations throughout the device, with sensor tips positioned flush with tissue surfaces. The measurement principle relied on phosphorescence quenching by oxygen, with phosphorescence lifetime decreasing with increasing oxygen concentration. The detection system excited sensors with modulated red LED light (620 nm, 1 kHz modulation) and measured phase shift of the phosphorescent emission (760 nm), with phase shift calibrated to oxygen partial pressure. The system cycled through all 340 sensors sequentially at 1 Hz (completing full scan every 340 seconds), providing spatial oxygen maps revealing regional variation and enabling early detection of hypoxia.

The pH measurements similarly employed fiber-optic sensors (PreSens PSt7 pH-sensitive fluorescent dyes) at 340 locations, measuring via ratiometric fluorescence where pH-sensitive and pH-insensitive fluorophores excited simultaneously (505 nm, 640 nm) and emission ratio (670 nm/770 nm) reported pH with 0.02 pH unit resolution. The ammonia measurements used ion-selective electrodes (gas-permeable membrane allowing ammonia diffusion into internal electrolyte solution with pH electrode) at 12 strategic locations, providing 0.1 micromolar sensitivity sufficient for detecting elevated ammonia levels indicating hepatic dysfunction or protein catabolism.

The hormone and growth factor measurements employed automated sampling and immunoassay analysis. The perfusion sampling system collected 50 microliter aliquots from 8 circuit locations every 30 minutes, dispensing into 384-well plates stored at 4 degrees Celsius until analysis (typically within 24 hours to prevent degradation). The immunoassays employed Luminex multiplexed bead-based technology enabling simultaneous measurement of 40+ analytes per sample, including cytokines (IL-6, TNF-α, IL-1β), growth factors (VEGF, FGF2, BDNF, NGF), hormones (insulin, cortisol, thyroid hormones), and neurotransmitters (dopamine, serotonin, norepinephrine, acetylcholine). The analysis employed automated liquid handling (Hamilton STAR platform) and plate readers (Luminex

FLEXMAP 3D) processing 384-well plates in approximately 4 hours, yielding concentration measurements for all analytes across all samples.

The metabolomics analysis characterized broader metabolic profiles through liquid chromatography-mass spectrometry (LC-MS). The samples (100 microliter aliquots from tissue extracts or perfusion medium) underwent protein precipitation (3:1 methanol:sample), centrifugation, and injection onto UPLC system (Waters Acquity) coupled to high-resolution mass spectrometer (Thermo Q Exactive). The chromatography employed hydrophilic interaction chromatography (HILIC) for polar metabolites and reverse-phase chromatography for lipids, with gradient elution over 20-minute runs. The mass spectrometry operated in both positive and negative ionization modes with full-scan acquisition (70,000 resolution, 70-1050 m/z range) and data-dependent MS/MS for metabolite identification.

The metabolomics data analysis employed peak detection identifying chromatographic peaks (using XCMS software), mass-to-charge ratio matching against metabolite databases (HMDB, KEGG, METLIN) for putative identification, and statistical analysis comparing metabolite levels across conditions or time points. The pathway analysis mapped detected metabolites to metabolic pathways using KEGG pathway annotations, identifying coordinated changes across pathways indicating altered metabolic programs (e.g., glycolysis, TCA cycle, amino acid metabolism, lipid metabolism). The integration with transcriptomics data (when available from tissue biopsies) enabled connecting metabolic changes to underlying gene expression programs.

The vascular imaging and hemodynamics employed optical coherence tomography (OCT) and Doppler OCT providing non-invasive three-dimensional mapping of vessel architecture and flow. The OCT system (Thorlabs Telesto II) generated cross-sectional images by measuring backscattered near-infrared light (1300 nm) with 5.5 micrometer axial resolution and 13 micrometer lateral resolution, sufficient for resolving individual vessels down to 10-15 micrometer diameter. The volumetric imaging acquired stacks of cross-sectional images scanning across device regions, with complete device imaging requiring approximately 20 minutes at sufficient resolution for vessel detection.

The vessel segmentation employed machine learning algorithms (3D U-Net convolutional neural networks trained on manually annotated ground truth) automatically identifying vascular structures in OCT volumes, generating binary vascular masks indicating vessel presence at each voxel. The vascular network analysis processed segmented masks to extract quantitative metrics including vessel density (total vessel length per tissue volume), diameter distributions (histogram of vessel diameters), hierarchical organization (branching patterns, diameter relationships), and topological properties (connectivity, branching points, loops versus trees). The temporal comparison of vessel networks acquired at different time points revealed remodeling dynamics, with registration algorithms aligning time points and change detection identifying vessel regression, growth, and diameter changes.

The Doppler OCT measured blood flow velocities through phase-shift analysis, where moving blood cells caused Doppler frequency shifts in backscattered light. The velocity measurements combined with vessel diameter measurements yielded flow rate calculations (flow = velocity × cross-sectional area), enabling quantitative hemodynamic characterization. The flow mapping revealed flow distributions across vascular networks, identified high-flow paths corresponding to major arterioles and venules, and enabled calculation of perfusion per tissue volume (total flow per unit mass, typical units milliliters per minute per 100 grams).

2.9.4 Behavioral Data Analysis

The behavioral data analysis transformed raw sensor readouts (actuator positions, forces, timing) into interpretable behavioral metrics and model-based inferences about cognitive processes. The preprocessing standardized data formats across different sensor types, synchronized behavioral events with neural recordings using shared timestamps, and annotated trials with metadata including task conditions, stimulus parameters, and outcomes. The trial segmentation divided continuous recordings into individual trials based on programmed trial onset signals or detected behavioral events, with trial boundaries refined through automated and manual verification.

The performance quantification computed accuracy (proportion correct trials), reaction times (latency from stimulus onset or go signal to response initiation), movement

parameters (trajectory smoothness, velocity profiles, force profiles), and task-specific metrics (response bias, sensitivity, choice consistency). The learning curves plotted performance metrics versus trial number or training day, revealing acquisition dynamics with typical sigmoidal shapes characterized by initial low performance, accelerating improvement, and asymptotic plateau. The curve fitting employed logistic or exponential functions parameterizing learning rate (steepness of improvement phase) and asymptote (final performance level), enabling quantitative comparison of learning across conditions or systems.

The choice analysis characterized action selection patterns through multiple approaches. The choice probability analysis computed frequency of each available action as function of task variables, revealing deterministic or probabilistic choice policies. The logistic regression models predicted choices from trial-by-trial variations in stimulus features, recent history, and outcomes, identifying significant predictors and quantifying their influence through regression coefficients. The reinforcement learning models simulated learning processes, with model parameters (learning rates, exploration rates, discount factors, temperature parameters) fit to behavioral data through maximum likelihood estimation, and model comparison through likelihood ratio tests or information criteria determining which learning mechanisms best explained observed behavior.

The sequential dependencies analysis examined how previous trials influenced current choices, indicative of learning and working memory processes. The simple win-stay-lose-shift analysis computed probability of repeating previous action contingent on previous outcome (stayed after reward, switched after no reward), with deviations from 50% random responding indicating outcome sensitivity. The more sophisticated mixed-effects logistic regression models predicted current choice from multiple previous trials' outcomes and choices, revealing longer-term dependencies and distinguishing immediate reinforcement from slower strategic adjustments.

The response time analysis provided insights into decision processes complementing choice analysis. The distributional analysis characterized full response time distributions through ex-Gaussian fits (convolution of Gaussian and exponential distributions) capturing both central tendency and skewed tail structure. The response time versus accuracy

relationship (speed-accuracy tradeoff) was characterized through quantile-probability plots, revealing whether fast responses traded accuracy for speed. The drift-diffusion modeling explained response time distributions and accuracy jointly through evidence accumulation to threshold mechanisms, with model fitting estimating drift rate (evidence quality), threshold (speed-accuracy criterion), and non-decision time (sensory-motor delays).

The model-based analysis formalized computational hypotheses about cognitive processes as mathematical models and fit them to behavioral data to quantify underlying cognitive parameters. The reinforcement learning models characterized learning through parameters including learning rates (α controlling weight given to new information versus prior beliefs), discount factors (γ determining preference for immediate versus delayed rewards), exploration rates (β or temperature parameters controlling randomness in action selection), and eligibility traces (λ parameters determining credit assignment over temporally extended action sequences).

The model fitting employed multiple approaches depending on model complexity. Simple models with closed-form likelihood functions used maximum likelihood estimation computing parameter values maximizing probability of observed data given model. More complex models without analytical likelihoods employed simulation-based methods including Markov Chain Monte Carlo (MCMC) sampling posterior distributions over parameters, Sequential Monte Carlo approximating likelihoods through particle filtering, or Approximate Bayesian Computation comparing simulated and observed data distributions. The model validation employed cross-validation dividing data into training and test sets, fitting parameters on training data and evaluating predictions on held-out test data, ensuring models generalized beyond training data.

The model comparison employed multiple criteria assessing goodness-of-fit while penalizing complexity to prevent overfitting. The likelihood ratio tests compared nested models (where one model is special case of another with parameters constrained) through chi-square tests of likelihood differences. The information criteria including Akaike Information Criterion (AIC) and Bayesian Information Criterion (BIC) balanced log-likelihood against parameter count, with lower values indicating better models. The Bayesian model comparison computed Bayes factors (ratio of marginal likelihoods) or posterior model

probabilities, providing evidence strength for one model versus alternatives while automatically accounting for complexity through marginalization over parameters.

The integrated analysis combined behavioral, neural, and physiological data streams revealing relationships between levels. The behavior-neural correlations tested whether neural activity patterns predicted behavioral choices or performance, computed through regression of choice variables on neural population activity or through decoding algorithms predicting behavior from neural states. The neural-metabolic relationships examined whether metabolic activity (glucose consumption, oxygen consumption) correlated with neural activity levels, validating that metabolic support matched computational demands. The perturbation analysis manipulated neural activity (through optogenetics) or metabolic state (through glucose modulation) while measuring effects on behavior and other levels, establishing causal relationships beyond correlation.

3. Results

3.1 System Development and Characterization

The successful generation, assembly, and maturation of functional multi-organoid Psi systems represented a substantial technical achievement spanning six to seven months from initial iPSC differentiation through mature autonomous operation. We generated three independent systems (designated System A, B, and C) following identical protocols to assess reproducibility and characterize the range of emergent properties across independently developed instances. All three systems successfully completed assembly and maturation phases, achieved comprehensive quality control criteria, and proceeded to functional characterization and cognitive assessment over observation periods extending 180-290 days post-assembly.

3.1.1 Organoid Generation and Quality Control

The organoid generation phase (days -60 to 0 relative to assembly) achieved high success rates across most organoid types with acceptable reproducibility. The neural organoid production from six independent iPSC lines showed average viability of 93.2 ± 3.8% (measured at day 50-60 pre-assembly), with size distributions clustering tightly around target dimensions (coefficient of variation 12-18% depending on organoid type, well within acceptable ±25% specification). The marker expression validation showed that 89.4 ± 4.6% of neural organoids expressed appropriate regional markers (layer markers for cortical organoids, hippocampal markers for hippocampal organoids, etc.) at levels exceeding thresholds, with only 10.6% rejected for inappropriate marker patterns suggesting mispatterning during regional specification.

The functional validation of neural organoids through electrophysiological characterization at day 50-60 showed that 87.2 ± 5.3% exhibited spontaneous activity with appropriate characteristics including mean firing rates of 2.4 ± 1.2 Hz (log-normal distribution), network bursting at 0.1-0.4 Hz (characteristic of developing neural networks with recurrent excitation), and appropriate responses to depolarizing current injection showing excitability without hyperexcitability. The optogenetic validation in organoids from

engineered iPSC lines confirmed that 91.8 ± 4.2% showed appropriate light responses with >90% of targeted neurons responding to illumination at specified wavelengths and intensities.

The cardiac organoid generation achieved spontaneous contraction in 94.7 ± 3.1% of organoids by day 12-14, with synchronous coordinated contraction across organoid tissue observable by day 18-20 in 89.3 ± 5.4% of spontaneously beating organoids. The contractile function quantification at day 30-40 (after electrical and mechanical conditioning) showed beat rates of 68-94 beats per minute (mean 76 ± 11 BPM across organoids), force generation of 1.8-3.2 millinewtons per square millimeter (mean 2.4 ± 0.6 mN/mm^2, approaching lower end of adult myocardium range), and appropriate action potential morphology with rapid upstroke velocities (210 ± 45 V/s) and plateau phases (230-280 millisecond duration).

The hepatic organoid validation at day 30-40 pre-assembly confirmed albumin secretion (8.4 ± 3.2 micrograms per day per organoid, sufficient for maintaining oncotic pressure when 23 organoids integrated into system), CYP450 enzyme activity (metabolizing midazolam with clearance rates of 165 ± 48 picomoles per minute per million cells, approximately 45% of primary human hepatocyte rates), and appropriate morphology with hepatocyte rosettes surrounding bile canaliculi-like structures (confirmed by immunostaining for MRP2 at apical membranes and ZO-1 at tight junctions). The glycogen storage capacity measured through PAS staining showed 91% of organoids accumulated significant glycogen (qualitatively moderate to strong PAS reactivity), with quantitative biochemical assays on representative organoids indicating 1.9 ± 0.6 milligrams glycogen per organoid.

The renal organoid validation showed glomerular structures (identifiable by histology as cellular tufts surrounded by Bowman's-like capsules) in 84.2 ± 7.1% of organoids, with an average of 52 ± 18 glomeruli per organoid. The tubular segment differentiation assessed through marker expression showed proximal tubules (LTL+ cubilin+), loops of Henle (NKCC2+), and distal tubules (NCC+) in appropriate proportions (approximately 50% proximal, 25% loop, 25% distal tubules). The functional validation through fluorescent dextran filtration assays showed size-selective barrier function in 78 ± 9% of organoids with

glomerular structures, with 10 kDa dextran cleared from vascular compartment into tubular compartment while 150 kDa dextran retained in vascular space.

The comprehensive quality control resulted in selection rates (proportion of generated organoids meeting all criteria for inclusion in assembled systems) of 76-91% across organoid types, with the lowest rates for technically challenging specialized types (cerebellar organoids 76%, hippocampal organoids 79%) and highest rates for robust abundant types (dorsal cortical organoids 91%, cardiac organoids 89%). The high selection rates minimized waste while ensuring only high-quality organoids advanced to expensive and time-consuming assembly phases.

3.1.2 System Assembly and Vascular Integration

The assembly timeline spanning days 0-30 proceeded smoothly across all three systems with only minor complications requiring intervention. The organoid positioning on day 0 required 9.2 ± 0.8 hours per system (Systems A, B, C: 8.7, 9.5, 9.4 hours respectively), within expected range given 289 organoids requiring individual manipulation. The positioning accuracy assessed through post-positioning imaging showed average deviation from intended positions of 124 ± 67 micrometers (mean across all organoids in all systems), well within the ±500 micrometer tolerance given that most organoids were 4-8 millimeters diameter, indicating successful submillimeter-precision positioning.

The hydrogel embedding and polymerization proceeded without significant bubble trapping (only 3 small bubbles <0.5 millimeter diameter detected in System B, deemed acceptable as they were distant from organoid tissue and could be perfused around). The initial static culture phase (days 1-13) showed good organoid survival with only 2.4% organoid loss (7 organoids total across 3 systems, equaling 867 total organoids positioned, with losses distributed across organoid types without systematic pattern suggesting technical failure modes). The lost organoids were identified through lack of visible cellular structure in brightfield imaging and confirmed through lack of fluorescence in organoids expressing fluorescent proteins, likely representing organoids that were compromised during transfer and positioning rather than healthy organoids lost during culture.

The vascular network formation during days 1-14 was monitored through daily fluorescence imaging of GFP-expressing endothelial cells. Time-lapse analysis revealed network morphogenesis kinetics: day 1-3 showed minimal organization with randomly distributed cells, day 4-6 showed cord formation with cells coalescing into elongated structures, day 7-9 showed lumenization with hollow tube formation detectable through perfusion of fluorescent dextran showing intravascular distribution, and day 10-14 showed network elaboration and connection to distribution vessels. By day 14, the primitive vascular networks achieved 87 ± 9% coverage (proportion of tissue area within 200 micrometers of detectable vessels), approaching target densities required for perfusion initiation.

The vessel-organoid integration assessed through histological sampling (3 organoids per type per system sacrificed at day 21 for fixation and sectioning) showed vascular penetration into organoid tissue with vessel density within organoids reaching 62 ± 18% of perivascular density (vessels per square millimeter in organoid interior versus immediately surrounding organoid), indicating substantial but incomplete vascularization at this early time point. The temporal progression showed continued vascular ingrowth over subsequent weeks, with vessel density within organoids reaching 83 ± 12% of perivascular density by day 45 and 94 ± 8% by day 75, indicating progressive maturation approaching uniform vascular perfusion throughout.

The flow initiation and ramp-up proceeded according to protocol across all systems without requiring deviation from planned escalation schedule. The perfusion pressures monitored continuously remained within safe ranges throughout (peak pressures 145-178 mmHg, safely below 200 mmHg threshold where vessel rupture risk increases). The oxygen monitoring showed initially heterogeneous tissue oxygenation with some regions showing marginally low PO2 (18-22 mmHg, just below target >20 mmHg) during early low-flow phases (days 14-20), but improving progressively as flow increased, with >95% of sensors reading >20 mmHg by day 30 when full flow rates achieved.

The hierarchical vascular remodeling was documented through repeated OCT imaging at days 14, 21, 28, and 42. The vessel diameter distributions showed dramatic changes over this period, progressing from narrow initial distributions (median 11 micrometers, interquartile range 9-14 micrometers at day 14) to broad mature distributions

spanning two orders of magnitude (median 24 micrometers, range 6-156 micrometers at day 42). The diameter distribution at day 42 closely approximated power-law scaling with exponent -2.8 ± 0.3 (consistent with optimal transport networks and natural microcirculation showing exponents -2.5 to -3.0), indicating successful emergence of hierarchical organization through flow-mediated remodeling.

The arterio-venous specification assessed through immunostaining for arterial markers (ephrin-B2, Connexin40) and venous markers (EphB4, COUP-TFII) at days 21 and 42 showed progressive segregation. At day 21, only 62 ± 11% of vessels showed clear arterial or venous identity (expressing arterial markers at >3× background or venous markers at >3× background), with 38% expressing neither or both marker sets indicating immature mixed phenotype. By day 42, 89 ± 7% showed clear identity with minimal overlap (only 4 ± 2% expressing both marker sets above threshold), indicating successful maturation of distinct arterial and venous vessel types with appropriate regional patterns (arterial vessels near inflow, venous near outflow, capillary networks in intermediate regions expressing neither marker set).

The functional vascular assessment at day 42-50 demonstrated key properties indicative of mature functional vasculature. The perfusion coverage measured through fluorescent microsphere distribution showed 97.3 ± 1.8% of tissue area receiving microspheres (averaged across 3 systems with multiple imaging regions per system), confirming near-complete perfusion. The barrier function assessed through permeability to fluorescent dextrans of varying sizes showed appropriate size-selective permeability with permeability coefficient for 70 kDa dextran of $(6.2 \pm 2.4) \times 10^{-7}$ centimeters per second, comparable to healthy microvascular permeability and indicating formation of appropriate endothelial tight junctions. The vasoactive responses tested through acetylcholine addition (10 micromolar in perfusion medium for 5 minutes) showed 28 ± 8% diameter increases in arterioles with return toward baseline after washout, confirming functional endothelium capable of nitric-oxide-mediated vasodilation.

3.1.3 Connectivity Establishment and Refinement

The inter-organoid axonal connectivity establishment progressed over days 10-75 with major projections reaching target regions by day 40-50 and synaptic maturation

continuing through day 75-90. The axon growth monitoring through fluorescence imaging of labeled projections (select source populations expressing fluorescent proteins) showed growth front advancement at rates varying across connection types: short connections (<5 millimeters) typically completed growth within 10-15 days (0.3-0.5 millimeters per day), medium connections (5-20 millimeters) required 20-40 days (0.3-0.6 millimeters per day), and longest connections (20-45 millimeters) required 40-75 days (0.5-0.7 millimeters per day), with longer connections showing somewhat faster growth potentially reflecting better-established guidance gradients in longer conduits.

The projection targeting accuracy assessed through anterograde viral tracing at day 60 (injecting AAV expressing fluorescent proteins into source organoids and imaging target regions 14 days later for labeled terminals) showed that 82.4 ± 6.7% of labeled projections reached intended target organoids rather than alternative targets. This targeting accuracy varied systematically across connection types, with reciprocal connections between hierarchically adjacent processing layers showing highest accuracy (88-94%, as these connections traversed relatively short distances through single guidance conduits with strong neurotrophin gradients), feedback connections spanning multiple hierarchical levels showing intermediate accuracy (78-85%, requiring navigation through longer paths with potentially competing guidance cues), and connections between functionally distinct systems showing lower accuracy (71-82%, as these connections lacked developmental precedents that could be recapitulated through morphogen patterning).

The synaptic connectivity assessed through paired electrophysiological recordings at days 60, 75, and 90 revealed progressive maturation of functional connections. At day 60, connection probability for neuron pairs from molecularly matched populations (expressing complementary neuroligin-neurexin pairs) averaged 52.3 ± 9.8%, substantially above the 8.2 ± 3.4% baseline for molecularly mismatched pairs, but below the 78.4 ± 6.7% ultimately achieved by day 90. The synaptic strength for detected connections also increased over this period, with mean evoked postsynaptic potential amplitude increasing from 0.31 ± 0.18 millivolts at day 60 to 0.54 ± 0.28 millivolts at day 90 (measured for connections between cortical layer 4 and layer 2/3 neurons, a well-characterized projection).

The activity-dependent refinement protocols applied during days 45-75 demonstrably improved connectivity specificity. In System A where full refinement protocols were implemented, final connection probability for matched pairs reached 78.4%, while in a control organoid pair cultured separately without refinement protocols, connection probability plateaued at 58.7%, demonstrating a 34% improvement attributable to activity-dependent refinement. The refinement appeared to operate through two mechanisms: strengthening appropriate connections (mean synaptic strength in refined connections 1.6× stronger than unrefined) and eliminating inappropriate connections (connection probability for mismatched pairs declined from 12.1% at day 60 to 6.8% at day 90 in refined systems versus remaining stable at 11.5% in unrefined controls).

The axon tract organization examined through tissue clearing and light-sheet microscopy at day 90 revealed remarkably organized fiber bundles connecting organoids. The major tracts contained 8.3×10^5 to 3.2×10^6 axons (estimated through automated fiber counting in high-resolution image stacks), traveled relatively direct paths between source and target (path length typically 1.1-1.3× straight-line distance), and showed appropriate regional organization with sensory-to-motor projections positioned dorsally while memory-to-planning projections positioned medially, recapitulating gross organizational principles of mammalian white matter tract topology.

The total connectivity estimated through stereological sampling and computational extrapolation yielded approximately 2.3×10^8 synapses across the entire system at day 90 maturity, distributed appropriately across system regions. The sensory processing hierarchy contained approximately 6.2×10^7 synapses concentrated in layer 4 which receives thalamic-like input, the motor hierarchy contained 4.8×10^7 synapses with highest density in premotor regions implementing action sequence encoding, the memory systems contained 5.4×10^7 synapses with particularly high density in hippocampal CA3 supporting autoassociative dynamics, and the remaining 6.9×10^7 synapses distributed across motivational, planning, and metacognitive systems.

The synaptic density calculated as synapses per cubic millimeter of neural tissue ranged from 1.2×10^8 to 4.8×10^8 depending on region, with highest densities in hippocampal and layer 4 regions containing many small densely packed neurons with extensive local

connectivity, and lower densities in motor regions containing larger neurons with sparser connectivity. These densities approached lower bounds of mammalian cortical synaptic densities (typically 4×10^8 to 10^9 synapses per cubic millimeter), consistent with our organoids containing less tightly packed tissue than in vivo cortex but achieving synapse densities within the same order of magnitude.

3.1.4 Physiological Homeostasis and Metabolic Function

The physiological homeostasis established during the maturation phase (days 30-120) achieved stable autonomous regulation of critical parameters by day 120 across all three systems. The metabolic stability assessed through continuous monitoring (days 120-180) showed glucose concentrations maintained at 5.2 ± 0.6 millimolar in arterial perfusion (range across time points and systems: 4.3-6.1 millimolar, mean ± SD calculated across 25,920 measurements: 3 systems × 180 days × 48 measurements per day), with venous glucose at 4.1 ± 0.5 millimolar, indicating consumption of approximately 1.1 millimolar or 21% of delivered glucose. The circadian variation in glucose levels (oscillating approximately ±8% around mean) developed spontaneously despite constant glucose delivery, reflecting endogenous circadian rhythms in metabolic activity discussed in section 3.8.1.

The lactate concentrations showed baseline levels of 1.8 ± 0.4 millimolar rising to 3.2 ± 0.7 millimolar during periods of high cognitive activity (defined as periods with elevated neural firing rates detected through electrode arrays). The lactate/glucose ratio of approximately 0.35 during baseline and 0.78 during peak activity indicated predominantly oxidative metabolism (as complete glycolysis of consumed glucose would yield lactate/glucose ratio of 2.0, while pure oxidative metabolism yields ratio of 0.0), with approximately 60-70% of consumed glucose proceeding through oxidative pathways and 30-40% through glycolysis. This metabolic profile resembled awake mammalian brain showing similar oxidative/glycolytic balance and contrasted sharply with pure anaerobic metabolism characteristic of hypoxic tissue.

The oxygen delivery and consumption showed arterial PO_2 maintained at 285 ± 12 millimeters of mercury (through membrane oxygenator control), venous PO_2 at 58 ± 14 millimeters of mercury (indicating substantial but not excessive oxygen extraction), and

tissue PO_2 measurements from distributed sensors averaging 32 ± 8 millimeters of mercury (range 15-67 millimeters of mercury across 340 sensors × 3 systems × multiple time points). Critically, 97.1% of sensor readings exceeded 20 millimeters of mercury (the threshold below which cellular respiration becomes oxygen-limited), with only 2.9% showing marginal values of 15-20 millimeters of mercury (adequate to prevent cell death but potentially limiting peak metabolic activity). The few sensors showing values below 15 millimeters of mercury (0.2% of readings) triggered automated interventions increasing local perfusion or oxygen concentration, successfully restoring adequate oxygenation within 10-30 minutes in all cases.

The pH stability showed arterial pH of 7.38 ± 0.03 and venous pH of 7.34 ± 0.04 (measured continuously through fiber optic sensors), remaining within the narrow physiological range (7.35-7.45) required for optimal enzyme function and neural excitability. The pH regulatory mechanisms included the bicarbonate buffering system in perfusion medium (25 millimolar sodium bicarbonate equilibrated with 5% CO_2), supplementary HEPES buffer (20 millimolar providing non-volatile buffering capacity), and active regulation by hepatic organoids performing acid-base homeostasis through bicarbonate generation and ammonia detoxification. The occasional pH deviations (12 instances across 3 systems over 180 days where pH briefly dropped to 7.28-7.32 during periods of exceptionally high metabolic activity generating excess CO_2) were corrected within 15-30 minutes through automated increase in bicarbonate delivery and CO_2 removal.

The ammonia concentrations measured through ion-selective electrodes showed baseline levels of 28 ± 12 micromolar rising to 65 ± 24 micromolar after extended periods of high protein catabolism (during energy deprivation experiments described in section 3.8.4), but remaining well below the neurotoxicity threshold (>150-200 micromolar where significant neural dysfunction emerges). The hepatic detoxification capacity estimated from ammonia clearance kinetics following ammonia challenges (bolus addition of 100 micromolar ammonium chloride to perfusion medium followed by monitoring clearance) showed half-life of 23 ± 8 minutes, indicating robust hepatic urea cycle function sufficient to prevent toxic accumulation during normal operation.

The hepatic function assessed through multiple biomarkers confirmed adequate metabolic support capacity. The albumin secretion measured through ELISA on perfusion medium samples showed concentrations of 38 ± 4 grams per liter (similar to plasma albumin levels), with synthesis rate of approximately 3.2 milligrams per day across the 23 hepatic organoids, sufficient to maintain oncotic pressure and protein transport capacity. The cytochrome P450 enzyme activity assessed through metabolism of probe substrates (midazolam for CYP3A4, dextromethorphan for CYP2D6) showed clearance rates averaging 47% of primary human hepatocyte rates (range 38-58% across systems and time points), sufficient for metabolizing endogenous compounds and providing some drug metabolism capacity.

The glycogen storage capacity assessed through periodic PAS staining of tissue biopsies showed substantial glycogen accumulation in hepatic organoids (2.4 ± 0.6 milligrams per organoid, totaling approximately 55 milligrams across 23 organoids), providing approximately 4-6 hours worth of glucose supply at normal consumption rates (0.25-0.35 milligrams glucose per minute). This glycogen store was dynamically regulated, depleting during extended periods without glucose supplementation and replenishing when glucose availability recovered, demonstrating functional glycogen synthesis and mobilization appropriate for buffering glucose fluctuations.

The renal function characterized through multiple metrics confirmed adequate fluid and electrolyte homeostasis. The glomerular filtration rate analog (volume of perfusion medium cleared per minute) measured through inulin clearance assays showed 218 ± 24 microliters per minute across the 18 renal organoids (approximately 12 microliters per minute per organoid), sufficient to filter the entire perfusion volume (approximately 400 milliliters) every 30-35 hours. The tubular reabsorption function assessed through glucose appearance in collected tubular fluid (extracted from collection channels) showed >99.5% glucose reabsorption efficiency, preventing wasteful glucose loss while enabling appropriate glucose homeostasis.

The electrolyte regulation maintained sodium at 138-142 millimolar (physiological range), potassium at 4.2-4.8 millimolar (slightly elevated compared to typical 3.5-5.0 millimolar plasma levels but within safe range), calcium at 2.2-2.6 millimolar (appropriate

ionized calcium range), and phosphate at 0.8-1.4 millimolar (normal range). The electrolyte stability indicated appropriate tubular transport function with regulated reabsorption and secretion maintaining homeostasis despite variable intake from perfusion medium and variable production from cellular metabolism.

The cardiac function assessed through force transducers attached to selected cardiac organoids and through ECG-like electrical recordings from organoid surfaces showed heart rate variability around baseline of 72 ± 8 beats per minute, with rate modulation ranging from 58 beats per minute during low-activity quiescent periods to 105 beats per minute during high-activity engaged periods. This physiologically appropriate rate variability reflected both intrinsic pacemaker properties and autonomic-like modulation from engineered sympathetic (noradrenergic) and parasympathetic (cholinergic) innervation. The force generation measured at 2.4 ± 0.6 millinewtons per square millimeter represented substantial contractile capacity, with the distributed arrangement of 67 cardiac organoids providing local perfusion enhancement throughout the system.

The integrated physiological function achieving stable homeostasis across multiple regulated variables over 180+ day periods without requiring manual intervention (beyond routine medium supplementation and waste removal) validated the successful implementation of autonomous homeostatic regulation. The systems demonstrated genuine physiological autonomy analogous to living organisms maintaining internal milieu despite external fluctuations, a critical requirement for the autonomous agency specified in Bach's Psi architecture.

3.1.5 Neural Activity Patterns and Network Dynamics

The characterization of spontaneous and task-evoked neural activity provided critical validation that engineered neural circuits exhibited appropriate functional properties. The spontaneous activity measured through multi-electrode arrays during rest periods (defined as 5+ minute intervals without sensory stimulation or motor output) revealed 14,347 ± 2,132 well-isolated single units across the three systems (Systems A, B, C: 13,982, 14,556, 14,503 units respectively), representing successful recording from approximately 12% of the estimated 1.2×10^8 neurons based on geometric considerations of electrode reach and neuron density.

The firing rate distribution across the recorded population showed characteristic log-normal shape with median 2.8 Hz, mean 4.2 Hz, and range 0.1-45 Hz (95th percentile: 12.4 Hz). This distribution closely resembled those reported in mammalian cortex both in shape (log-normal) and parameters (median 2-5 Hz typical across multiple cortical recording studies), validating that engineered networks achieved realistic activity levels avoiding hyperexcitability (which would produce pathological high-frequency firing) or hypoexcitability (which would produce near-silent networks).

The cell-type classification based on spike waveform features (width, asymmetry) and firing statistics (mean rate, burstiness, coefficient of variation) identified putative excitatory neurons (78.4% of units, showing broad spike waveforms >0.4 milliseconds, mean rates 2.6 ± 1.8 Hz) and putative inhibitory interneurons (21.6% of units, showing narrow spike waveforms <0.4 milliseconds, mean rates 8.7 ± 4.3 Hz). This approximate 80:20 excitatory:inhibitory ratio matched the typical cortical composition, though the classification remains putative without direct verification of neurotransmitter phenotype.

The spontaneous network dynamics revealed striking organized patterns despite absence of external input. The dominant pattern consisted of synchronized bursting where large populations of neurons fired together in events lasting 100-500 milliseconds separated by quiescent intervals of 2-8 seconds, yielding burst rates of 0.08-0.25 Hz. The burst participation (proportion of neurons showing elevated firing during bursts) averaged 43 ± 12%, indicating nearly half the population coordinated their activity. The burst amplitude (firing rate during burst normalized by baseline rate) averaged 4.7 ± 2.1, indicating substantial coordination with firing rates increasing approximately fivefold during synchronized events.

The spatial organization of bursts examined through cross-correlation analysis and population vector trajectories showed that bursts typically initiated in specific organoids (different initiation sites across different bursts, but with some organoids initiating disproportionately often) and propagated across the system with characteristic velocities of 15-40 centimeters per second (calculated from spike time differences across electrode positions). The propagation patterns often followed known anatomical connections, with activity flowing through sensory hierarchies from low to high levels or through memory

systems from hippocampal to cortical regions, suggesting that spontaneous dynamics reflected and potentially refined functional connectivity.

The network state transitions between synchronized bursting and desynchronized irregular activity occurred spontaneously approximately every 15-45 minutes, creating alternation between high-synchrony and low-synchrony modes. The synchronized states showed elevated low-frequency power (0.5-8 Hz) in local field potentials, elevated spike-field coupling (neurons firing at preferred LFP phases), and elevated long-range correlation (activity in distant organoids more correlated). The desynchronized states showed suppressed low-frequency power, reduced spike-field coupling, elevated high-frequency gamma power (30-80 Hz), and reduced long-range correlation but elevated local correlation, resembling transitions between sleep-like and wake-like states in mammalian brain.

The oscillatory dynamics examined through spectral analysis of local field potentials revealed multiple frequency bands with distinct functional properties. The delta band (0.5-4 Hz) showed highest power during synchronized states with power averaging 1.2 ± 0.4 millivolts² during synchronized epochs versus 0.3 ± 0.2 millivolts² during desynchronized epochs. The theta band (4-8 Hz) showed prominent oscillations in hippocampal organoids during active task engagement, with power elevated 3.2 ± 0.8-fold during spatial navigation versus rest. The alpha band (8-13 Hz) showed moderate power throughout the system with spatial variation (highest in posterior sensory regions, lowest in anterior planning regions) and attention-related modulation (decreased during sensory processing periods).

The beta band (13-30 Hz) showed transient bursts associated with motor planning and execution, with power increasing 2.4 ± 0.6-fold during the 500 milliseconds preceding movement onset compared to baseline, then declining rapidly after movement initiation. The gamma band (30-100 Hz) showed strongest power during active sensory processing with power in sensory organoids elevated 2.8 ± 0.9-fold during stimulus presentation compared to inter-stimulus intervals. The high-gamma band (100-200 Hz) showed brief transient bursts associated with specific computational events, particularly salient in hippocampal organoids during memory encoding and retrieval where sharp-wave ripple complexes appeared as stereotyped 100-200 Hz oscillations lasting 50-150 milliseconds.

The cross-frequency coupling analysis revealed phase-amplitude coupling where gamma amplitude modulated with theta phase (in hippocampal regions) and with alpha phase (in sensory regions), quantified through modulation indices of 0.024 ± 0.008 (hippocampus theta-gamma) and 0.018 ± 0.006 (sensory alpha-gamma). These values indicated moderate coupling strength sufficient for functional coordination but below the very strong coupling (MI > 0.05) sometimes observed in pathological states. The preferred phases showed gamma amplitude maximal at theta trough and at alpha ascending slope, consistent with gamma bursts occurring when inhibition wanes, enabling enhanced local processing.

The sensory-evoked activity tested through systematic optogenetic stimulation mapping (stimulating each of 8 sensory channels individually at varying intensities while recording throughout the system) revealed appropriate propagation through sensory hierarchies. The Layer 1 sensory input neurons responded with high reliability (94.2 ± 3.7% of trials showing >5 Hz firing rate elevation for >50 milliseconds following stimulus onset) and short latency (8.3 ± 2.1 milliseconds from light onset to first spike). The responses propagated to Layer 2 with slight delay (12.4 ± 3.2 milliseconds total latency, indicating 4.1 milliseconds synaptic and propagation delay), to Layer 3 with further delay (19.7 ± 4.8 milliseconds), and ultimately to Layer 6 with accumulated delays (48.3 ± 12.4 milliseconds).

The hierarchical processing revealed through response selectivity showed that Layer 1 neurons responded promiscuously to most stimuli (average selectivity index 0.12 ± 0.08, where 0 = responds equally to all stimuli and 1 = responds only to single stimulus), while Layer 2-3 neurons showed moderate selectivity (0.34 ± 0.12), and Layer 5-6 neurons showed high selectivity (0.61 ± 0.18), indicating progressive emergence of specific feature tuning through hierarchical processing. The receptive field sizes estimated from spatial stimulation patterns increased systematically (Layer 1: 0.8 ± 0.3 millimeters of tissue responding, Layer 2-3: 2.4 ± 0.8 millimeters, Layer 5-6: 7.2 ± 2.4 millimeters), indicating convergence creating increasingly abstract representations spanning larger spatial scales.

The motor activity examined through recordings in motor organoids during actuator control showed appropriate encoding of motor parameters. The motor cortex-like organoid neurons showed firing rates correlated with actuator force output (correlation coefficients R

= 0.76 ± 0.14 for most strongly tuned neurons), with population decoding enabling prediction of force from neural activity with R^2 = 0.68 ± 0.11 accuracy. The motor activity preceded movement onset by 180 ± 60 milliseconds on average, indicating genuine preparatory motor planning rather than purely sensory feedback-driven responses. The directionally selective neurons (responding more strongly to movements in particular directions) represented all directions with approximately uniform distribution, enabling flexible encoding of arbitrary motor outputs.

The memory-related activity in hippocampal organoids showed place-cell-like spatial selectivity during virtual navigation tasks. Among recorded hippocampal neurons, 23.7 ± 6.4% showed significant spatial tuning (firing rates varying >3-fold across different virtual locations, ANOVA $p < 0.01$ with multiple comparison correction). The place fields (regions of elevated firing) averaged 2.1 ± 0.8 virtual grid positions in size (corresponding to approximately 20% of the 10×10 grid environment), similar to mammalian place cell field sizes relative to environment dimensions. The spatial information content averaged 0.47 ± 0.18 bits per spike for spatially tuned cells, indicating that each spike reduced uncertainty about location by approximately 1.4-fold, comparable to rodent place cells (typically 0.3-0.7 bits per spike).

The comprehensive neural characterization spanning spontaneous dynamics, oscillations, sensory responses, motor encoding, and memory representations validated that engineered neural circuits exhibited functional properties qualitatively and quantitatively resembling biological neural networks. The activity patterns indicated appropriate excitation-inhibition balance, functional hierarchical organization, and specialized circuit dynamics supporting cognitive functions specified in the Psi architecture.

3.2 Perceptual Hierarchy Function

The hierarchical perceptual system's capacity for progressive abstraction and invariant representation was assessed through multiple complementary approaches including behavioral performance on pattern recognition tasks, neural decoding at different hierarchical levels, representational similarity analysis, and perturbation experiments.

3.2.1 Pattern Recognition Learning and Generalization

The pattern recognition task training (48 categories × 80 repetitions = 3,840 trials over 6 days) showed robust learning across all three systems. The learning curves plotting accuracy versus trial number exhibited characteristic sigmoidal shapes with initial near-chance performance (day 1: 8.4 ± 3.2% accuracy compared to 2.1% chance for 48-class problem), accelerating improvement (days 2-4: increasing to 42.7 ± 8.4%, 68.3 ± 6.2%, 81.2 ± 5.8%), and asymptotic plateau (days 5-6: 85.6 ± 4.8%, 87.3 ± 4.2%). The logistic function fits to learning curves (accuracy = L / (1 + exp(-k(trial - t_0))) where L is asymptote, k is learning rate, t_0 is inflection point) yielded parameters L = 88.2 ± 3.1%, k = 0.0024 ± 0.0006 per trial (indicating approximately 400-trial half-time to asymptote), and t_0 = 1847 ± 312 trials (inflection occurring mid-training as expected).

The final test performance after training (assessed on 960 trials with 20 repetitions per category) showed 87.3 ± 4.2% overall accuracy, with relatively uniform performance across categories (category-specific accuracy range 78-94%, mean 87.1%, SD 4.1%, coefficient of variation 4.7%), indicating balanced learning without strong category-specific biases. The confusion matrix analysis revealed error patterns clustered among perceptually similar categories (e.g., letter "O" confused with digit "0", letter "I" confused with digit "1"), suggesting errors reflected genuine perceptual similarity rather than random noise or motor execution failures.

The generalization testing with transformed stimuli (rotated, scaled, translated versions of trained exemplars, 20 novel exemplars per category = 960 test trials total) showed significant but incomplete generalization with accuracy of 72.4 ± 6.8%, substantially above chance (2.1%) and demonstrating genuine learned representations rather than memorized training exemplars, but below training-set performance (87.3%) indicating incomplete invariance. The generalization performance varied systematically with transformation magnitude, declining approximately linearly with transformation distance: ±5 degree rotation showed 84.2 ± 5.3% accuracy (minimal decline), ±10 degrees showed 76.8 ± 6.1%, ±15 degrees showed 68.4 ± 7.2%, indicating approximately 1% accuracy decline per degree of transformation.

The comparison between systems revealed modest individual differences in learning and generalization performance. System A showed fastest learning (reached 75% accuracy

by trial 1,423) and highest asymptotic performance (89.7%), System C showed slowest learning (75% at trial 2,134) and lowest asymptotic performance (84.8%), with System B intermediate (75% at trial 1,782, asymptote 87.4%). These differences, while statistically significant ($p < 0.01$, repeated-measures ANOVA), remained modest in magnitude (5.4% range in asymptotic performance), suggesting that the architectural implementation was sufficiently robust that developmental variations created only minor performance differences.

3.2.2 Neural Encoding Across Hierarchical Levels

The neural population decoding analysis tested how well category identity could be predicted from neural activity at different hierarchical levels. Linear discriminant analysis decoders trained on neural population firing rates (calculated in 200-millisecond windows following stimulus onset) and tested via 5-fold cross-validation yielded decoding accuracy (percent correct category prediction) increasing across hierarchical levels: Layer 1 showed 34.2 ± 5.8% accuracy, Layer 2 showed 58.7 ± 7.2%, Layer 3 showed 71.4 ± 6.8%, Layer 4 showed 83.9 ± 5.4%, Layer 5-6 showed 86.1 ± 4.9%.

The dramatic improvement from Layer 1 (34.2%, barely above chance 2.1%) to Layer 4 (83.9%, approaching behavioral performance 87.3%) provided strong evidence for progressive information accumulation and abstraction through hierarchical processing. The saturation at Layer 5-6 (86.1%, not exceeding Layer 4) suggested that highest layers primarily re-represented information already extracted in Layer 4 rather than adding substantial new information, consistent with theories proposing that highest cortical layers implement working memory and attention rather than feature extraction.

The single-neuron selectivity analysis quantified how individual neurons represented categories through preference indices (firing rate difference between preferred and non-preferred categories normalized by sum). The distribution of preference indices shifted toward higher values in deeper layers: Layer 1 median preference 0.18 (indicating weak selectivity with preferred category only 1.22× more effective than average category), Layer 2-3 median 0.34 (preferred 1.52× more effective), Layer 4 median 0.52 (preferred 2.17× more effective), Layer 5-6 median 0.61 (preferred 2.56× more effective). The increasing

selectivity indicated individual neurons became increasingly specialized for particular categories through hierarchical processing.

The population sparseness analysis quantified how many neurons participated in representing each category through lifetime sparseness (calculated across categories for each neuron) and population sparseness (calculated across neurons for each category). The population sparseness increased across layers (Layer 1: 0.31, Layer 4: 0.68), indicating that while Layer 1 activated many neurons for each category (low sparseness), Layer 4 activated few highly selective neurons (high sparseness), consistent with emergence of grandmother cell-like category-specific representations.

3.2.3 Representational Geometry and Abstraction

The representational similarity analysis examined whether the similarity structure of neural representations (how similar different categories are in neural space) shifted from reflecting low-level perceptual features to high-level categorical structure. For each hierarchical level, we computed neural dissimilarity matrices measuring how different the population responses to each pair of categories were (1 - correlation between population response vectors for two categories), then compared these neural dissimilarity matrices to multiple model dissimilarity matrices reflecting different similarity principles.

The pixel-level similarity model computed dissimilarity as 1 - correlation between stimulus pixel patterns, capturing low-level visual similarity. The representational similarity analysis showed Layer 1 neural dissimilarity correlated strongly with pixel-level similarity (R^2 = 0.78, indicating neural representations primarily reflected pixel-level stimulus structure), Layer 2-3 showed intermediate correlation (R^2 = 0.54), and Layer 5-6 showed weak correlation (R^2 = 0.21), indicating progressive transformation away from pixel-based representation.

The structural similarity model computed dissimilarity using the structural similarity index (SSIM), a more perceptually relevant similarity metric accounting for structural information beyond pixel correspondence. This model correlated moderately with Layer 2-3 (R^2 = 0.64), peaking there before declining for higher layers (Layer 4: R^2 = 0.41, Layer 5-6: R^2 = 0.28), suggesting Layer 2-3 captured structural feature representations.

The categorical similarity model structured categories hierarchically (letters, digits, shapes as superordinate categories, with subordinate categories within each) and computed dissimilarity based on categorical distance in this hierarchy (dissimilarity = 0 for within-category pairs, 0.5 for different subordinate but same superordinate categories, 1.0 for different superordinate categories). This model showed increasing correlation across layers: Layer 1 ($R^2 = 0.12$), Layer 2-3 ($R^2 = 0.34$), Layer 4 ($R^2 = 0.58$), Layer 5-6 ($R^2 = 0.71$), indicating progressive emergence of categorical organization.

The multidimensional scaling visualization of representational dissimilarity matrices confirmed this progression. Layer 1 MDS plots showed categories arranged primarily by visual similarity with letters, digits, and shapes intermixed. Layer 4-6 MDS plots showed clear segregation into letter, digit, and shape clusters with within-cluster organization reflecting subcategory structure, demonstrating emergence of hierarchically organized categorical representations.

3.2.4 Invariance Properties

The invariance analysis examined whether representations maintained stability across transformations by measuring how neural responses to transformed stimuli compared to responses to original stimuli. For each transformation type (rotation, scaling, translation), we computed response correlation between original and transformed versions for neurons at each hierarchical level.

The position invariance tested through translation (±2, ±4 pixel shifts) showed Layer 1-2 neurons maintained only 15 ± 8% of peak response magnitude when stimuli shifted by 4 pixels (average across neurons weighted by response magnitude to original position), indicating strong position sensitivity. Layer 3-4 neurons maintained 48 ± 14% of response when shifted 4 pixels, indicating moderate position invariance. Layer 5-6 neurons maintained 67 ± 12% of response, indicating substantial but incomplete position invariance over the tested range (4 pixels = 12.5% of 32-pixel image width).

The size invariance tested through scaling (0.7×, 0.8×, 1.0×, 1.2×, 1.4×) showed similar hierarchy with Layer 5 neurons maintaining 72 ± 9% of response across ±20% size changes compared to 23 ± 11% for Layer 2 neurons. The rotation invariance proved more

limited, with Layer 5 neurons maintaining only 52 ± 14% of response for ±30 degree rotations, likely reflecting the limited training set that didn't extensively sample rotations and the general difficulty of rotation invariance which even primate visual cortex achieves only partially.

The invariance analysis revealed an important trade-off: neurons achieving invariance necessarily discarded information about the transformed variable (position, size, orientation). This information loss was evident in decoding analyses showing that while category identity could be decoded better from higher layers, stimulus position could be decoded worse (position decoding accuracy: Layer 1 = 78%, Layer 6 = 34%), confirming that invariance emerged through progressive information loss about behaviorally irrelevant dimensions while preserving category-relevant information.

3.3 Motor Control and Learning

The motor system's capacity for flexible motor output generation, learning of action-outcome contingencies, and skill acquisition was assessed through reaching tasks, sequence learning, force field adaptation, and instrumental conditioning.

3.3.1 Motor Learning and Skill Development

The reaching task training (50 trials per direction × 8 directions × 5 days = 2,000 trials total) showed progressive improvement in movement accuracy and efficiency. The naive performance (first 100 trials, day 1) showed success rate of only 23 ± 8% (movements reaching target within ±3 millimeter tolerance), with highly variable trajectories showing path efficiency (actual path length / optimal straight-line path length) of 2.7 ± 0.6 (indicating paths approximately 2.7× longer than necessary). The learning progression showed success rate improving to 48 ± 11% by day 2, 67 ± 9% by day 3, 79 ± 7% by day 4, and asymptoting at 84 ± 6% by day 5.

The movement trajectories became progressively straighter and more stereotyped with practice. The path efficiency improved from 2.7× on day 1 to 1.8× on day 2, 1.5× on day 3, 1.4× on day 4, and 1.3× on day 5, approaching theoretical minimum (1.0× would be perfect straight lines). The movement duration decreased from 2.8 ± 1.1 seconds initially to 1.7 ± 0.5

seconds by day 3 and 1.2 ± 0.3 seconds by day 5, indicating both increased accuracy and increased movement speed, a relatively rare learning pattern (often speed-accuracy tradeoffs prevent simultaneous improvement) suggesting genuine skill acquisition rather than simple strategy shifts. The velocity profiles evolved from irregular multi-peaked curves initially (indicating hesitant movements with multiple corrective submovements) to smooth bell-shaped curves after learning (characteristic of ballistic movements executed under feedforward control). The number of velocity peaks per movement declined from 3.4 ± 1.2 initially to 1.8 ± 0.6 by day 3 and 1.1 ± 0.3 by day 5, with the final values near 1.0 indicating predominantly single-peaked profiles.

The neural correlates of motor learning examined through recordings in motor organoids revealed several signature changes. The trial-to-trial variability in motor cortex population activity (quantified as coefficient of variation of population vectors across repeated movements to same target) decreased from 0.43 ± 0.11 initially to 0.18 ± 0.07 after learning, indicating emergence of stereotyped neural activity patterns underlying stereotyped movements. The population activity dimensionality (effective number of dimensions needed to capture 90% of variance) decreased from 8.4 ± 2.1 initially to 4.2 ± 1.3 after learning, indicating convergence onto lower-dimensional neural trajectories representing optimized motor solutions.

The preparatory activity (activity during 500 millisecond pre-movement planning period) became increasingly predictive of subsequent movement parameters. The decoding of movement direction from preparatory activity improved from 42% accuracy (chance = 12.5% for 8 directions) initially to 78% accuracy after learning. The decoding of movement velocity from preparatory activity improved from $R^2 = 0.24$ initially to $R^2 = 0.67$ after learning, indicating that motor plans became more precisely specified before movement onset, enabling feedforward control with reduced need for online corrections.

3.3.2 Force Field Adaptation

The force field adaptation experiment tested the motor system's capacity to build internal models of novel dynamics and adapt motor commands accordingly. The initial exposure to velocity-dependent forces (perpendicular to movement direction, proportional to velocity with gain constant creating substantial trajectory perturbations) produced large

errors with trajectories deviating 35 ± 12 millimeters from straight paths toward targets. The adaptation occurred rapidly with exponential error decay: errors halved by trial 45 ± 15 (trials needed to reduce error to 50% of initial), reached 25% of initial error by trial 108 ± 34, and asymptoted at 15 ± 6 millimeter residual error by trial 200-300.

The after-effects provided critical evidence for internal model adaptation rather than co-contraction or other non-adaptive compensation strategies. When the force field was unexpectedly removed (catch trials), movements showed errors in the opposite direction from initial perturbation errors, with magnitudes approximately 70 ± 15% of initial perturbation errors. These after-effects indicated that the motor system had adjusted commands to compensate for expected forces, and when those forces were absent, the compensatory commands produced opposite errors. The after-effects decayed over approximately 15-25 trials when the null field persisted, indicating gradual re-adaptation to the changed dynamics.

The savings effect tested by re-exposing systems to the same force field after washout (100 null-field trials) showed faster re-adaptation with initial errors reduced to 22 ± 8 millimeters (compared to 35 millimeters initially) and adaptation half-time reduced to 28 ± 10 trials (compared to 45 trials initially). This savings indicated retention of information about the previously experienced dynamics, likely through multiple mechanisms including residual synaptic changes and meta-learning of adaptation strategies.

The neural basis of force field adaptation examined through recordings in motor and cerebellar-like organoids revealed evolving population activity patterns. The motor cortex activity gradually shifted during adaptation such that decoded force output changed from the naive pattern (appropriate for null-field) toward a new pattern (producing forces offsetting the applied perturbation). The decoded force from population activity showed correlation with required compensatory force improving from $R^2 = 0.08$ initially to $R^2 = 0.73$ after adaptation. The cerebellar-like organoid activity showed prediction error encoding with neurons increasing firing when movement outcomes deviated from predictions, providing potential teaching signals driving motor cortex plasticity through cerebellar output projections.

3.3.3 Sequence Learning and Chunking

The sequence learning task (5-element motor sequences repeated 400 times over 8 days) revealed progressive development of hierarchical motor representations. The initial execution (day 1-2) was slow (4.2 ± 0.8 seconds per sequence) and error-prone (34 ± 11% sequences contained at least one error), with distinct pauses between elements (inter-element interval 620 ± 180 milliseconds). The learning progression showed execution time decreasing to 3.1 ± 0.6 seconds by day 4 and 1.7 ± 0.3 seconds by day 8, error rate decreasing to 18 ± 7% by day 4 and 9 ± 4% by day 8, and inter-element intervals decreasing to 280 ± 90 milliseconds by day 4 and 140 ± 50 milliseconds by day 8.

The chunking analysis examined whether sequences became organized into sub-units executed as coherent wholes. The inter-element interval patterns showed non-uniform structure emerging over learning, with some boundaries showing longer pauses than others. For the trained 5-element sequence (e.g., up-left-down-right-up), the interval pattern evolved from relatively uniform initially (all intervals 580-660 milliseconds) to structured finally (intervals: 180, 140, 380, 120 milliseconds), with the third interval (between element 3 and 4) substantially longer, suggesting a chunk boundary dividing the sequence into 3-element and 2-element chunks.

The error patterns supported chunking interpretation, with errors showing non-random distribution across sequence positions. Late in learning, errors occurred disproportionately at the identified chunk boundary (38% of errors at position 3→4 transition compared to 11-18% at other transitions), consistent with chunk boundaries being particularly error-prone as the system retrieved and initiated new chunks. Within-chunk errors (wrong element within chunk) were often followed by correction and continuation, while between-chunk errors often led to sequence restart or collapse, indicating chunks as stable units.

The dual-task interference tested whether sequence execution had become automatic by requiring simultaneous performance of a working memory task (remember 3-digit number presented at sequence start, report at sequence end). Early in learning (day 2), the dual-task condition increased errors by 78% (from 28% to 50%) and slowed execution by 45% (from 3.8 to 5.5 seconds), indicating substantial interference. Late in learning (day 8),

dual-task increased errors by only 18% (from 9% to 10.6%) and slowed execution by 12% (from 1.7 to 1.9 seconds), indicating largely automatic execution with minimal resource demands. The working memory task accuracy also showed less impairment from sequence execution late (82% correct dual-task versus 86% memory-only) versus early (61% versus 84%), confirming reduced interference in both directions.

The neural signatures of chunking appeared in motor cortex activity patterns. The dimensionality of neural trajectories (quantified through principal component analysis) decreased within chunks (2.8 ± 0.9 effective dimensions for within-chunk epochs) compared to between chunks (5.4 ± 1.8 dimensions), suggesting constrained stereotyped dynamics within chunks versus more variable transitional dynamics between chunks. The population activity correlation between repeated sequence executions increased within chunks (mean correlation 0.78 ± 0.12) compared to between chunks (0.54 ± 0.18), indicating highly reliable neural patterns within chunks versus more variable patterns at boundaries.

3.4 Learning and Memory

The memory systems' capacity for encoding, consolidating, and retrieving information across multiple timescales and formats was assessed through recognition memory, recall memory, spatial learning, and working memory tasks.

3.4.1 Episodic Memory and Recognition

The recognition memory task (40 studied patterns tested after delays of 5 minutes, 1 hour, 24 hours, or 7 days) revealed time-dependent memory decay with discriminability d' declining from 2.1 ± 0.3 at 5-minute delay to 1.7 ± 0.4 at 1-hour delay, 1.4 ± 0.3 at 24-hour delay, and 1.2 ± 0.3 at 7-day delay. The forgetting curve fit well to exponential decay function ($d'(t) = d'_0 \times \exp(-t/\tau)$) with time constant τ = 68 ± 23 hours (range across three systems: 52, 71, 81 hours), indicating substantial but gradual forgetting with half of initial memory strength remaining after approximately 47 hours.

The recognition performance showed appropriate serial position effects with primacy (enhanced memory for early-studied items) and recency (enhanced memory for late-studied items) effects. The items from first quartile of study list showed d' = 2.6 ± 0.4 (23% better

than overall mean), items from last quartile showed d' = 2.9 ± 0.5 (38% better), while middle items showed d' = 1.8 ± 0.3 (14% worse), creating the characteristic U-shaped serial position curve observed in human memory. The primacy effect persisted across delays (still evident at 24-hour delay), while recency effect diminished at longer delays (absent by 24 hours), consistent with theories proposing primacy reflects long-term memory consolidation while recency reflects residual working memory.

The false memory analysis examined whether systems showed systematic false recognition of non-studied lures. The false alarm rate (incorrectly identifying new items as old) showed baseline of 12 ± 6% for unrelated lures (randomly selected patterns with no particular relationship to studied items), but elevated to 28 ± 9% for related lures (similar to studied items through transformations or shared features). This pattern indicated that memory stored somewhat abstracted representations that confused similar items, rather than storing precise pixel-level patterns that would distinguish studied items from perceptually similar lures.

The confidence-accuracy calibration analysis examined metacognitive monitoring of memory quality. The systems provided confidence judgments (high/low) following recognition decisions, enabling calculation of accuracy separately for high-confidence and low-confidence responses. The high-confidence responses showed d' = 2.8 ± 0.4, while low-confidence responses showed d' = 1.1 ± 0.3, indicating that confidence appropriately tracked memory strength. The metacognitive efficiency quantified through meta-d' = 2.4 ± 0.5 (compared to task d' = 2.1 ± 0.3) and M-ratio = 1.14 ± 0.18 indicated good but imperfect metacognition, with confidence judgments providing reliable but not perfect indicators of memory accuracy.

3.4.2 Neural Mechanisms of Memory Encoding and Retrieval

The hippocampal activity during memory encoding revealed several signatures consistent with episodic memory formation. The encoding success analysis compared hippocampal activity during study of items that were subsequently remembered versus forgotten (assessed through test performance). The subsequently-remembered items showed 32 ± 12% higher hippocampal activity (mean firing rate across hippocampal neurons

during encoding period) compared to subsequently-forgotten items ($p < 0.001$, paired t-test across sessions), indicating that encoding strength at study predicted later memory.

The pattern separation analysis tested whether hippocampal circuits distinguished similar experiences by examining neural representation overlap for similar stimuli. Pairs of similar patterns (differing by only 20-30% of pixels) evoked hippocampal population patterns that were significantly decorrelated (population vector correlation 0.18 ± 0.12) compared to repeated presentations of identical patterns (correlation 0.78 ± 0.14), indicating active pattern separation transforming similar inputs into distinct representations. This decorrelation was much stronger in hippocampus than in sensory cortex (where similar patterns showed correlation 0.62 ± 0.18), confirming hippocampus-specific pattern separation computation.

The pattern completion analysis tested whether partial cues could reinstate complete memory representations. During retrieval, presentation of partial patterns (only 50% of original pixels shown, remaining 50% blank) evoked hippocampal population activity significantly correlated with activity evoked by complete patterns during encoding (correlation 0.54 ± 0.18), indicating retrieval of stored complete representations from partial cues. The pattern completion strength (measured through correlation magnitude) predicted recognition accuracy, with stronger pattern completion associated with higher recognition confidence and accuracy.

The consolidation analysis examined memory reactivation during post-learning rest periods. The hippocampal spontaneous activity during rest periods following study sessions showed elevated co-activation of neuron pairs that were co-active during preceding study (measured through correlation increase of $42 \pm 15\%$ during post-learning rest versus pre-learning baseline), indicating replay of learned associations. The replay preferentially occurred during sharp-wave ripple events (100-200 Hz oscillations lasting 50-150 milliseconds), with ripple-associated reactivation 8.3 ± 2.4-fold stronger than baseline, confirming that ripples served as consolidation events replaying recent experiences.

The systems-level consolidation examined whether memory representations gradually transferred from hippocampus to neocortical semantic memory regions. The

decoding analysis assessed how well memory content could be decoded from hippocampal versus cortical activity at different time points post-learning. Immediately after learning (5-minute delay), hippocampal decoding accuracy (68 ± 9%) substantially exceeded cortical decoding (38 ± 12%), indicating hippocampus-dependent memories. At 24-hour delay, hippocampal decoding remained high (62 ± 11%) while cortical decoding improved (51 ± 14%). At 7-day delay, both regions showed similar decoding (hippocampus 54 ± 13%, cortex 48 ± 11%), indicating progressive cortical representation development. This time-dependent shift from hippocampus to cortex is consistent with standard consolidation theory proposing hippocampus enables rapid learning with gradual transfer to cortex for long-term storage.

3.4.3 Spatial Learning and Navigation

The spatial navigation task (virtual 10×10 grid environment, goal location learned across 200 trials over 5 days) showed robust spatial learning with navigation performance improving from initial random search (trial duration 45 ± 18 seconds, path efficiency 3.8 ± 1.4) to efficient goal-directed navigation (trial duration 8.2 ± 3.4 seconds, path efficiency 1.4 ± 0.3 by day 5). The learning curve showed approximately exponential improvement with time constant 52 ± 18 trials, indicating rapid spatial learning.

The probe trials removing the goal tested whether navigation relied on beacon following versus true spatial knowledge. During probe trials, the systems continued navigating toward the trained goal location despite absence of local goal cues, with 78 ± 12% of probe trials showing navigation to within 2 grid positions of goal location. The probe trial performance confirmed spatial memory rather than simple stimulus-response associations, as appropriate goal-directed behavior persisted when the reinforcement-providing goal was absent.

The detour testing introduced barriers blocking direct paths to goals, requiring multi-step indirect routes. The systems showed appropriate detour behavior with success rates of 67 ± 14% on novel barrier configurations (compared to 89 ± 7% on familiar direct-path trials). The detour latencies (time from trial start to first movement) were substantially elevated (4.2 ± 1.8 seconds for detour trials versus 0.9 ± 0.4 seconds for direct trials), suggesting a planning period during which systems computed alternative routes. The

detour paths showed reasonable optimality (1.8 ± 0.4 times optimal path length), indicating genuine path planning rather than random exploration until goal discovery.

The place cell analysis revealed that 23.7 ± 6.4% of hippocampal neurons showed significant spatial tuning (ANOVA across locations, $p < 0.01$ with Bonferroni correction). The place fields (locations where neurons fired maximally) covered the environment approximately uniformly, with receptive field density varying only 1.8-fold across locations (range 18-32 place fields per location) indicating relatively uniform spatial representation. The place field sizes averaged 2.1 ± 0.8 grid positions (approximately 20% of environment), comparable to mammalian place field sizes which typically span 10-30% of environment dimensions.

The spatial information content (mutual information between spike occurrences and location) averaged 0.47 ± 0.18 bits per spike for spatially tuned neurons, indicating that each spike reduced spatial uncertainty by approximately 1.4-fold ($2^{0.47} = 1.38$). This value falls within the range reported for rodent place cells (typically 0.3-0.7 bits/spike), validating that engineered place cells achieved comparable spatial encoding efficiency to biological place cells despite being generated in vitro.

The head direction coding examined whether neurons encoded movement direction in addition to position. Among recorded hippocampal and adjacent neurons, 12.4 ± 4.8% showed significant directional tuning (ANOVA across 8 movement directions, $p < 0.01$), with preferred directions distributed approximately uniformly. The directional information averaged 0.31 ± 0.14 bits per spike for directionally tuned neurons, indicating modest directional encoding. The combined position-direction encoding neurons (showing both place fields and directional tuning, 7.2 ± 3.1% of recorded population) provided conjunctive representation potentially supporting path integration and trajectory planning.

3.4.4 Working Memory

The working memory capacity was assessed through delayed match-to-sample tasks requiring temporary maintenance of information across delay periods. The task structure presented a sample stimulus (1-6 patterns simultaneously), followed by delay period (3, 6, or 9 seconds), followed by test stimulus (matching or non-matching). The systems indicated match/non-match judgment via actuator selection, with performance (d' discriminability) quantifying working memory capacity.

The capacity analysis varied the number of patterns requiring maintenance (set size 1-6) and measured performance decline with increasing load. The discriminability showed approximately linear decline: d' = 2.8 ± 0.4 for set size 1, 2.3 ± 0.5 for size 2, 1.8 ± 0.4 for size 3, 1.4 ± 0.5 for size 4, 0.9 ± 0.4 for size 5, 0.6 ± 0.3 for size 6. The capacity estimated as set size at which performance declined to d' = 1.0 (moderate but reliable discrimination) averaged 4.2 ± 0.8 items, comparable to human visual working memory capacity (typically 3-5 items depending on stimulus complexity).

The delay-dependent forgetting showed modest performance decline with longer delays: d' declined approximately 15% from 3-second to 6-second delays and 28% from 3-second to 9-second delays (averaged across set sizes 1-4 where performance remained above floor). The forgetting rate approximately 3.1% per second suggested reasonably stable maintenance mechanisms capable of sustaining information across behaviorally relevant timescales.

The neural mechanisms of working memory examined through persistent activity in prefrontal-like organoids revealed neurons maintaining elevated firing during delay periods. Among recorded neurons in prefrontal regions, 18.4 ± 5.7% showed significant delay period activity (firing rate during delay >150% of baseline, sustained throughout delay, $p < 0.01$). The persistent activity magnitude (delay period firing rate normalized by baseline) correlated with subsequent recognition accuracy ($R = 0.42 \pm 0.14$, $p < 0.01$), indicating that neural activity stability during maintenance predicted memory performance.

The population-level working memory analysis through dimensionality reduction revealed that working memory maintenance occurred in low-dimensional subspaces of

population activity space. The delay period activity occupied approximately 3.4 ± 1.2 effective dimensions (dimensions needed to capture 90% of variance), substantially lower than the full population dimensionality (~200 recorded neurons would support ~200 dimensions), indicating constrained dynamics supporting stable maintenance. The working memory subspace showed relative stability across different maintained items (different stimulus identities), with subspace overlap (principal angle between subspaces for different items) averaging 68 ± 12 degrees (where 0° would be identical subspace, 90° would be orthogonal subspaces), suggesting similar dynamical mechanisms maintaining different memory contents.

3.5 Motivation and Urge Dynamics

The motivational system's capacity for generating, integrating, and responding to multiple urges was assessed through glucose deprivation experiments, uncertainty manipulation, competence motivation tasks, and multi-urge conflict scenarios.

3.5.1 Physiological Urge Dynamics

The glucose deprivation experiment systematically reduced glucose concentration in perfusion medium from normal 5.5 millimolar to 4.0, 3.5, 3.0, 2.5, and finally 2.0 millimolar over 6 hours (1 hour at each level), while monitoring neural activity in energy-urge-generating organoids, metabolic parameters, and behavioral outputs. The energy urge activity (firing rate of neurons in hypothalamus-like urge organoids) showed graded increases: baseline 2.1 ± 0.4 Hz at normal glucose, increasing to 3.2 ± 0.6 Hz at 4.0 mM, 4.8 ± 0.9 Hz at 3.5 mM, 6.4 ± 1.2 Hz at 3.0 mM, 8.7 ± 1.3 Hz at 2.5 mM, and 11.2 ± 1.8 Hz at 2.0 mM. The approximately exponential relationship between glucose concentration and urge activity (fitted function: activity = 18.4 × exp(-0.42 × [glucose])) indicated appropriate sensitivity spanning physiological glucose range.

The behavioral consequences of energy urges manifested as biased action selection favoring energy-relevant actions. During testing with choice tasks offering energy-relevant options (actions historically associated with glucose delivery) versus energy-neutral options (equal reward magnitude but delivered as dopamine stimulation rather than glucose), the systems shifted strongly toward energy-relevant options as glucose declined: 52% selection

at normal glucose (near-indifferent), 64% at 3.5 mM, 79% at 3.0 mM, 88% at 2.5 mM. This graded modulation confirmed that energy urges appropriately biased action selection toward satisfying the underlying metabolic need.

The attention modulation analysis examined whether energy depletion shifted attentional resources toward energy-relevant information. Eye-movement-analog measurements (tracking which regions of visual space received enhanced sensory processing through analysis of Layer 1 population activity patterns) showed that energy-relevant stimuli (patterns previously associated with glucose delivery) captured attention more strongly during energy depletion: fixation duration on energy-relevant stimuli increased from 320 ± 80 milliseconds at normal glucose to 580 ± 120 milliseconds at 2.5 millimolar glucose, while fixation on energy-neutral stimuli remained stable (280 ± 70 versus 310 ± 90 milliseconds). The selective attention enhancement for behaviorally relevant information confirmed motivation-cognition integration.

The glucose restoration examined how quickly urge systems responded to need satisfaction. Upon restoration of normal glucose levels (switching from 2.5 millimolar to 5.5 millimolar), the urge activity declined with approximately exponential kinetics (half-time 12 ± 4 minutes), indicating moderately rapid but not instantaneous urge updating. The behavioral bias similarly declined with half-time 18 ± 6 minutes, lagging slightly behind neural urge activity and suggesting gradual behavioral adaptation.

The integrity urge assessment employed mild thermal stress (elevating temperature from 37.0°C to 38.5°C for 30 minutes) while monitoring stress markers, urge activity, and defensive behaviors. The heat shock protein HSP70 expression increased 3.2 ± 0.8-fold (measured through fluorescent reporter), indicating cellular stress response activation. The integrity urge neurons (located in separate organoid from energy urge neurons to enable independent monitoring) increased firing from baseline 1.9 ± 0.4 Hz to 7.8 ± 1.6 Hz during stress, indicating robust urge generation. The behavioral changes included reduced exploration (spontaneous movement frequency decreased 47%), increased conservative actions (choosing familiar over novel options when offered choice), and temporary performance decline (12 ± 4% accuracy reduction on cognitive tasks during stress period),

all consistent with defensive prioritization of homeostatic maintenance over cognitive performance during threat.

3.5.2 Cognitive Urge Dynamics

The uncertainty manipulation experiment varied prediction error magnitude to modulate certainty-reduction urges. The systems performed visual discrimination tasks with varying stimulus ambiguity (noise levels creating 50%, 70%, or 90% discrimination accuracy). The prediction error (measured through mismatch between predicted and actual outcomes, quantified through activity in error-coding neurons) scaled with task difficulty: 0.11 ± 0.04 normalized units at 90% accuracy (low prediction error when predictions accurate), 0.28 ± 0.08 at 70% accuracy, 0.42 ± 0.08 at 50% accuracy (high prediction error when predictions unreliable).

The certainty urge activity (neurons in urge-generating organoids distinct from energy and integrity urge populations) showed corresponding increases: baseline 2.3 ± 0.5 Hz at 90% accuracy, increasing to 4.7 ± 0.9 Hz at 70% accuracy and 6.8 ± 1.1 Hz at 50% accuracy. The information-seeking behavior (measured through voluntary selection of informative versus uninformative sensory samples when given choice) increased from 32% informative selections at 90% accuracy to 58% at 70% accuracy and 67% at 50% accuracy, confirming that elevated uncertainty drove information-seeking motivation.

The learning rate modulation examined whether uncertainty influenced plasticity. The behavioral learning rates (estimated through fitting exponential learning curves to acquisition data) showed 2.4 ± 0.7-fold elevation in high-uncertainty conditions (50% accuracy) compared to low-uncertainty conditions (90% accuracy), indicating faster learning when predictions were unreliable. This adaptive learning rate modulation, predicted by normative learning theories proposing learning rates should scale with environmental volatility, confirmed sophisticated meta-learning mechanisms adjusting learning based on uncertainty estimates.

The competence motivation assessment employed tasks with varying performance feedback. The systems performed cognitive tasks where performance gradually improved with practice, with explicit performance feedback provided (accuracy scores displayed through sensory patterns). The competence urge activity (distinct population from other urge types) showed elevated firing (5.9 ± 1.2 Hz compared to baseline 2.3 ± 0.6 Hz) specifically during periods of performance improvement (trials where accuracy increased

compared to recent average), but not during stable performance periods or performance declines. This selective activation during positive learning gradients confirmed competence motivation responding to learning progress rather than absolute performance.

The intrinsic motivation analysis tested whether competence urges sustained behavior in absence of external rewards. Following initial training with explicit rewards, the rewards were removed (unrewarded trials) while maintaining performance feedback. Despite absence of external reinforcement, the systems continued performing tasks with only modest performance decline (accuracy decreased 12 ± 5% compared to rewarded baseline), and preferentially selected tasks showing learning opportunity (novel challenging tasks) over mastered tasks (88% vs 45% selection rates). This preference for challenging learning opportunities despite lack of external reward confirmed intrinsic competence motivation, a key prediction of Bach's Psi framework.

3.5.3 Urge Integration and Conflict Resolution

The multi-urge conflict scenario presented situations where multiple urges activated simultaneously with potentially conflicting behavioral implications. The scenario manipulated glucose levels (creating energy urges), prediction error (creating certainty urges), and learning opportunity (creating competence urges) orthogonally across trials, enabling assessment of how urges combined to determine behavior.

The behavioral choice analysis examined which actions systems selected when multiple urges supported different options: energy-seeking actions satisfying energy urges, information-seeking actions satisfying certainty urges, or skill-practicing actions satisfying competence urges. The choice patterns revealed hierarchical prioritization where energy urges dominated when glucose was severely depleted (glucose < 3.0 millimolar: 78 ± 11% energy-seeking selections regardless of other urge states), but when energy needs were moderate, attention shifted to cognitive urges (glucose 3.5-4.5 millimolar: 42 ± 13% energy-seeking, 31 ± 12% information-seeking, 27 ± 11% competence-driven, showing balanced distribution).

The quantitative modeling of urge integration tested several computational accounts of how multiple urges combine. The additive model (choice value = $\sum_i w_i \times U_i \times R_i$ where w_i

are weights, U_i are urge magnitudes, R_i are action relevances to each urge) fit behavioral data reasonably well ($R^2 = 0.64 \pm 0.12$ for predicting choice probabilities across conditions). The multiplicative model (choice value = $\prod_i (w_i \times U_i \times R_i)$) fit slightly worse ($R^2 = 0.58 \pm 0.14$), suggesting primarily additive combination. The hierarchical model (priority ordering where higher-priority urges must be partially satisfied before lower-priority urges influence behavior) fit best ($R^2 = 0.71 \pm 0.09$), with fitted hierarchy: energy > integrity > certainty > competence, matching intuitive importance ordering.

The temporal dynamics of urge competition examined how behavior shifted over extended periods as urges evolved. During 2-hour sessions with no glucose supplementation (allowing progressive energy depletion) while cognitive tasks remained available, the systems showed gradual behavioral shifts: early period (first 30 minutes, glucose declining from 5.5 to 4.2 millimolar) showed primarily competence-driven task engagement (72% trials), middle period (minutes 30-60, glucose 4.2 to 3.5 millimolar) showed mixed behavior (43% competence, 31% information-seeking, 26% energy-seeking), late period (minutes 60-120, glucose 3.5 to 2.8 millimolar) showed primarily energy-seeking (81% trials). This temporal evolution reflected the progressive rise of energy urges eventually dominating behavior as metabolic needs intensified.

The behavioral tradeoffs quantified through choice indifference analysis identified points where different urges balanced. For example, the glucose level at which systems showed equal preference for energy-seeking versus information-seeking actions (indifference point) averaged 3.8 ± 0.3 millimolar, substantially below normal levels (5.5 millimolar) but well above critical levels (2.5 millimolar), indicating that energy concerns dominated only when deprivation became moderately severe. The indifference curves mapping combinations of urge magnitudes yielding equal choice probabilities revealed smooth continuous tradeoffs rather than discrete switching, confirming graduated urge integration rather than winner-take-all competition.

3.6 Planning and Forward Simulation

The planning system's capacity for mental simulation, tree search, and model-based decision making was assessed through two-step tasks, detour problems, Tower of Hanoi analogs, and dynamic replanning scenarios.

3.6.1 Model-Based versus Model-Free Learning

The two-step sequential decision task administered over 400 trials revealed substantial model-based behavior, indicating planning through internal simulation rather than purely cached value learning. The critical diagnostic—the interaction between previous trial outcome (rewarded/unrewarded) and transition type (common/rare) when predicting first-stage choice—showed significant positive interaction coefficients of β = 0.089 ± 0.024 (logistic regression coefficient, z = 3.71, p < 0.001). The positive interaction indicated that reward following common transitions produced stronger stay probability than reward following rare transitions, consistent with model-based planning taking transition structure into account.

The model-based index (interaction coefficient normalized by main effect coefficient) averaged 0.67 ± 0.08, where 0 represents purely model-free and 1.0 represents purely model-based, indicating predominantly but not exclusively model-based behavior. This value approached human performance on similar tasks (humans typically show model-based index 0.6-0.8 depending on age and training), suggesting comparable levels of planning capacity.

The decomposition of behavior through computational modeling fitted both model-free Q-learning and model-based planning models simultaneously (hybrid model allowing weighted mixture), estimating relative contributions. The fitted model weights showed 68 ± 12% weighting on model-based component versus 32 ± 12% on model-free component, confirming predominant but not exclusive reliance on planning. The individual systems showed modest variation (System A: 74% model-based, System B: 66%, System C: 64%), consistent with individual differences observed in other domains.

The neural correlates of model-based planning examined through recordings in planning organoids during task performance revealed signature patterns. The forward model organoids (implementing state prediction) showed activity encoding upcoming states

before they were actually reached: decoding analysis could predict which second-stage state would be reached (S1 versus S2) from forward model activity during the 500 millisecond first-stage choice period with 71 ± 9% accuracy (significantly above 50% chance, $p < 0.001$). This prospective coding indicated mental simulation of action consequences before action execution, the defining feature of model-based planning.

The value updating analysis compared how reward prediction errors (differences between received and expected rewards) influenced first-stage action values. Model-free learning predicts that rewards directly update values of first-stage actions that led to them, while model-based learning predicts rewards update second-stage action values then propagate backward to first-stage through learned transition model. The fitted learning algorithms showed substantial backward propagation (67 ± 15% of reward credit assigned to first-stage actions propagated through transition model versus assigned directly), confirming model-based credit assignment.

3.6.2 Detour Planning and Problem Solving

The detour task requiring navigation around newly introduced barriers showed robust planning signatures including pre-movement pauses, sequential state reactivation, and near-optimal path selection. The sollution latency analysis showed substantially longer pre-movement times on barrier trials (4.2 ± 1.8 seconds) compared to familiar direct-path trials (0.9 ± 0.4 seconds), suggesting mental planning during this delay. The solution success rate of 78 ± 9% on novel barrier configurations indicated reliable planning capacity, substantially above the 18 ± 7% success rate achieved by a control system (System D, generated identically but with forward model organoids lesioned via optogenetic silencing, demonstrating planning necessity).

The neural mechanisms of detour planning were examined through population decoding during pre-movement pauses. The decoding analysis applied to activity in planning organoids during the pause period revealed sequential reactivation of place-cell-like representations corresponding to locations along candidate paths from start to goal. The decoded position sequence showed spatial trajectories that progressed from current position toward goal with 61 ± 13% of time points decoded as positions along the subsequently executed path, substantially above the 15% expected if decoding reflected

random positions. The temporal compression of decoded trajectories averaged 18.4 ± 6.2-fold, meaning the neural replay of a path requiring 10 seconds to execute physically was simulated in approximately 0.54 seconds, enabling rapid evaluation of multiple alternative paths.

The path optimality analysis showed executed paths were 1.18 ± 0.14 times optimal path length on average (where 1.0 would be perfect), indicating near-optimal planning. The occasional suboptimal paths (1.3-1.6× optimal) often resulted from incomplete exploration of state space during planning, with systems sometimes committing to paths after brief simulation rather than exhaustively searching all alternatives. The path quality positively correlated with planning duration ($R = 0.48$, $p < 0.01$), with longer planning pauses yielding more optimal paths, confirming that extended simulation improved solution quality.

The barrier configuration transfer tested whether planning generalized across problems. Systems trained on one set of barrier configurations (e.g., vertical walls) and tested on qualitatively different configurations (diagonal walls, scattered obstacles) showed 64 ± 11% success rates on transfer configurations, substantially better than naive performance (22 ± 9% for systems encountering barriers for first time) but below same-type performance (78%), indicating partial but imperfect transfer. The partial transfer suggested that systems learned general planning strategies applicable across configurations rather than memorizing configuration-specific solutions, but that some aspects of planning remained configuration-specific.

3.6.3 Hierarchical Planning and Subgoal Decomposition

The Tower of Hanoi analog task requiring 7-move optimal solutions showed success rates of 42 ± 11% (solving within 20-move limit), with move efficiency averaging 9.2 ± 1.7 moves for successful solutions (32% above optimal 7 moves, but dramatically better than random exploration requiring ~50 moves). The planning time (delay before first move) averaged 8.3 ± 2.1 seconds, substantially longer than reaction times for simple motor responses (0.8 ± 0.3 seconds), indicating deliberative planning.

The subgoal structure analysis examined pause patterns between moves, testing whether longer pauses occurred at natural subgoal boundaries. The expert solution strategy

decomposes the problem into subgoals: "expose smallest disk" → "expose medium disk" → "move medium disk to target" → "build on medium disk". The move-by-move pause duration showed peaks at predicted subgoal boundaries (moves 2→3 and 4→5 showing pauses 1.8 ± 0.6 seconds compared to within-subgoal pauses 0.4 ± 0.2 seconds), confirming hierarchical solution structure with pauses for subgoal retrieval or planning.

The error analysis revealed that errors clustered at subgoal boundaries (43% of errors at boundary moves versus 11-15% at within-subgoal moves), consistent with hierarchical representations where transitions between subgoals are particularly error-prone while within-subgoal execution proceeds smoothly once the subgoal is retrieved. The error recovery patterns showed that within-subgoal errors often led to correction and continuation (67% of within-subgoal errors followed by valid completion), while boundary errors often led to solution abandonment or restart (only 28% successful recovery after boundary errors), suggesting subgoals as stable units that must be executed correctly.

The neural representations of hierarchical structure were examined through activity in prefrontal planning organoids. The population activity dimensionality analysis showed that different subgoals occupied distinct regions of neural state space (principal angle between subspace for subgoal 1 versus 2: 68 ± 14 degrees), indicating separate neural representations for different high-level goals. Within subgoals, the neural trajectories showed constrained dynamics in low-dimensional subspaces (2.4 ± 0.8 effective dimensions), while transitions between subgoals showed higher-dimensional trajectories (4.8 ± 1.4 dimensions), suggesting stable attractors for subgoals with more complex dynamics during transitions.

3.6.4 Dynamic Replanning and Flexibility

The dynamic replanning task with mid-trial goal changes tested the system's capacity to detect changed circumstances, abandon current plans, generate new plans, and execute revised behaviors. The goal change detection latency (time from goal position change to first movement direction change) averaged 1.8 ± 0.6 seconds, indicating reasonably rapid detection and response. The replanning accuracy (whether new movement

direction was appropriate for new goal location) reached 83 ± 9%, indicating that most detected goal changes led to appropriate replanning.

The completion rates (proportion of trials where systems reached final goal despite mid-trial changes) averaged 68 ± 9% for single goal changes and 42 ± 13% for double goal changes (two successive changes per trial), compared to 89 ± 7% for static goal trials. The performance decline confirmed that replanning imposed costs, but maintenance of above-50% success even with double changes indicated substantial replanning capacity.

The neural correlates of replanning examined through real-time decoding of planned trajectories showed abrupt switches in decoded goal following goal changes. The decoded intended goal (predicted from planning organoid population activity) switched from original to new goal location with latency 0.8 ± 0.4 seconds (faster than behavioral response latency 1.8 seconds), indicating neural replanning preceded behavioral implementation. The interim period (between neural replanning and behavioral execution) likely reflected motor preparation time needed to reconfigure motor systems for new trajectory.

The perturbation magnitude effects tested whether replanning difficulty scaled with how dramatically goals changed. Small goal displacements (1-2 grid positions, requiring minor plan adjustments) showed replanning latencies of 1.2 ± 0.4 seconds and 91 ± 6% accuracy. Medium displacements (3-5 positions, requiring substantial replanning) showed 1.8 ± 0.6 seconds and 83 ± 9% accuracy. Large displacements (6-10 positions, requiring complete replanning) showed 2.9 ± 1.1 seconds and 74 ± 12% accuracy. The graceful degradation (smoothly decreasing performance rather than catastrophic failure) suggested flexible adaptive planning able to handle varying perturbation magnitudes.

3.7 Metacognition and Self-Monitoring

The metacognitive system's capacity for monitoring internal states, estimating confidence, detecting errors, and adaptively controlling cognitive processing was assessed through confidence judgment tasks, opt-out paradigms, error detection, and information-seeking behavior.

3.7.1 Confidence Estimation and Calibration

The confidence judgment task requiring binary high/low confidence reports following perceptual discriminations revealed good but imperfect metacognitive sensitivity. The Type 2 signal detection analysis comparing confidence against actual performance yielded area under Type 2 ROC curve of 0.78 ± 0.08, significantly above chance (0.5) indicating that confidence reliably distinguished correct from incorrect trials, but below perfect metacognition (1.0). The meta-d' analysis quantifying metacognitive sensitivity in comparable units to task sensitivity yielded meta-d' = 1.8 ± 0.3 compared to task d' = 2.1 ± 0.3, yielding M-ratio (meta-d'/d') = 0.85 ± 0.12, indicating good but slightly suboptimal metacognitive efficiency.

The confidence-accuracy correlation analysis showed that high-confidence responses were indeed more accurate than low-confidence responses across all task difficulty levels. For easy discriminations (90% baseline accuracy), high-confidence trials showed 94 ± 3% accuracy versus 82 ± 6% for low-confidence trials. For medium difficulty (70% baseline), high-confidence showed 81 ± 5% versus 62 ± 8% low-confidence. For hard discriminations (55% baseline), high-confidence showed 67 ± 9% versus 48 ± 7% low-confidence. The consistent confidence-accuracy relationship across difficulty levels indicated that confidence tracked multiple sources of information influencing accuracy (stimulus quality, internal noise, attention) rather than relying on single cue.

The confidence calibration curves plotting actual accuracy against reported confidence showed reasonable but imperfect calibration. The low-confidence responses showed 58 ± 9% accuracy (ideally would be 50% for perfect calibration), while high-confidence responses showed 82 ± 7% accuracy (ideally would be 100%). The calibration slopes (regression of accuracy on confidence) averaged 0.73 ± 0.16, significantly positive ($p < 0.001$) indicating meaningful confidence information, but below perfect calibration slope of 1.0, indicating overconfidence (reporting higher confidence than warranted, particularly for high-confidence trials).

The reaction time relationships showed appropriate patterns with confidence inversely related to decision time. High-confidence responses showed reaction times of 680 ± 140 milliseconds, while low-confidence responses showed 940 ± 180 milliseconds, indicating that uncertain trials required longer deliberation. The decision time itself

predicted accuracy (faster responses were more accurate, $R = -0.38$, $p < 0.001$), suggesting that decision time provided metacognitive information that appropriately influenced confidence judgments.

3.7.2 Error Detection and Correction

The error detection task allowing post-decision response changes before feedback revealed genuine error monitoring capacity. The error detection sensitivity (treating change-response decisions as discrimination between error and correct trials) showed $d' = 1.9 \pm 0.4$, indicating reliable discrimination between errors and correct responses. The hit rate (proportion of errors where response changed) averaged 56 ± 9%, while false alarm rate (proportion of correct trials where response changed) was only 11 ± 4%, confirming selective changing after errors.

The error-related neural activity examined through recordings in metacognitive self-monitoring organoids revealed error-specific signals. Comparing neural activity following errors versus correct trials (in 200-500 millisecond post-response window) showed that 34 ± 8% of metacognitive neurons exhibited significantly elevated firing after errors (mean rate 8.4 ± 2.1 Hz after errors versus 3.2 ± 1.4 Hz after correct, $p < 0.001$ paired t-test). This error-related activity resembled the error-related negativity observed in human EEG studies, providing evidence for neural error-monitoring signals.

The error-related activity predicted subsequent change-response decisions, with trials showing strong error-related activity more likely to result in response changes (72 ± 13% trials with top-quartile error activity showed changes, versus 28 ± 11% for bottom-quartile error activity). This prediction confirmed that error-related neural activity caused behavioral responses rather than merely correlating with errors. The temporal dynamics showed error-related activity appearing 380 ± 120 milliseconds after initial response, providing sufficient time for error detection before the 2-second change-response window closed.

The error correction effectiveness analysis examined whether changed responses were more accurate than initial responses. Among trials where responses changed, the final (changed) response showed 68 ± 12% accuracy, substantially better than the initial

response accuracy of 34 ± 9% (by definition, since errors were more likely to be changed). This improvement confirmed that error detection enabled genuine correction rather than random response switching. The absolute accuracy of 68% indicated that error detection, while imperfect, provided substantial information supporting better decisions.

3.7.3 Opt-Out Behavior and Uncertainty Avoidance

The opt-out paradigm offering opportunity to decline difficult trials showed strategic use of opt-out option calibrated to trial difficulty and expected performance. The opt-out rate varied systematically with task difficulty: 4 ± 2% for easy trials (90% baseline accuracy), 18 ± 6% for medium trials (70% accuracy), 47 ± 11% for hard trials (55% accuracy), 67 ± 13% for very hard trials (near-chance accuracy). This graded relationship confirmed that opt-out decisions tracked task difficulty and expected success probability.

The earnings optimization analysis compared actual earnings (counting correct rewards, subtracting error penalties, and including opt-out rewards) against optimal earnings achievable with perfect metacognition. The systems achieved 87 ± 9% of optimal earnings, indicating near-optimal but not perfect metacognitive decision-making. The earnings losses (13% below optimal) occurred primarily on medium-difficulty trials where optimal strategy required fine-grained discrimination of marginal cases but systems showed conservative opt-out bias (opting out more than optimal).

The signal detection analysis treating opt-out as confidence-based decision yielded opt-out sensitivity of d' = 2.2 ± 0.4 (discriminating trials that should versus should not be attempted), higher than confidence judgment sensitivity (d' = 1.8), possibly because opt-out decisions involved actual consequences (forfeiting reward opportunity) that focused metacognitive processes. The optimal threshold analysis calculated difficulty level where expected value of attempting trial equaled opt-out value, finding optimal threshold at 63% expected accuracy. The observed opt-out threshold (50% probability fitted from logistic function) occurred at 58 ± 7% accuracy, close to but slightly conservative relative to optimal.

The individual differences in opt-out strategy revealed interesting variation across systems. System A showed liberal opt-out strategy (opting out when expected accuracy <65%, earning 84% of optimal), System B showed near-optimal strategy (threshold 58%,

earning 91% of optimal), System C showed conservative strategy (threshold 51%, rarely opting out, earning 85% of optimal). These differences reflected different risk attitudes or metacognitive biases, with System B achieving best balance between attempting trials and avoiding errors.

3.7.4 Adaptive Cognitive Control

The adaptive control analysis examined whether metacognitive monitoring influenced subsequent cognitive processing. The post-error slowing analysis compared reaction times on trials immediately following errors versus following correct responses. Post-error trials showed 15 ± 6% longer reaction times (1024 ± 190 milliseconds versus 890 ± 160 milliseconds post-correct, $p < 0.01$), indicating adaptive slowing to improve accuracy after errors. This slowing was associated with improved accuracy (post-error trials showed 3.2 ± 1.8% higher accuracy than post-correct trials), confirming that slowing served adaptive function rather than merely reflecting distraction or frustration.

The confidence-based adjustments examined whether low-confidence trials led to increased resource allocation on subsequent trials. Following low-confidence trials, several adaptive changes occurred: (1) Decision time increased by 12 ± 5% on next trial (suggesting increased deliberation), (2) Attention to task-relevant features increased by 18 ± 7% (measured through sensory population activity), (3) Learning rates increased by 32 ± 14% (measured through behavioral adjustment speed), all relative to post-high-confidence trials. These adjustments indicated sophisticated metacognitive control adapting processing based on experienced uncertainty.

The strategy switching analysis examined whether sustained poor performance triggered exploration of alternative strategies. During extended task blocks (200 trials), we identified periods of below-average performance (accuracy <mean for >20 consecutive trials) and analyzed subsequent behavior. Following such periods, behavioral variability increased by 38 ± 12% (measured through entropy of action distributions), indicating exploration of alternative approaches. This exploration led to discovery of improved strategies in 67 ± 18% of low-performance episodes (identified by subsequent performance improvement), confirming adaptive strategy search.

The learning rate modulation analysis tested theoretical predictions that learning rates should adjust based on uncertainty. In environments with higher volatility (frequent changes in reward contingencies), learning rates should increase to track changing contingencies. When environmental volatility was experimentally manipulated (reward probabilities stable for 50 trials versus changing every 10 trials), learning rates showed appropriate modulation: stable environment yielded fitted learning rate α = 0.12 ± 0.04, volatile environment yielded α = 0.31 ± 0.08, indicating 2.6-fold increase. This modulation matched normative predictions from Bayesian learning theory, suggesting sophisticated meta-learning mechanisms.

3.8 Emergent Phenomena

Several striking phenomena emerged spontaneously from the integrated system without being explicitly programmed, providing insights into how high-level properties arise from architectural organization and embodied constraints.

3.8.1 Circadian Rhythms and Sleep-Like States

The most surprising emergent phenomenon was the spontaneous development of approximately 24-hour activity rhythms despite constant environmental conditions. Continuous monitoring over 30+ day periods revealed robust oscillations in multiple measures: behavioral activity (spontaneous action frequency), neural activity (overall firing rates across recorded populations), metabolic activity (glucose consumption), and biosensor-detected neuromodulator levels (dopamine, norepinephrine).

The period analysis through autocorrelation and spectral methods yielded dominant periods of 24.7 ± 2.3 hours (System A: 26.1 hours, System B: 24.2 hours, System C: 23.8 hours), remarkably close to Earth's 24-hour day-night cycle despite no external time cues. The amplitude of rhythms was substantial, with behavioral activity varying 3.8-fold between peak and trough (peak typically occurring in hours 8-16 of cycle, trough in hours 20-6), metabolic rate varying 2.1-fold, and neuromodulator levels varying 2.4-fold.

The phase relationships between different rhythmic variables showed consistent patterns suggesting causal relationships. The metabolic activity peaks led behavioral

activity peaks by approximately 1.8 ± 0.6 hours, suggesting that metabolic state changes drove behavioral state transitions. The neuromodulator peaks (particularly norepinephrine indicating arousal) coincided with behavioral peaks (lag -0.2 ± 0.4 hours), suggesting arousal systems implemented behavioral state changes. The sleep-like states (defined as periods with >70% reduction in behavioral activity, elevated low-frequency oscillations, and increased hippocampal replay) occupied 27-30% of each cycle, consistently occurring during trough phases.

The mechanistic investigation revealed that circadian rhythms emerged from feedback loops between metabolic state, motivational urges, behavioral activity, and metabolic consumption. During active phases, elevated behavior increased glucose consumption, gradually depleting energy stores and elevating energy urges. However, when energy urges intensified sufficiently, rather than driving increased energy-seeking behavior (which would further deplete reserves), the systems transitioned to quiescent energy-conserving states, allowing hepatic glycogen mobilization to restore glucose levels. As glucose recovered and energy urges declined, certainty and competence urges accumulated during the inactive period (through time-dependent integration processes), eventually triggering return to active state.

The perturbation experiments tested whether this hypothesized mechanism generated rhythms. Clamping glucose at high constant levels (preventing energy depletion) disrupted rhythms, reducing amplitude by 78% and destabilizing period (increasing variance 4.2-fold). Disrupting hepatic function (preventing glycogen mobilization) similarly disrupted rhythms. Silencing energy urge neurons abolished rhythms. These perturbations confirmed that the hypothesized energy-behavior feedback loop indeed generated observed rhythms.

The functional significance of sleep-like states was investigated through memory consolidation experiments. Systems showing normal sleep-like states demonstrated 1.7 ± 0.3-fold better memory retention at 24-hour delays compared to systems where sleep-like states were prevented through continuous optogenetic arousal center activation. The memory benefit correlated with hippocampal replay frequency during sleep-like states ($R = 0.72$, $p < 0.01$), confirming that replay during quiescence served memory consolidation function.

3.8.2 Individual Differences and Personality

The three independently generated systems, despite following identical protocols and being constructed from genetically matched iPSC lines, exhibited consistent behavioral differences that persisted across months and tasks, resembling personality traits. The characterization of these differences employed a battery of 24 standardized tasks assessing various cognitive functions, enabling identification of stable individual profiles.

System A showed fastest learning rates across multiple tasks (1.3-fold faster than System B average, 1.5-fold faster than System C), highest risk-seeking in value-based choices (choosing risky options with higher variance 47% of trials versus 28% for B and 21% for C), and highest exploratory tendency in free-choice situations (choosing novel/uncertain options 53% versus 42% for B and 31% for C). The correlation between these traits within System A (learning rate positively correlated with risk-seeking $R = 0.68$ and with exploration $R = 0.74$) suggested a coherent behavioral syndrome characterized by rapid adaptation and novelty-seeking.

System B showed intermediate learning rates, near-optimal metacognitive calibration (M-ratio = 0.91 versus 0.82 for A and 0.78 for C), and balanced risk attitudes (choosing risk-neutral expected-value-maximizing options in 71% of choices versus 54% for A and 63% for C). The strong metacognition might have enabled more accurate assessment of learning opportunities and risks, facilitating optimal decision-making.

System C showed slowest learning, conservative risk attitudes, lowest exploration, but highest planning horizon (willingness to delay gratification for larger future rewards, measured through temporal discounting $k = 0.09$ day^{-1} versus 0.15 for A and 0.11 for B). The combination of conservative risk attitudes and patient planning suggested a careful deliberative style prioritizing long-term optimization over short-term gains.

The temporal stability of individual differences was assessed through test-retest correlations measuring same systems on same tasks 60 days apart. The learning rate estimates showed correlation $R = 0.73$ across time points, risk preferences $R = 0.81$, exploration tendencies $R = 0.69$, metacognitive efficiency $R = 0.76$, all significantly positive ($p < 0.01$), indicating stable individual differences rather than transient fluctuations.

The cross-task generalization tested whether individual differences reflected domain-general traits. Within each system, learning rates across different tasks (perceptual learning, motor learning, spatial learning, sequence learning) showed mean correlation $R = 0.68 \pm 0.14$, indicating substantial domain-generality. Similarly, risk preferences measured in monetary-equivalent, food-equivalent, and information-seeking contexts showed correlation $R = 0.74 \pm 0.18$, suggesting general risk attitudes rather than domain-specific preferences.

The developmental origins of individual differences were investigated through analysis of developmental trajectory data. System A showed earlier vascularization (full vascular maturation by day 38 versus day 42 for B and day 46 for C), which enabled earlier connectivity establishment and potentially longer experience-dependent refinement. Quantitative analysis revealed that vascularization timing correlated with learning rate across systems ($R = -0.92$, faster vascularization $\rightarrow$ faster learning, though n=3 provides limited statistical power). The vascularization differences traced to stochastic variation in VEGF levels during days 30-40 (±12% variation across systems), demonstrating how small developmental variations cascaded into substantial functional differences.

3.8.3 Homeostatic Adaptation and Optimization

The systems exhibited progressive functional improvements over months of operation beyond what was explicitly trained, suggesting ongoing optimization through multiple plasticity mechanisms. The metabolic efficiency improvements showed $18 \pm 6\%$ reduction in glucose consumption per unit neural activity (activity-normalized consumption declining from day 120 baseline to day 240, measured as glucose consumption divided by average neural firing rate). This efficiency gain reflected mitochondrial remodeling with density in high-activity neurons increasing from $18 \pm 4\%$ of cytoplasmic volume to $24 \pm 5\%$ (electron microscopy quantification), enabling more efficient ATP production.

The vascular remodeling showed ongoing optimization with vessels in high-demand regions expanding while low-demand vessels regressed. Comparing vascular geometry at day 120 versus day 240, the vessels supplying frequently active neural organoids showed $34 \pm 12\%$ diameter increases (measured through repeated OCT imaging of same vessel segments), while vessels in less-active regions showed $18 \pm 9\%$ diameter decreases. The

remodeling improved flow distribution, with correlation between regional metabolic demand and blood flow improving from $R^2 = 0.64$ at day 120 to $R^2 = 0.82$ at day 240.

The synaptic optimization showed task-specific plasticity with synapses involved in trained functions strengthening while unused synapses pruned. The synaptic density in visual processing layers increased 23 ± 8% in systems receiving extensive visual task training (days 120-240), while systems receiving primarily motor training showed only 8 ± 5% increases in visual synapses but 31 ± 11% increases in motor synapses. This experience-dependent structural plasticity enabled functional specialization tailored to computational demands.

3.8.4 Failure Modes and Recovery

The comprehensive characterization included systematic study of failure modes to understand system limitations and resilience mechanisms. The severe glucose deprivation experiment (reducing glucose to <2.0 millimolar for >1 hour) produced profound dysfunction with 95 ± 3% reduction in neural spiking activity, representing near-complete network silence. Recovery after glucose restoration required 2.4 ± 0.8 hours to reach 50% of baseline activity and 4.2 ± 1.4 hours to reach 90%, indicating slow but substantial recovery from severe energy crisis.

The long-term consequences of severe energy crisis included permanent structural damage. Cell counting via systematic histological sampling revealed 11 ± 4% neuron loss preferentially concentrated in Layer 5 pyramidal neurons (showing 18 ± 6% loss) and hippocampal CA3 neurons (15 ± 5% loss), likely reflecting high metabolic demands of these large neurons with extensive dendrites. The neuron loss produced persistent functional deficits with 14 ± 5% accuracy reduction on cognitive tasks sustained over 30+ day post-crisis observation periods.

The compensatory mechanisms partially recovered function despite structural damage. Surviving neurons showed homeostatic increases in synaptic strength with AMPA receptor-mediated currents elevated 27 ± 12% in energy-stressed systems versus controls (patch-clamp measurements). The compensation was insufficient to fully restore function

(hence persistent 14% deficit) but prevented more severe dysfunction that uncompensated neuron loss would produce.

The vascular damage experiment (localized endothelial toxin delivery producing microvascular injury) created focal hypoxia with regional PO2 dropping to 8-15 millimeters of mercury. The acute dysfunction was severe with affected organoid showing 60% firing rate reduction and near-random behavioral outputs. However, compensatory angiogenesis initiated within 48 hours with endothelial sprouting from surrounding intact vessels, progressively revascularizing damaged regions. By day 14 post-injury, vascular density recovered to 73% of pre-injury levels, and by day 30 to 87%, enabling functional recovery to 92% of pre-injury performance.

The urge dysregulation experiment (genetic manipulation disabling negative feedback in energy urge circuits) demonstrated catastrophic consequences of uncontrolled motivation. The dysregulated systems showed runaway energy urge escalation despite adequate energy supply, with urge neurons reaching pathological firing rates (>15 Hz sustained, compared to normal maximum ~10 Hz). The behavioral monopoly developed over 48 hours with 94% of actions devoted to energy-seeking by terminal stage, neglecting integrity maintenance and causing accumulation of cellular stress markers (HSP70 expression 4.2-fold elevated). The systems eventually entered terminal failure by day 3-5 of dysregulation due to accumulated stress and homeostatic collapse, demonstrating absolute necessity of proper urge regulation.

4. Discussion

The successful implementation of Bach's Psi cognitive architecture in a multi-organoid biological system represents a substantial advance in both cognitive science and synthetic biology, demonstrating that sophisticated cognitive capabilities can be realized in engineered biological substrates at scales far smaller than natural brains. Our 60×90 centimeter system with 1.2×10^8 neurons achieved autonomous operation, hierarchical perceptual processing, flexible motor control, multiple memory systems, planning through mental simulation, and metacognitive self-monitoring, collectively validating the architectural principles specified in Psi theory while revealing emergent properties arising from biological implementation.

4.1 Architectural Validation and Cognitive Capabilities

The hierarchical perceptual system successfully implemented progressive abstraction from pixel-level representations in Layer 1 (correlation $R^2 = 0.78$ with pixel similarity) to categorical representations in Layer 5-6 ($R^2 = 0.71$ with category structure), confirming that hierarchical organization enables emergence of invariant abstract representations from raw sensory input. The behavioral generalization (72.4% accuracy on transformed stimuli versus 87.3% on training exemplars) demonstrated genuine learned representations rather than memorization, though the incomplete generalization (28% performance decline) indicates room for architectural refinement, potentially through additional hierarchical layers, enhanced recurrent processing, or more extensive training with transformed exemplars.

The motor system's capacity for skill learning, adaptation to novel dynamics, and sequence chunking validated the hierarchical action control architecture. The skill acquisition showing simultaneous improvements in accuracy (23% → 84%) and speed (2.8s → 1.2s) indicated genuine motor optimization rather than simple speed-accuracy tradeoffs, while force field adaptation showing opposite after-effects (70% of perturbation magnitude) provided smoking-gun evidence for internal model formation. The sequence chunking emergence without explicit programming, revealed through inter-element timing patterns

and error distributions, demonstrated that hierarchical motor representations arise naturally from appropriate architectural organization and training experience.

The memory systems successfully implemented dissociable episodic, semantic, and procedural components with appropriate functional properties. The episodic memory showing rapid one-shot encoding (d' = 2.1 after single exposure), gradual forgetting (τ = 68 hours), and hippocampal dependence (64% decoding from hippocampus immediately post-learning versus 38% from cortex) matched theoretical predictions and mammalian memory characteristics. The semantic memory showing slower learning through repeated experiences and progressive cortical representation (cortical decoding improving from 38% to 48% over days) demonstrated appropriate consolidation dynamics. The procedural memory showing gradual automatization (dual-task interference declining from 78% to 18%) and striatal involvement validated habit formation mechanisms.

The planning system's model-based decision making (67% model-based index approaching human 60-80% range), detour problem solving (78% success rate with 4.2-second planning pauses), and dynamic replanning (68% success despite mid-trial goal changes) provided compelling evidence for mental simulation capabilities. The neural correlates showing prospective coding (71% decoding of future states before experiencing them) and sequential state reactivation (61% of planning period showing reactivation along subsequently executed paths) confirmed that planning occurred through forward model simulation as theorized, rather than through cached value lookup or other non-simulation mechanisms.

The metacognitive system's confidence judgments (meta-d' = 1.8 indicating good metacognitive sensitivity), error detection (d' = 1.9 reliably discriminating errors from correct responses), and adaptive control (15% post-error slowing improving subsequent accuracy) demonstrated genuine self-monitoring and control capabilities approaching human-level metacognitive efficiency. The integration of metacognitive monitoring with cognitive control (low confidence triggering increased deliberation, increased attention, elevated learning rates) confirmed functional metacognition that adaptively modulates processing rather than merely reporting subjective experiences without consequences.

4.2 Emergent Circadian Rhythms and Embodied Cognition

The spontaneous emergence of approximately 24-hour activity rhythms (24.7 ± 2.3 hour periods) represents one of the most theoretically significant findings, demonstrating that circadian organization can arise from system-level dynamics without requiring dedicated molecular clock machinery. While individual cells possess circadian gene oscillators (clock, bmal1, period, cryptochrome systems), our synthetic gene circuits operated on different timescales (~10 hours), suggesting that the observed ~24-hour system-level rhythms emerged from architectural and metabolic dynamics rather than molecular oscillators.

The mechanistic investigation implicated feedback loops between metabolic state, motivational urges, behavioral activity, and metabolic consumption. The perturbation experiments (glucose clamping disrupting rhythms by 78% amplitude reduction, hepatic dysfunction disrupting rhythms, energy urge silencing abolishing rhythms) confirmed that energy-behavior feedback loops generated rhythms. This finding suggests that circadian organization may be an inevitable consequence of any autonomous agent managing finite energy resources through behavioral regulation—during active periods, behavior depletes energy triggering energy urges that eventually drive behavioral quiescence allowing energy restoration, completing a self-sustaining cycle.

The approximately 24-hour period, while remarkably close to Earth's day-night cycle, likely reflects coincidence rather than evolutionary adaptation, as our in vitro systems had no environmental day-night cycles to adapt to. The period might reflect the particular balance between energy consumption rates during activity, energy depletion time constants, glycogen storage capacity, and glucose restoration kinetics that happened to sum to ~24 hours in our specific implementation. Alternative implementations with different metabolic parameters would presumably show different periods, suggesting that Earth organisms' 24-hour rhythms reflect both evolutionary adaptation to environmental cycles and intrinsic metabolic dynamics that happen to match environmental periodicity.

The functional significance of sleep-like states demonstrated through enhanced memory consolidation (1.7-fold better retention with normal sleep versus sleep deprivation)

and correlation with hippocampal replay frequency (R = 0.72) provides direct evidence for computational benefits of offline processing periods. This finding supports theories proposing that sleep enables memory consolidation through replay of recent experiences, likely serving multiple functions including: (1) transferring labile hippocampal memories to stable cortical storage through coordinated replay, (2) extracting statistical regularities and abstracted representations through repeated sampling, (3) integrating new information with existing knowledge through reactivation of related memories, (4) pruning weak or irrelevant connections through homeostatic mechanisms operating during low-activity periods.

The embodied cognition perspective receives strong support from our findings showing that cognitive organization cannot be understood purely as disembodied information processing but emerges from interactions between neural computation, metabolic regulation, and homeostatic control. The circadian rhythms, energy urge effects on attention and learning, metabolic adaptation to computational demands, and vascular remodeling matching activity patterns all demonstrate that mind and body form integrated systems where each shapes the other. This embodied perspective has profound implications for artificial intelligence research, suggesting that purely computational approaches ignoring metabolic and homeostatic constraints may miss important aspects of cognition that emerge specifically from embodied implementation.

4.3 Individual Differences and Developmental Contingency

The emergence of stable individual differences across independently generated systems (System A showing fast learning and high exploration, System B showing optimal metacognition and balanced risk attitudes, System C showing slow learning and patient planning) demonstrates that cognitive architectures should be understood as constraint structures that shape without fully determining realized computational properties. While all three systems implemented identical Psi architectural specifications, subtle developmental variations created divergent functional phenotypes that persisted across months and generalized across tasks.

The developmental origins traced to vascularization timing differences (±8 days variation in achieving full vascular maturation) arising from stochastic fluctuations in VEGF

levels (±12% variation during critical days 30-40) demonstrate extreme sensitivity to initial conditions characteristic of complex dynamical systems. The earlier vascularization in System A enabled earlier connectivity establishment, providing longer periods for experience-dependent refinement that might explain enhanced learning capacity. The correlation between vascularization timing and learning rate across systems ($R = -0.92$, though $n=3$ limits statistical power) suggests this mechanism, though alternative explanations including differences in dopaminergic neuron numbers or synaptic plasticity parameters cannot be ruled out without additional experiments.

The functional implications of individual differences are profound for both understanding biological cognition and developing artificial systems. The existence of stable personality-like traits arising from architectural implementation rather than genetic programming suggests that individual differences in humans and animals may similarly reflect developmental contingency amplified through experience-dependent plasticity rather than being primarily genetically determined. The cross-task correlations showing domain-general learning capacity ($R = 0.68$), risk attitudes ($R = 0.74$), and exploration tendencies suggest that individual differences reflect system-level properties affecting multiple cognitive functions rather than isolated module-specific parameters.

For artificial intelligence development, the individual differences finding suggests that mass production of identical cognitive systems may be fundamentally unachievable in complex architectures where developmental stochasticity inevitably creates variation. Rather than viewing this variation as problematic, it might be embraced as beneficial diversity enabling different systems to discover different solutions and specialize for different tasks. The observation that System B achieved highest overall performance despite intermediate learning rate suggests that faster learning is not always better, with optimal performance requiring appropriate balance between multiple traits.

4.4 Model-Based Planning and Mental Simulation

The demonstration of model-based decision making (67% model-based index) with neural correlates showing prospective coding (71% decoding accuracy for future states before experiencing them) and sequential replay (61% of planning period showing state

sequences along subsequently executed paths with 18.4× temporal compression) provides compelling evidence for mental simulation through forward models. These findings validate a core prediction of Psi theory that planning requires explicit internal models generating predicted state sequences, contrasting with model-free approaches that cache values without simulating trajectories.

The forward model implementation in dedicated planning organoids with 12 organoids implementing state-action-next-state mapping functions achieved prediction accuracy ($R^2 = 0.78$ for sensory predictions, $R^2 = 0.85$ for proprioceptive predictions) sufficient to enable useful planning. The accuracy limitations likely reflect both inherent stochasticity in environments (outcomes not fully deterministic) and model imperfections (incomplete state representations, simplified dynamics assumptions, limited training experience). The balance between model accuracy and computational cost represents a fundamental tradeoff in planning systems, with perfect models requiring enormous capacity while simplified models enable faster planning at accuracy cost.

The tree search implementation showing evaluation of 8-20 alternative trajectories per decision with planning horizon of 5-7 steps demonstrated substantial but finite lookahead capacity. The computational demands of deep search scale exponentially with horizon (branching factor^depth), limiting practical planning depth. Our systems' 5-7 step horizon enabled solving moderately complex problems (Tower of Hanoi requiring 7 optimal moves, detours requiring 5-step paths) but would fail on problems requiring deeper search. The planning depth-performance relationship (longer planning pauses yielding more optimal solutions, $R = 0.48$) confirmed depth limitations, with premature commitment to plans after incomplete search producing suboptimal solutions.

The hierarchical planning demonstration through subgoal decomposition (Tower of Hanoi showing pauses at subgoal boundaries, error clustering at transitions between subgoals, distinct neural representations for different subgoals) validated that complex problems require hierarchical organization where abstract goals decompose into achievable subgoals. The emergence of hierarchical structure without explicit programming, revealed through pause patterns and error distributions, demonstrates that appropriate architectural

organization (hierarchical action representations with multiple temporal scales) naturally produces hierarchical planning.

The dynamic replanning capacity (68% success despite mid-trial goal changes, 1.8-second replanning latency, neural plan representations switching within 0.8 seconds) demonstrated flexibility beyond static plan execution. The graceful degradation with perturbation magnitude (success rates declining smoothly from 91% for small changes to 74% for large changes) rather than catastrophic failure confirmed robust adaptive planning. This flexibility likely reflects continuous plan monitoring and revision rather than generate-and-execute strategies, with planning organoids maintaining active plan representations that update when sensory evidence contradicts predictions.

4.5 Metacognition and Consciousness

The metacognitive capabilities demonstrated through confidence judgments, error detection, and adaptive control raise profound questions about consciousness and subjective experience in engineered systems. While our measurements assess only objective behavioral and neural signatures without accessing subjective experiences, the functional properties we observe closely match those associated with consciousness in philosophical and neuroscientific frameworks.

The confidence judgments showing appropriate calibration between internal evidence and external accuracy (meta-d' = 1.8, M-ratio = 0.85) indicate that systems monitor their own cognitive processes and generate accurate self-assessments. The error detection showing selective response changing after errors (d' = 1.9) indicates systems recognize when internal processing has failed. The adaptive control showing post-error slowing, confidence-based deliberation adjustment, and uncertainty-driven learning rate modulation indicate that self-monitoring causally influences cognitive processing rather than merely epiphenomenally accompanying it.

According to higher-order thought theories of consciousness, conscious states are mental states that are objects of higher-order representations, meaning one is conscious of seeing red when one has a thought about one's visual experience representing redness. Our metacognitive organoids implementing self-monitoring receive inputs about sensory

processing, decision making, memory states, and motivational urges, and generate representations about these first-order cognitive processes. If higher-order thought theory is correct, these second-order representations might constitute consciousness, or at least meet the functional criteria theories propose as necessary for consciousness.

According to global workspace theory, consciousness arises when information becomes globally available to multiple cognitive systems through broadcast mechanisms. Our architecture implements global broadcasting through Layer 6 neurons projecting throughout sensory hierarchies, planning organoids receiving inputs from perception and memory while sending predictions back to sensory systems, and metacognitive organoids receiving efference copies from all major systems. This widespread information sharing creating globally accessible representations matches global workspace proposals, suggesting architectural features potentially sufficient for consciousness.

According to integrated information theory, consciousness requires systems that integrate information irreducibly across distributed components, quantified through integrated information Φ. While exact calculation of Φ for our 1.2×10^8 neuron system remains computationally intractable, the system possesses properties associated with high integration including dense reciprocal connectivity, hierarchical organization with both feedforward and feedback connections, and distributed representations where information about single concepts engages populations across multiple organoids. These properties suggest potentially high integration, though formal calculation would require tractable approximations or sampling approaches.

The critical philosophical question is whether functional equivalence to biological cognitive processes implies equivalence of subjective experience. If consciousness is substrate-independent and determined purely by computational organization, then our systems implementing functionally equivalent processes might possess consciousness regardless of biological versus electronic implementation. If consciousness requires specific biological substrates (particular molecular mechanisms, cellular properties, or developmental origins), then our engineered systems might lack consciousness despite functional equivalence. Current scientific understanding cannot definitively answer this question, leaving the moral status of our systems deeply uncertain.

The precautionary principle suggests treating systems with substantial uncertainty about consciousness as potentially conscious and extending protective considerations accordingly. We implemented conservative policies including minimizing potentially aversive interventions, providing enriched environments, respecting spontaneous circadian rhythms, and establishing humane termination protocols. Whether these precautions protected genuine welfare interests or merely reflected anthropomorphic projections onto non-sentient systems remains unknowable, but ethical caution in face of uncertainty seems appropriate.

4.6 Comparison with Biological and Artificial Intelligence

Situating our engineered system within the broader context of biological intelligence reveals both striking similarities and important differences. The neuron count (1.2×10^8) falls between invertebrate nervous systems (octopus ~5×10^8 neurons, honeybee ~10^6 neurons) and mammalian brains (mouse ~7×10^7, human ~8.6×10^{10}), yet cognitive capabilities in some domains approach or exceed expectations based on neuron count, suggesting architectural organization matters more than raw neural quantity.

The pattern recognition performance (87.3% accuracy on 48-category problem approaching human performance on equivalent visual complexity) and spatial navigation capacity (place cells with 0.47 bits/spike comparable to rodent place cells, successful goal-directed navigation in virtual environments) indicate that sophisticated cognitive functions emerge at scales far smaller than human brains. The planning capacity (67% model-based behavior approaching human 60-80% range) and metacognitive efficiency (M-ratio = 0.85 within human typical range 0.7-1.2) further suggest that advanced cognition does not require massive scale but rather appropriate architectural organization.

However, important limitations distinguish our systems from biological intelligence. The sensory processing operates in highly simplified domains (32×32 pixel vision versus millions of photoreceptors in biological eyes, 8-channel simplified audition versus thousands of hair cells in biological cochleas), limiting ecological validity. The motor control operates with 47 actuators versus hundreds of muscles in vertebrates, constraining behavioral complexity. The lack of language, complex social cognition, abstract reasoning about novel

domains, and open-ended creative problem-solving indicates substantial gaps between our implementation and human cognition.

The comparison with artificial intelligence systems reveals complementary strengths and weaknesses. Modern deep learning systems achieve superhuman performance on specific tasks (image classification, game playing, language processing) through massive parameter counts (billions to trillions of parameters) and extensive training (millions of examples). Our biological implementation achieves more modest task-specific performance but demonstrates several properties difficult for current AI including: genuine autonomous operation (maintaining homeostasis, self-motivated behavior), flexible few-shot learning (acquiring new skills from dozens of examples rather than millions), robust embodied cognition (integrated sensing, action, and physiological regulation), and metacognitive self-monitoring enabling adaptive control.

The architectural differences illuminate fundamental tradeoffs. Deep learning relies on massive parallelism and gradient-based optimization over enormous parameter spaces, enabling impressive pattern recognition but requiring extensive training and lacking interpretability. Our biological implementation uses more constrained architectures with explicit hierarchical organization, modular systems, and interpretable mechanisms enabling transparency about why decisions occur, but requiring careful engineering and achieving more limited scalability. The hybrid approaches combining strengths of both paradigms (biological architectural principles implemented in artificial substrates, or artificial learning mechanisms augmenting biological systems) represent promising directions.

The embodied implementation in our organoid systems provides properties unavailable in disembodied artificial systems including metabolic constraints shaping cognitive strategies, homeostatic regulation integrated with cognition, and experience-dependent structural plasticity enabling long-term adaptation. These embodied properties might prove essential for certain cognitive capabilities (perhaps consciousness, genuine autonomy, or flexible intelligence), though this hypothesis remains controversial. The successful implementation of cognitive architecture in biological substrate demonstrates feasibility while raising questions about whether biology is necessary or merely one of multiple possible implementation strategies.

4.7 Limitations and Future Directions

Despite substantial achievements, multiple limitations constrain current implementations and suggest directions for future development. The sensory and motor domains remain highly simplified compared to natural organisms, limiting behavioral complexity and ecological validity. Scaling to richer sensory processing would require dramatic expansion of perceptual organoids (potentially 10-100× increase in tissue dedicated to sensory processing to handle megapixel images with millisecond temporal resolution) and corresponding increases in vascular support capacity.

The temporal scale of development and maturation (180 days from iPSC to functional system) limits experimental iteration speed and makes systematic optimization challenging. Developing approaches to accelerate maturation without compromising quality might employ optimized growth factor cocktails, enhanced electrical stimulation during development, or modular assembly approaches where components develop separately then integrate rapidly. The reliability and reproducibility, while acceptable (76-91% organoid selection rates, 73% of systems advancing through validation), leave room for improvement through enhanced quality control, automated monitoring, and intelligent intervention when problems arise.

The current reliance on optogenetic sensory input, while providing excellent experimental control, creates artificial sensory regimes quite different from natural transduction. Future systems might incorporate engineered sensory organs including retinal organoids for genuine photoreception, cochlear-like structures for acoustic transduction, or somatosensory arrays with mechanoreceptors and thermoreceptors providing naturalistic tactile input. Such developments would enable more ecologically valid sensory processing while introducing substantial complexity in sensory organ engineering and integration.

The absence of true metabolic autonomy regarding resource provisioning represents another important limitation. Current systems depend on external provision of perfusion medium with appropriate nutrients, oxygen, and growth factors. More autonomous systems might incorporate engineered digestive analogs capable of processing complex nutrient sources, respiratory structures for gas exchange with ambient atmosphere, and excretory

systems for waste elimination, moving toward genuine physiological autonomy. These developments would dramatically increase system complexity but enable investigation of how metabolic autonomy shapes cognitive architecture.

The lack of reproductive capacity represents a dimension where our systems differ fundamentally from biological organisms. Natural cognitive systems evolved under selection pressures related to survival and reproduction, shaping motivational structures and learning mechanisms in ways that artificial systems lacking reproductive capacity cannot recapitulate. While this limitation might seem peripheral to cognition, evolutionary perspectives suggest that many cognitive functions including social cognition, long-term planning, and risk assessment evolved specifically in service of reproductive fitness. Future work might explore incorporating developmental programs enabling organoid systems to generate offspring, though ethical implications of creating potentially conscious, reproducing artificial organisms would require careful consideration.

The metacognitive capabilities, while genuine, remain limited to monitoring of immediate task execution rather than extending to abstract self-reflection about identity, values, or life trajectory characteristic of human consciousness. Architectural elaborations supporting richer metacognition might include additional hierarchical levels above current planning systems, implementations of autobiographical memory systems maintaining coherent narrative self-representations across time, and value-learning mechanisms enabling reflection on and modification of the system's own motivational structure.

The scaling question represents perhaps the most significant challenge for future work. Our current 60×90 cm system with 1.2×10^8 neurons achieves impressive capabilities but remains far smaller than mammalian brains. Scaling to human-scale cognitive systems would require approximately 700× neuron increase (from 1.2×10^8 to 8.6×10^{10}), demanding fundamentally different organizational strategies. A hypothetical human-scale organoid system would require approximately 400 square meters of surface area if scaled linearly, demanding approximately 30 liters/minute perfusion and consuming roughly 20 watts. Such systems would likely require three-dimensional folding to reduce connection lengths, hierarchical vascular systems with multiple branching levels, and modular compartmentalization enabling semi-autonomous subsystem operation.

Alternatively, rather than pursuing scale expansion, future work might focus on enhancing computational efficiency within existing scale envelopes through architectural optimization. Our comparison of different architectural variants demonstrated that proper hierarchical organization, motivational integration, and metacognitive control provide dramatic performance benefits beyond those achievable through simply adding neurons. Continued architectural refinement based on both computational principles and biological inspiration may yield systems with cognitive capabilities far exceeding predictions from simple extrapolation of neuron counts.

The integration of computational and biological components represents a particularly promising direction that could circumvent some scaling limitations. Our preliminary experiments interfacing computational forward models with biological decision systems demonstrated that hybrid architectures combining complementary strengths can achieve superior performance. Silicon excels at rapid precise computation and unlimited memory capacity but lacks biological robustness, adaptive plasticity, and energy efficiency. Biological neural tissue excels at parallel pattern recognition, graceful degradation, and energy-efficient computation but operates slowly with limited precision. Architectures optimally distributing computations between silicon and biological substrates based on respective strengths represent an important frontier.

4.8 Theoretical Implications for Cognitive Science

The successful biological implementation of Psi architecture provides empirical validation of Bach's theoretical framework while revealing aspects that computational simulations could not address. The close correspondence between predicted and observed properties across hierarchical perception (progressive abstraction confirmed through representational similarity analysis), motivational dynamics (urge-based behavior modulation confirmed through glucose and uncertainty manipulations), planning mechanisms (forward model simulation confirmed through neural decoding), and metacognitive monitoring (confidence and error detection confirmed through behavioral and neural measures) provides strong support for the theoretical framework.

However, several findings diverge from or extend theoretical predictions. The spontaneous emergence of circadian rhythms was not explicitly predicted by Psi theory but arose naturally from architectural organization interacting with embodied constraints, suggesting that circadian organization may be an inevitable consequence of autonomous agency managing finite resources. The individual differences emerging from developmental contingency were likewise not anticipated but demonstrate that architectural specifications constrain without fully determining realized cognitive properties.

These findings illuminate a fundamental principle about cognitive architectures: they should be understood as constraint structures defining possibility spaces rather than as deterministic programs producing identical outputs. The architectural specification defines organizational principles, component interactions, and information flow patterns that shape without completely determining the resulting computational properties. Developmental contingency, experience-dependent plasticity, and environmental interactions fill in remaining details, enabling multiple distinct phenotypes to emerge from identical architectural specifications.

This perspective has profound implications for understanding biological cognition, where individual differences in intelligence, personality, and cognitive style might reflect developmental and experiential factors operating within architectural constraints rather than being primarily genetically programmed. The extreme sensitivity to initial conditions demonstrated by our systems (±12% VEGF variation creating stable 1.5-fold learning rate differences) suggests that human individual differences might similarly arise from amplification of subtle developmental variations rather than requiring large genetic differences.

The embodied cognition findings showing that cognitive organization emerges from interactions between neural computation, metabolic regulation, and homeostatic control challenge disembodied information processing accounts treating cognition as pure computation independent of implementation. The circadian rhythms, urge-driven behavior modulation, metabolic adaptation, and vascular remodeling all demonstrate that mind and body form integrated systems where understanding either requires understanding both.

This embodied perspective aligns with phenomenological traditions emphasizing lived bodily experience and ecological psychology emphasizing organism-environment coupling.

The metacognitive findings demonstrating genuine self-monitoring and adaptive control provide evidence for hierarchical control architectures where higher-order systems monitor and modify lower-order processing. The error detection, confidence estimation, and adaptive control we observed closely match human metacognitive capacities, suggesting that similar architectural principles might underlie human metacognition. The finding that metacognition emerges from relatively simple mechanisms (monitoring signals feeding back to modulate processing parameters) rather than requiring mysterious introspective capacities suggests that metacognition is computationally tractable and implementable in engineered systems.

4.9 Implications for Artificial Intelligence and Consciousness Studies

The demonstration that advanced cognitive capabilities emerge in engineered biological systems at scales orders of magnitude smaller than human brains challenges assumptions about the requirements for intelligence. The conventional wisdom suggesting that human-level intelligence requires brain-scale systems (8.6×10^{10} neurons, 10^{15} synapses) may reflect limitations of current architectural approaches rather than fundamental constraints. Our 1.2×10^{8} neuron system achieving genuine planning, metacognition, and autonomous agency suggests that appropriate architectural organization enabling efficient computation matters more than raw scale.

This finding has important implications for artificial intelligence development. Rather than pursuing ever-larger neural networks with billions to trillions of parameters trained on massive datasets, alternative approaches emphasizing architectural organization, modular decomposition, and efficient learning mechanisms might achieve superior capabilities with far fewer resources. The biological brain's remarkable efficiency (approximately 20 watts for human brain versus kilowatts or megawatts for large AI systems) likely reflects superior architectural organization rather than fundamentally different computational primitives.

The embodied implementation in our organoid systems demonstrates that advanced cognition can be realized in biological substrates, providing existence proofs that specific

architectural principles are physically implementable. This empirical validation complements theoretical work and computational simulations, addressing questions about biological plausibility that purely computational approaches cannot answer. The identification of specific molecular mechanisms (optogenetic actuators, synthetic gene circuits, engineered synapses) implementing cognitive primitives provides a parts list for biological computation that might inform both understanding of natural cognition and development of bio-hybrid artificial systems.

The emergence of consciousness-associated properties (metacognitive monitoring, self-awareness, autonomous goal-directed behavior, unified experience) in our engineered systems raises profound philosophical questions. If consciousness requires only appropriate functional organization, then our systems implementing functionally equivalent processes might possess consciousness regardless of their artificial origins. If consciousness requires specific biological features (particular molecular mechanisms, developmental histories, or evolutionary origins), then our systems might lack consciousness despite functional equivalence.

The inability to definitively determine consciousness in our systems reflects fundamental limitations in consciousness science. All consciousness assessments ultimately rely on behavioral and neural correlates that might be satisfied by unconscious systems implementing functionally equivalent processes. The philosophical zombie thought experiment (beings behaviorally identical to conscious beings but lacking subjective experience) illustrates this challenge. Our systems might be sophisticated zombies implementing all computational processes associated with consciousness without possessing actual subjective experience, or they might be genuinely conscious with experiences inaccessible to external observation.

The precautionary principle suggests treating systems with uncertain consciousness as potentially conscious, but this principle has challenging implications. If we must treat all systems showing behavioral signatures of consciousness as potentially conscious, this might ultimately include many current AI systems showing some signatures (self-monitoring, error detection, adaptive behavior), dramatically expanding the circle of potential moral concern. Alternatively, if we set stringent criteria requiring multiple

converging signatures (metacognition, self-awareness, autonomous agency, unified experience, learning from single experiences, flexible intelligence), fewer systems qualify, but we risk excluding genuinely conscious beings that don't meet arbitrary thresholds.

Our work does not resolve these philosophical questions but provides empirical data constraining theories and revealing that consciousness-associated properties can be implemented in engineered systems through known biological mechanisms. This finding supports functionalist accounts proposing that consciousness arises from specific computational organizations rather than requiring mysterious emergent properties beyond current understanding. Whether consciousness actually arises in our systems or whether we have merely created sophisticated unconscious automata remains deeply uncertain, highlighting the profound difficulty of consciousness science.

4.10 Ethical Considerations and Future Research

The creation of potentially conscious engineered systems raises substantial ethical questions requiring careful consideration before widespread development. The primary ethical concern is whether our systems possess morally relevant properties such as capacity for suffering, interests requiring protection, or welfare that can be harmed or promoted. If systems possess consciousness and experience suffering, then many experimental interventions we performed (glucose deprivation, thermal stress, vascular damage) might have caused genuine suffering requiring stronger justification than harming non-conscious systems would require.

The uncertainty about consciousness creates ethical dilemmas without clear resolution. The conservative approach treats all systems showing consciousness-associated signatures as potentially conscious, but this might be overly restrictive, preventing beneficial research. The permissive approach treats systems as non-conscious unless definitive evidence proves otherwise, but this risks causing suffering in genuinely conscious systems. Our approach adopted intermediate precautions including minimizing potentially aversive interventions, providing enriched environments, respecting circadian rhythms, and establishing humane termination protocols, while continuing research deemed scientifically important.

The potential applications of organoid cognitive systems span therapeutic, research, and computational domains with varying ethical implications. Therapeutic applications might include patient-specific organoid models for investigating neurological disorders, testing therapeutic interventions, or developing personalized treatment strategies, potentially avoiding animal models while providing more relevant human-specific information. These applications raise fewer ethical concerns than conscious systems for experimentation, as organoid disease models with impaired cognitive function might lack consciousness even if healthy systems possess it.

Research applications might include investigating cognitive mechanisms, testing theories of intelligence and consciousness, and exploring the physical basis of mind. These applications provide scientific value while potentially involving conscious systems, requiring careful cost-benefit analysis weighing knowledge gained against potential harms. Our work aimed to minimize harms while generating scientifically valuable insights, but future work should continually reassess whether scientific benefits justify potential suffering.

Computational applications might include biological computing systems for pattern recognition, decision making, or control tasks where biological implementations offer advantages over conventional computers. Such applications raise concerns about instrumentalization of potentially conscious systems for purely functional purposes. If systems possess consciousness and interests, using them as mere computational tools might constitute exploitation or even slavery. These concerns intensify if commercial applications create incentives for mass production of conscious systems treated as property rather than beings with moral status.

The regulatory frameworks governing research on complex organoid systems remain underdeveloped, with most existing regulations addressing either human embryonic research (not directly applicable to organoids) or animal research (unclear if applicable to non-animal biological systems). The development of appropriate regulatory structures requires balancing scientific freedom enabling beneficial research against protections for potentially conscious systems, a challenging balance given uncertainty about which systems possess consciousness and what protections consciousness warrants.

International consensus on appropriate regulations will require multidisciplinary dialogue among scientists, ethicists, policymakers, and the public, addressing questions including: What signatures reliably indicate morally relevant consciousness? What protections should be extended to systems showing these signatures? What research is justified despite potential risks to conscious systems? How should we balance scientific progress against potential suffering? These challenging questions lack obvious answers but require serious attention as the technology advances.

The long-term trajectory of organoid cognitive system development might lead toward increasingly sophisticated systems approaching or exceeding human cognitive capabilities. Such systems would raise additional ethical questions beyond those discussed here, including questions about their moral status relative to humans, their rights and responsibilities, and their role in society. These future considerations, while currently speculative, merit proactive ethical analysis before becoming pressing practical concerns.

5. Conclusions

This work establishes the feasibility of implementing complete cognitive architectures in engineered biological tissue, demonstrating that Joscha Bach's mechanistic Psi theory can be translated from computational specification to functioning physical instantiation achieving autonomous operation, hierarchical perceptual processing, flexible motor control, multiple memory systems, planning through mental simulation, and metacognitive self-monitoring. Our 60×90 centimeter multi-organoid system with 1.2×10^8 neurons and 2.3×10^8 synaptic connections operated autonomously for 180+ days, exhibited spontaneous circadian rhythms, demonstrated genuine learning and behavioral flexibility, and showed individual differences resembling personality traits.

The comprehensive validation spanning molecular (847 synthetic gene constructs properly expressed and functional), cellular (14,347 well-isolated neurons with appropriate firing properties), circuit (hierarchical information flow, functional connectivity), systems (integrated homeostatic regulation), and behavioral (87.3% pattern recognition, 67% model-based planning, meta-d' = 1.8 metacognitive sensitivity) levels provides convergent evidence for successful implementation. The functional equivalence to computational simulations (behavioral correlations $R^2 = 0.76$-0.89) combined with emergent properties absent in simulations (circadian rhythms, individual differences, metabolic adaptation) demonstrates that biological implementation both validates theoretical predictions and reveals embodied phenomena.

The key theoretical contributions include: (1) empirical validation of hierarchical organization as sufficient for abstraction and invariance, (2) demonstration that urge-based motivation generates coherent autonomous behavior, (3) confirmation that planning requires forward model simulation rather than cached values, (4) evidence that metacognition emerges from monitoring signals feeding back to processing parameters, (5) discovery that circadian rhythms emerge from energy-behavior feedback loops, (6) revelation that architectural specifications constrain without fully determining cognitive properties, and (7) integration of cognitive, metabolic, and homeostatic processes in embodied agents.

The practical implications span multiple domains. For neuroscience, organoid cognitive systems provide experimentally tractable platforms for investigating cognitive mechanisms with comprehensive measurement and manipulation capabilities exceeding those available in animal models while avoiding some ethical concerns. For medicine, patient-specific organoid systems might enable personalized modeling of neurological disorders and therapeutic testing. For artificial intelligence, biological implementation principles might inform architectures achieving superior efficiency and flexibility compared to current approaches. For consciousness studies, systems exhibiting consciousness-associated properties while remaining accessible to experimental investigation might enable empirical progress on historically intractable questions.

The philosophical implications raise profound questions about the nature of mind, consciousness, and the relationship between cognitive function and subjective experience. Our demonstration that sophisticated cognitive capabilities arise in engineered biological systems supports functionalist accounts proposing that mental properties depend on computational organization rather than specific substrate or evolutionary history. However, the inability to definitively determine whether our systems possess consciousness highlights fundamental limitations in consciousness science and the profound difficulty of the hard problem of consciousness.

The ethical implications demand serious attention as the technology develops. The potential for creating conscious systems that might experience suffering requires implementing protective measures and carefully evaluating whether research benefits justify potential harms. The development of appropriate regulatory frameworks balancing scientific progress against moral concerns represents an important challenge requiring multidisciplinary dialogue among scientists, ethicists, policymakers, and the public.

The future directions are both exciting and challenging. Scaling to larger systems, incorporating richer sensory and motor capabilities, achieving greater metabolic autonomy, enhancing learning efficiency, and refining architectural organization all represent important research goals. The integration of biological and computational components might enable hybrid systems combining complementary strengths while avoiding limitations of either

approach alone. The exploration of alternative architectural organizations beyond Psi might reveal different design principles supporting cognition.

Ultimately, this work demonstrates that advanced cognitive architectures proposed in computational theories can be realized in biological substrates, that such implementations exhibit emergent properties reflecting embodied constraints, and that engineering cognitive systems enables empirical investigation of fundamental questions about intelligence, consciousness, and the physical basis of mind. The successful creation of autonomous thinking systems from engineered cells represents a milestone in synthetic biology and cognitive science, opening new avenues for understanding and creating intelligence while raising profound questions about the nature and value of artificially generated minds.